



Regression 2

Industrial AI Lab.

Changyun Choi, Juhyeong Jeon

Linear Regression: Advanced

- Overfitting
- Regularization (Ridge and Lasso)

Overfitting: Start with Linear Regression

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline


# 10 data points
n = 10
x = np.linspace(-4.5, 4.5, 10).reshape(-1, 1)
y = np.array([0.9819, 0.7973, 1.9737, 0.1838, 1.3180, -0.8361, -0.6591, -2.4701, -2.8122, -6.2512]).reshape(-1, 1)

plt.figure(figsize=(10, 8))
plt.plot(x, y, 'o', label = 'Data')
plt.xlabel('X', fontsize = 15)
plt.ylabel('Y', fontsize = 15)
plt.grid(alpha = 0.3)
plt.show()
```

```
A = np.hstack([x**0, x])
A = np.asmatrix(A)
```

```
theta = (A.T*A).I*A.T*y
print(theta)
```

```
[[ -0.7774    ]
 [ -0.71070424]]
```



Recap: Nonlinear Regression

- Polynomial (here, quad is used as an example)

$$y = \theta_0 + \theta_1 x + \theta_2 x^2 + \text{noise}$$

$$\phi(x_i) = \begin{bmatrix} 1 \\ x_i \\ x_i^2 \end{bmatrix}$$


$$\Phi = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ \vdots & & \\ 1 & x_m & x_m^2 \end{bmatrix} \implies \hat{y} = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_m \end{bmatrix} = \Phi\theta$$

$$\implies \theta^* = (\Phi^T \Phi)^{-1} \Phi^T y$$

Nonlinear Regression

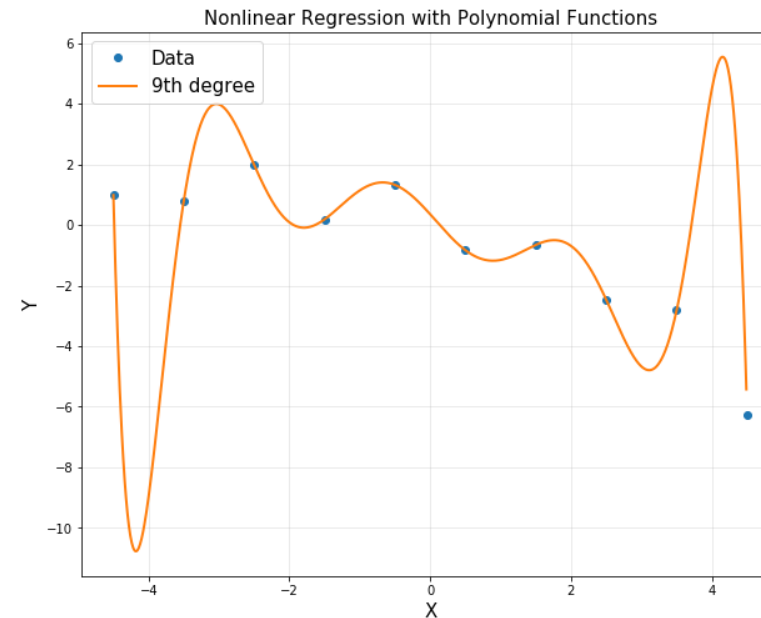
```
A = np.hstack([x**0, x, x**2])  
A = np.asmatrix(A)  
  
theta = (A.T*A).I*A.T*y  
print(theta)
```

```
[[ 0.33669062]  
 [-0.71070424]  
 [-0.13504129]]
```







```
A = np.hstack([x**i for i in range(10)])  
A = np.asmatrix(A)  
  
theta = (A.T*A).I*A.T*y
```

10 input points with degree 9 (or 10)



Polynomial Fitting with Different Degrees


Regression



Low error on input data points,
but high error nearby

Loss

- Loss: Residual Sum of Squares (RSS)



$$\min_{\theta} \|\hat{y} - y\|_2^2$$



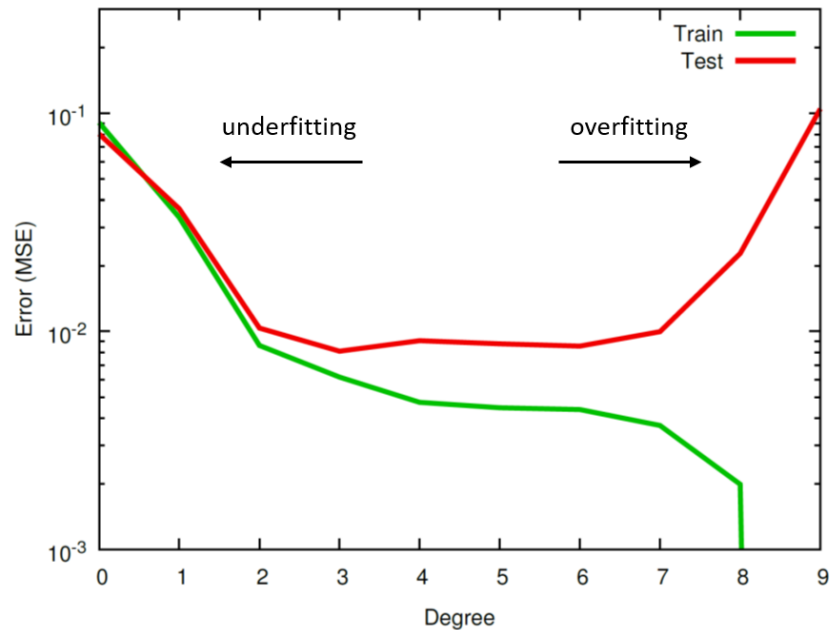
Minimizing loss in training data is often not the best



Low error on input data points, but high error nearby

Issue with Rich Representation

- Low error on input data points, but high error nearby
- Low error on training data, but high error on testing data



Linear Regression with RBF

```
xp = np.arange(-4.5, 4.5, 0.01).reshape(-1, 1)

d = 10
u = np.linspace(-4.5, 4.5, d)
sigma = 0.2


A = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d)])
rbfbasis = np.hstack([np.exp(-(xp-u[i])**2/(2*sigma**2)) for i in range(d)])

A = np.asmatrix(A)
rbfbasis = np.asmatrix(rbfbasis)

theta = (A.T*A).I*A.T*y
yp = rbfbasis*theta
```

$$\theta = (A^T A)^{-1} A^T y$$


- With many features, our prediction function becomes very expensive
- Can lead to overfitting



Regularization

Issue with Rich Representation

- Low error on input data points, but high error nearby
- Low error on training data, but high error on testing data




Generalization Error

- Fundamental problem: we are optimizing parameters to solve

$$\min_{\theta} \sum_{i=1}^m \ell(y_i, \hat{y}_i) = \min_{\theta} \sum_{i=1}^m \ell(y_i, \Phi\theta)$$

- But what we really care about is loss of prediction on new data (x, y)
 - also called generalization error
- Divide data into training set, and validation (testing) set



Representational Difficulties

- With many features, prediction function becomes very expressive (model complexity)
 - Choose less expressive function (e.g., lower degree polynomial, fewer RBF centers, larger RBF bandwidth)
 - Keep the magnitude of the parameter small
 - Regularization: penalize large parameters θ

$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_2^2$$

- λ : regularization parameter, trades off between low loss and small values of θ

With Less Basis Functions: Fewer RBF Centers

```
d = [2, 4, 6, 10]
sigma = 1

plt.figure(figsize=(12, 10))

for k in range(4):
    u = np.linspace(-4.5, 4.5, d[k])

    A = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d[k])])
    rbfbasis = np.hstack([np.exp(-(xp-u[i])**2/(2*sigma**2)) for i in range(d[k])])

    A = np.asmatrix(A)
    rbfbasis = np.asmatrix(rbfbasis)

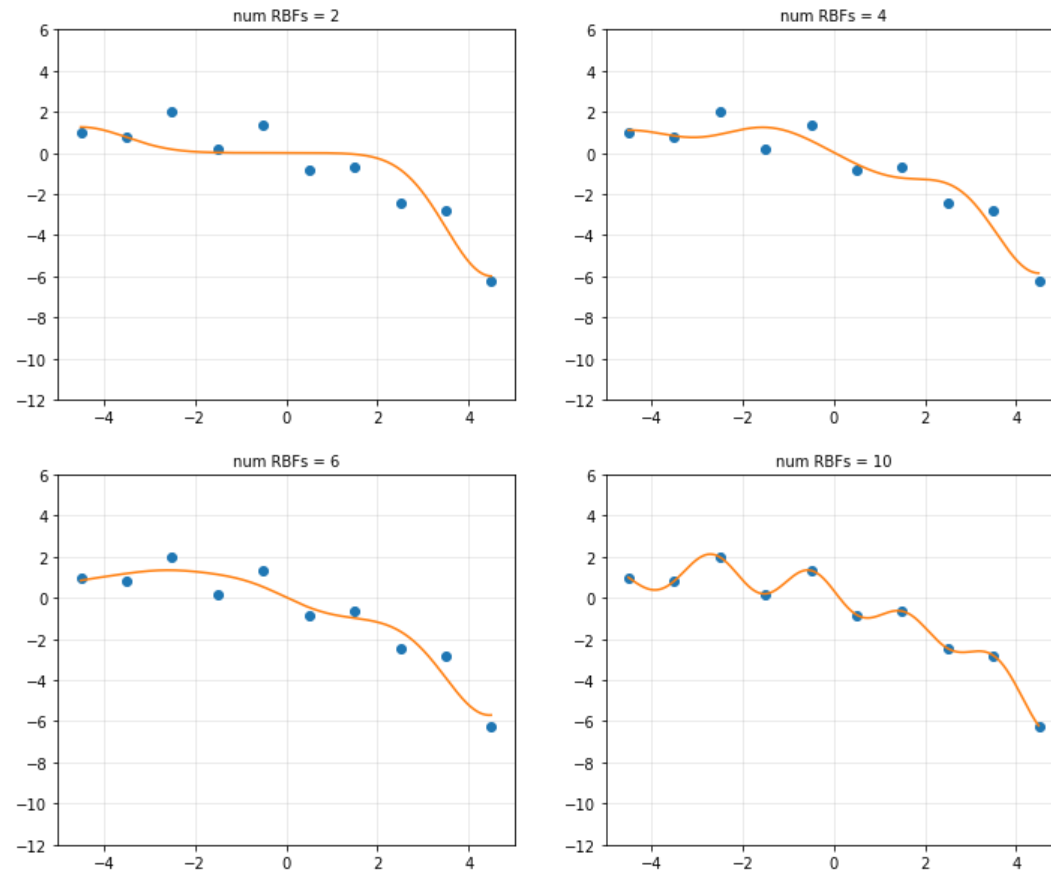
    theta = (A.T*A).I*A.T*y
    yp = rbfbasis*theta

    plt.subplot(2, 2, k+1)
    plt.plot(x, y, 'o')
    plt.plot(xp, yp)
    plt.axis([-5, 5, -12, 6])
    plt.title('num RBFs = {}'.format(d[k]), fontsize = 10)
    plt.grid(alpha = 0.3)
```

With Less Basis Functions: Fewer RBF Centers

- Least-squares fits for different numbers of RBFs

Nonlinear Regression with RBF Functions



Representational Difficulties

- With many features, prediction function becomes very expressive (model complexity)
 - Choose less expensive function (e.g., lower degree polynomial, fewer RBF centers, larger RBF bandwidth)
 - Keep the magnitude of the parameter small
 - Regularization: penalize large parameters θ

$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_2^2$$

- λ : regularization parameter, trades off between low loss and small values of θ

Regularization (Shrinkage Methods)

- Often, overfitting associated with very large estimated parameters
- We want to balance
 - how well function fits data
 - magnitude of coefficients

$$\text{Total cost} = \underbrace{\text{measure of fit}}_{RSS(\theta)} + \lambda \cdot \underbrace{\text{measure of magnitude of coefficients}}_{\lambda \cdot \|\theta\|_2^2}$$

$$\implies \min \|\Phi\theta - y\|_2^2 + \lambda \|\theta\|_2^2$$

- multi-objective optimization
- λ is a tuning parameter

Regularization (Shrinkage Methods)

- the second term, $\lambda \cdot \|\theta\|_2^2$, called a shrinkage penalty, is small when $\theta_1, \dots, \theta_d$ are close to zeros, and so it has the effect of shrinking the estimates of θ_j towards zero
- the tuning parameter λ serves to control the relative impact of these two terms on the regression coefficient estimates
- known as a *ridge regression*

RBF: Start from Rich Representation

```
d = 10
u = np.linspace(-4.5, 4.5, d)

sigma = 1


A = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d)])
rbfbasis = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d)])

A = np.asmatrix(A)
rbfbasis = np.asmatrix(rbfbasis)

theta = cvx.Variable([d, 1])
obj = cvx.Minimize(cvx.sum_squares(A*theta-y))
prob = cvx.Problem(obj).solve()

yp = rbfbasis*theta.value
```

$$\min \|\Phi\theta - y\|_2^2$$

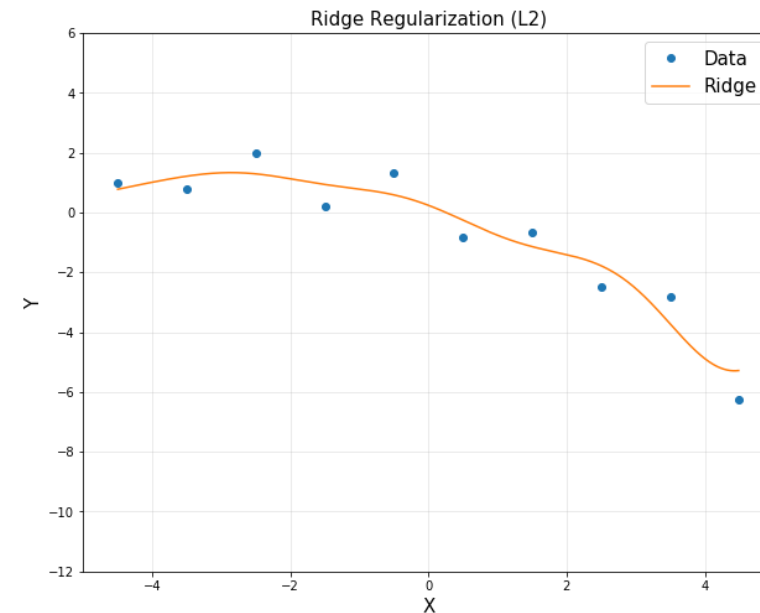




RBF with Regularization

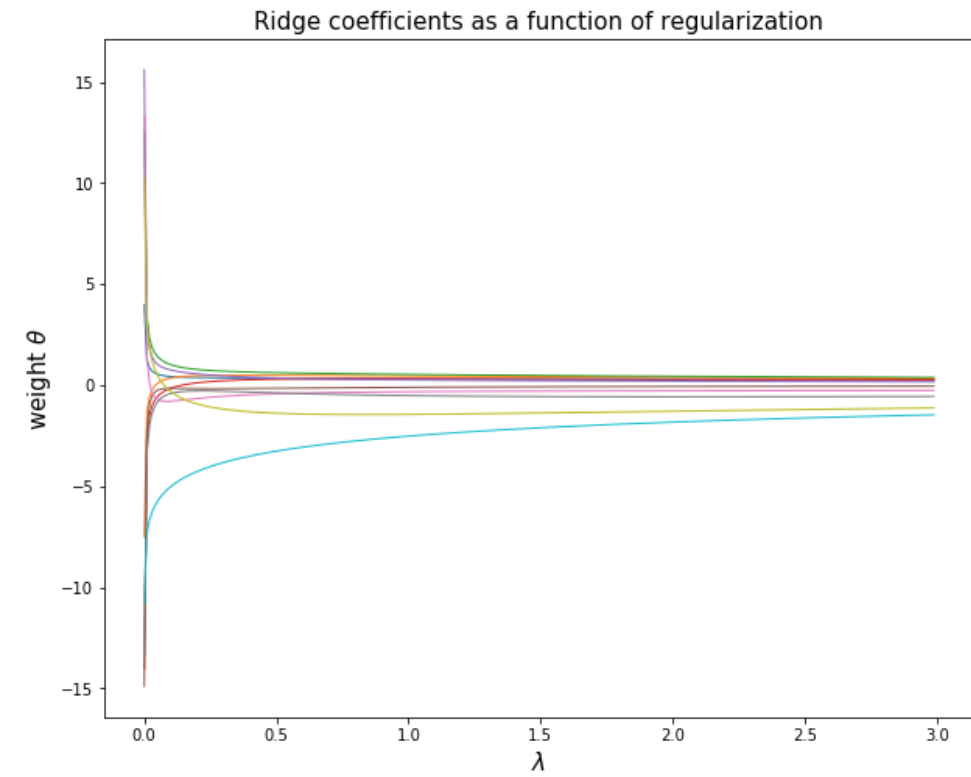

- Start from rich representation. Then, regularize coefficients θ

```
# ridge regression  
  
lamb = 0.1  
theta = cvx.Variable([d, 1])  
obj = cvx.Minimize(cvx.sum_squares(A*theta - y) + lamb*cvx.sum_squares(theta))  
prob = cvx.Problem(obj).solve()  
  
yp = rbfbasis*theta.value
```

$$\min \|\Phi\theta - y\|_2^2 + \lambda \|\theta\|_2^2$$



Coefficients θ



Let's Use L_1 Norm

- Ridge regression

$$\text{Total cost} = \underbrace{\text{measure of fit}}_{RSS(\theta)} + \underbrace{\lambda \cdot \text{measure of magnitude of coefficients}}_{\lambda \cdot \|\theta\|_2^2}$$

$$\implies \min \|\Phi\theta - y\|_2^2 + \boxed{\lambda \|\theta\|_2^2}$$

- Try this loss instead of ridge...

$$\text{Total cost} = \underbrace{\text{measure of fit}}_{RSS(\theta)} + \underbrace{\lambda \cdot \text{measure of magnitude of coefficients}}_{\lambda \cdot \|\theta\|_1}$$

$$\implies \min \|\Phi\theta - y\|_2^2 + \boxed{\lambda \|\theta\|_1}$$

- λ is a tuning parameter = balance of fit and sparsity
- Known as **LASSO**
 - least absolute shrinkage and selection operator


RBF with LASSO

```
# LASSO regression
```

```
lamb = 2  
theta = cvx.Variable([d, 1])  
obj = cvx.Minimize(cvx.sum_squares(A*theta - y) + lamb*cvx.norm(theta, 1))  
prob = cvx.Problem(obj).solve()  
  
yp = rbfbasis*theta.value
```


$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_1$$

- Approximated function looks similar to that of ridge regression

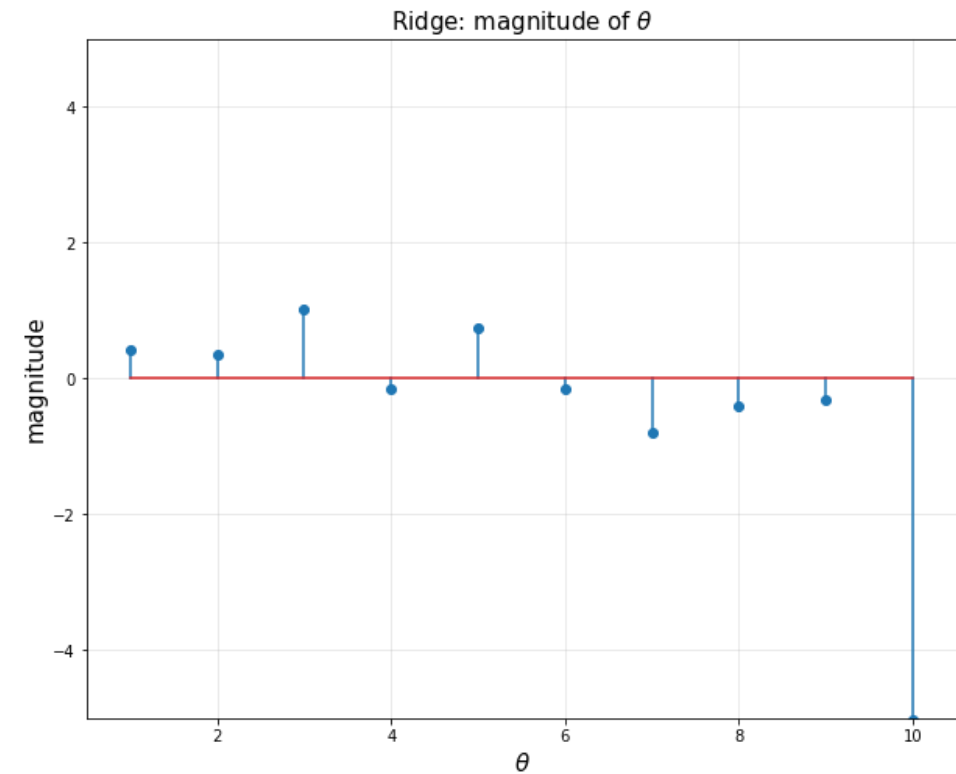


Coefficients θ with LASSO

- Non-zero coefficients indicate 'selected' features



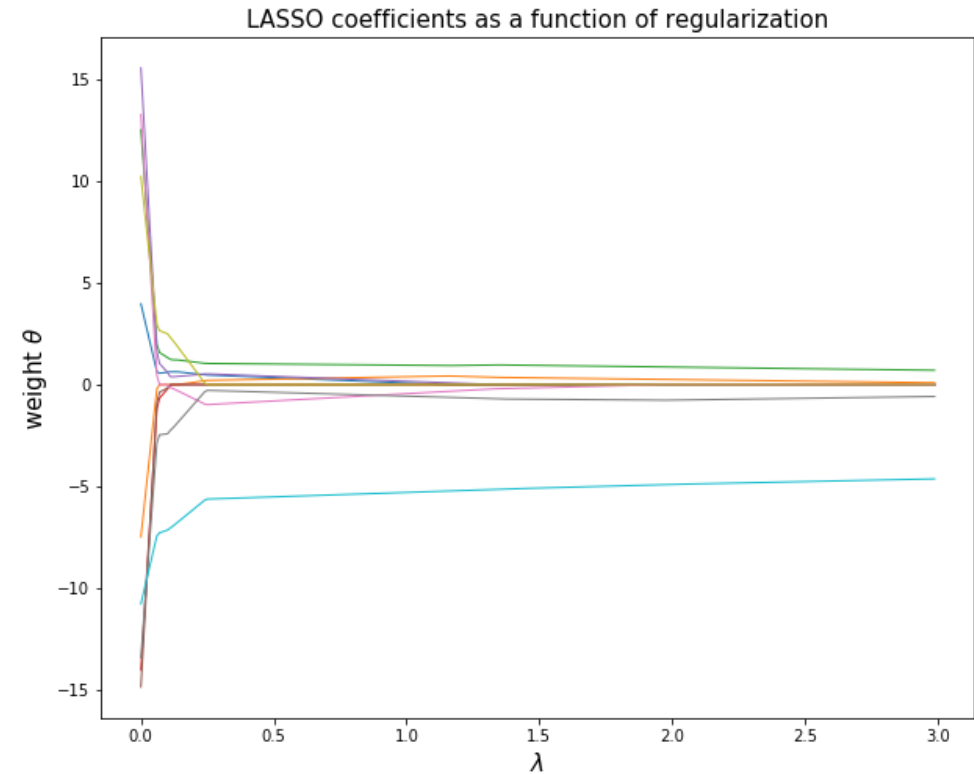

LASSO



Ridge

Coefficients θ with LASSO

- Non-zero coefficients indicate 'selected' features



Sparsity for Feature Selection using Lasso

- Least squares with a penalty on the L_1 norm of the parameters
- Start with full model (all possible features)
- ‘Shrink’ some coefficients exactly to 0
 - *i.e.*, knock out certain features
 - The L_1 penalty has the effect of forcing some of the coefficient estimates to be exactly equal to zero
- Non-zero coefficients indicate ‘selected’ features

Regression with Selected Features

```
# reduced order model
# we will use only theta 2, 3, 8, 10

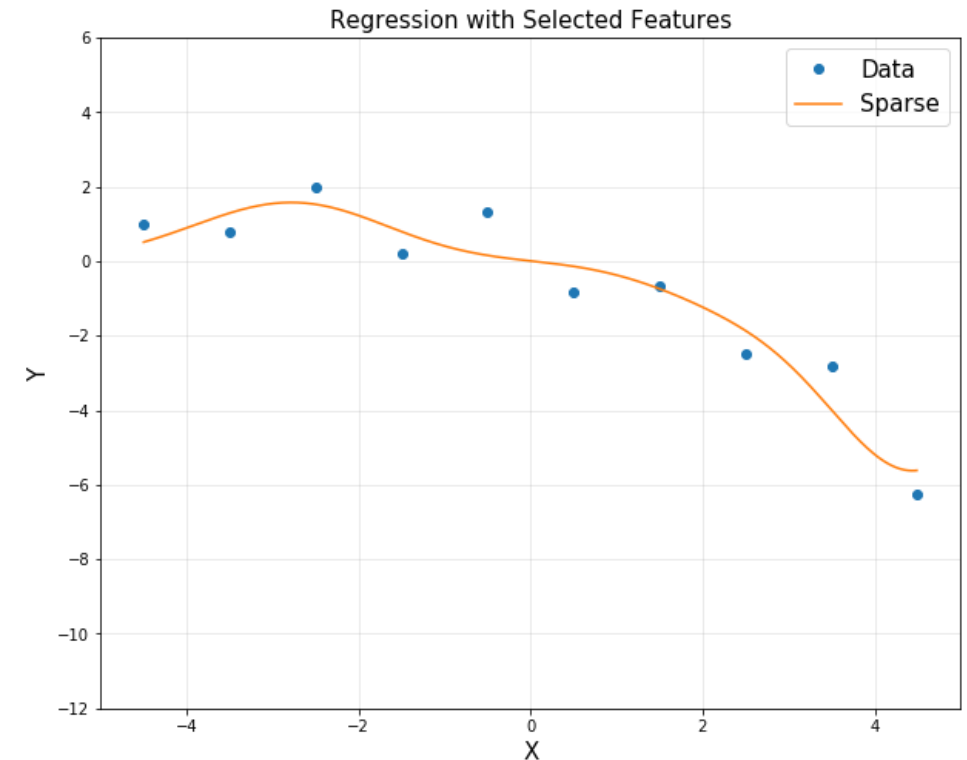
d = 4
u = np.array([-3.5, -2.5, 2.5, 4.5])
sigma = 1

rbfbasis = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d)])
A = np.hstack([np.exp(-(x-u[i])**2/(2*sigma**2)) for i in range(d)])

rbfbasis = np.asmatrix(rbfbasis)
A = np.asmatrix(A)

theta = cvx.Variable([d, 1])
obj = cvx.Minimize(cvx.norm(A*theta-y, 2))
prob = cvx.Problem(obj).solve()

yp = rbfbasis*theta.value
```



LASSO vs. Ridge

- Another equivalent forms of optimizations

$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_1$$

$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_2^2$$

\implies

$$\begin{aligned} & \min_{\theta} \|\Phi\theta - y\|_2^2 \\ & \text{subject to } \|\theta\|_1 \leq s_1 \end{aligned}$$

$$\begin{aligned} & \min_{\theta} \|\Phi\theta - y\|_2^2 \\ & \text{subject to } \|\theta\|_2 \leq s_2 \end{aligned}$$

LASSO vs. Ridge

- Another equivalent forms of optimizations



$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_1$$

$$\min \|\Phi\theta - y\|_2^2 + \lambda\|\theta\|_2^2$$

\Rightarrow

$$\begin{aligned} \min_{\theta} \quad & \|\Phi\theta - y\|_2^2 \\ \text{subject to} \quad & \|\theta\|_1 \leq s_1 \end{aligned}$$


$$\begin{aligned} \min_{\theta} \quad & \|\Phi\theta - y\|_2^2 \\ \text{subject to} \quad & \|\theta\|_2 \leq s_2 \end{aligned}$$



L2 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss alone.


Since the derivative of $\|x\|_2^2$ is zero at zero, the optimal will never move there if it was not already there.



L2 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss alone.


Since the derivative of $\|x\|_2^2$ is zero at zero, the optimal will never move there if it was not already there.



L2 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss alone.


Since the derivative of $\|x\|_2^2$ is zero at zero, the optimal will never move there if it was not already there.



L2 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss alone.


Since the derivative of $\|x\|_2^2$ is zero at zero, the optimal will never move there if it was not already there.



L2 Regularizers: Simple Example


Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss alone.

Since the derivative of $\|x\|_2^2$ is zero at zero, the optimal will never move there if it was not already there.




L1 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss without penalty.




L1 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss without penalty.




L1 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss without penalty.




L1 Regularizers: Simple Example

Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss without penalty.



L1 Regularizers: Simple Example


Increasing the λ parameter moves the optimal closer to 0, and away from the optimal for the loss without penalty.



Evaluation

- Adding more features will always decrease the loss
- How do we determine when an algorithm achieves “good” performance?

- A better criterion:
 - Training set (e.g., 70 %)
 - Testing set (e.g., 30 %)



- Performance on testing set called *generalization* performance