



Gradient Descent

Industrial AI Lab.

Prof. Seungchul Lee

Yunseob Hwang, Juwon Na

How do we Find $\nabla_x f(x) = 0$

Analytic Approach

- Direct solution
 - In some cases, it is possible to analytically compute x^* such that $\nabla_x f(x^*) = 0$

$$f(x) = 2x_1^2 + x_2^2 + x_1 x_2 - 6x_1 - 5x_2$$

$$\implies \nabla_x f(x) = \begin{bmatrix} 4x_1 + x_2 - 6 \\ 2x_2 + x_1 - 5 \end{bmatrix}$$

$$\implies x^* = \begin{bmatrix} 4 & 1 \\ 1 & 2 \end{bmatrix}^{-1} \begin{bmatrix} 6 \\ 5 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

Gradients

- Matrix derivatives

| y | $\frac{\partial y}{\partial x}$ |
|----------|---------------------------------|
| Ax | A^T |
| $x^T A$ | A |
| $x^T x$ | $2x$ |
| $x^T Ax$ | $Ax + A^T x$ |

How to Find $\nabla_x f(x) = 0$

- Direct solution
 - In some cases, it is possible to analytically compute x^* such that $\nabla_x f(x^*) = 0$

| y | $\frac{\partial y}{\partial x}$ |
|----------|---------------------------------|
| Ax | A^T |
| $x^T A$ | A |
| $x^T x$ | $2x$ |
| $x^T Ax$ | $Ax + A^T x$ |

$$f(x) = 2x_1^2 + x_2^2 + x_1 x_2 - 6x_1 - 5x_2$$

$$\Rightarrow \nabla_x f(x) = \begin{bmatrix} 4x_1 + x_2 - 6 \\ 2x_2 + x_1 - 5 \end{bmatrix}$$

$$\Rightarrow x^\star = \begin{bmatrix} 4 & 1 \\ 1 & 2 \end{bmatrix}^{-1} \begin{bmatrix} 6 \\ 5 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

Examples

- affine function $g(x) = a^T x + b$

$$\nabla g(x) = a, \quad \nabla^2 g(x) = 0$$

| y | $\frac{\partial y}{\partial x}$ |
|----------|---------------------------------|
| Ax | A^T |
| $x^T A$ | A |
| $x^T x$ | $2x$ |
| $x^T Ax$ | $Ax + A^T x$ |

- quadratic function $g(x) = x^T P x + q^T x + r$, $P = P^T$

$$\nabla g(x) = 2Px + q, \quad \nabla^2 g(x) = 2P$$


- $g(x) = \|Ax - b\|^2 = x^T A^T A x - 2b^T A x + b^T b$

$$\nabla g(x) = 2A^T A x - 2A^T b, \quad \nabla^2 g(x) = 2A^T A$$

Iterative Approach

- Iterative methods

- More commonly the condition that the gradient equal zero will not have an analytical solution, require iterative methods





- The gradient points in the direction of “steepest ascent” for function f

Gradient Descent

Descent Direction (1D)


- It motivates the *gradient descent* algorithm, which repeatedly takes steps in the direction of the negative gradient

$$x \leftarrow x - \alpha \nabla_x f(x) \quad \text{for some step size } \alpha > 0$$



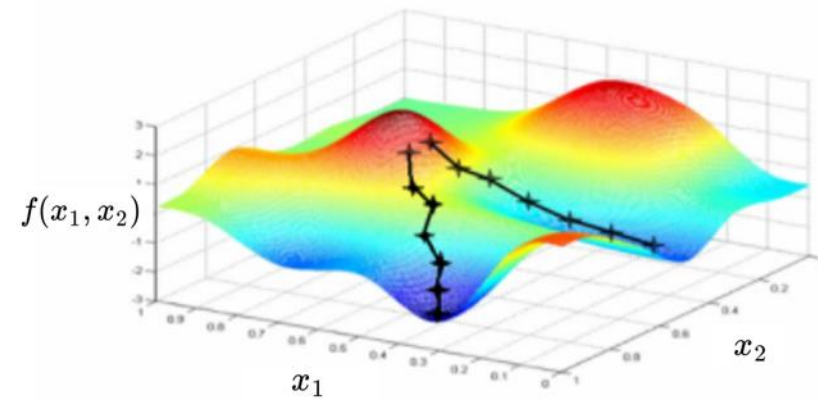
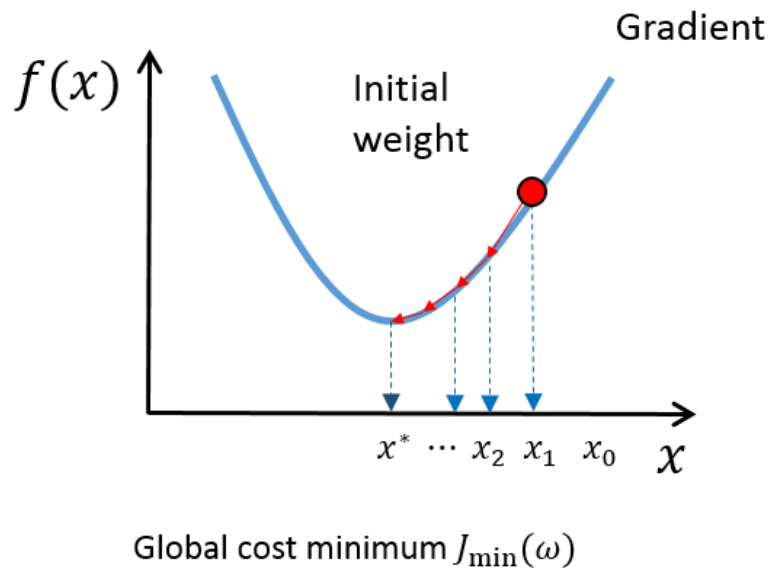
Gradient Descent

Repeat: $x \leftarrow x - \alpha \nabla_x f(x)$ for some *step size* $\alpha > 0$




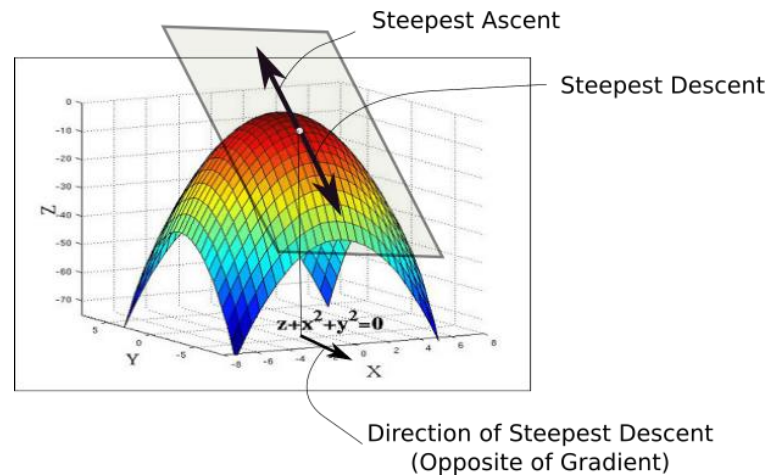
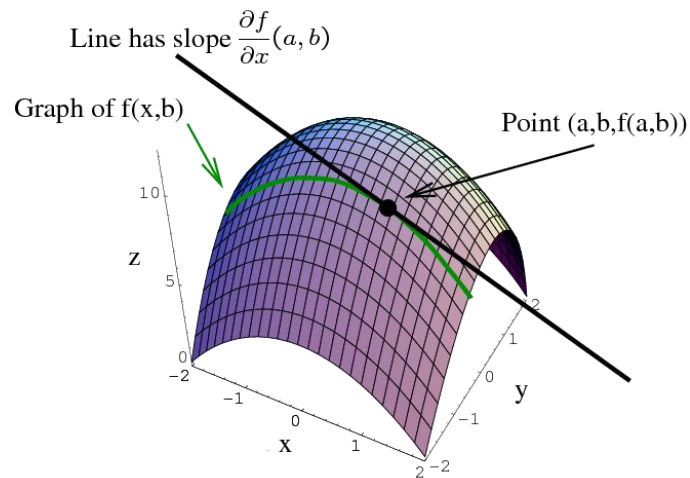
Gradient Descent in High Dimension

Repeat: $x \leftarrow x - \alpha \nabla_x f(x)$ for some *step size* $\alpha > 0$



Gradient Descent in High Dimension

Repeat: $x \leftarrow x - \alpha \nabla_x f(x)$ for some *step size* $\alpha > 0$



Gradient Descent

$$\begin{aligned} \min \quad & (x_1 - 3)^2 + (x_2 - 3)^2 \\ = \min \quad & \frac{1}{2} \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} - \begin{bmatrix} 6 & 6 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} + 18 \end{aligned}$$

$$\begin{aligned} f &= \frac{1}{2} X^T H X + g^T X \\ \nabla f &= H X + g \end{aligned}$$

- Update rule: $X_{i+1} = X_i - \alpha_i \nabla f(X_i)$

```
H = np.matrix([[2, 0],[0, 2]])
g = -np.matrix([[6],[6]])

x = np.zeros((2,1))
alpha = 0.2

for i in range(25):
    df = H*x + g
    x = x - alpha*df


print(x)
```

```
[[ 2.99999147]
 [ 2.99999147]]
```


| y | $\frac{\partial y}{\partial x}$ |
|----------|---------------------------------|
| Ax | A^T |
| $x^T A$ | A |
| $x^T x$ | $2x$ |
| $x^T Ax$ | $Ax + A^T x$ |

Choosing Step Size α


- Learning rate



Too small: converge
very slowly




Too big: overshoot and
even diverge




Reduce size over time

Where will We Converge?



Any local minimum is a global minimum



Multiple local minima may exist

- Random initialization
- Multiple trials

Gradient Descent vs. Analytical Solution


- Analytical solution for MSE
- Gradient descent
 - Easy to implement
 - Very general, can be applied to any differentiable loss functions
 - Requires less memory and computations (for stochastic methods)
- Gradient descent provides a general learning framework
- Can be used both for classification and regression
- Training Neural Networks: Gradient Descent

Practically Solving Optimization Problems

- The good news: for many classes of optimization problems, people have already done all the “hard work” of developing numerical algorithms
 - A wide range of tools that can take optimization problems in “natural” forms and compute a solution
- CVX (or CVXPY) as an optimization solver
 - Only for convex problems
 - Download: <https://www.cvxpy.org/>
- Gradient descent
 - Neural networks/deep learning
 - TensorFlow

Summary: Training Neural Networks

- Optimization procedure



- It is not easy to numerically compute gradients in network in general.
 - The good news: people have already done all the "hard work" of developing numerical solvers (or libraries)
 - There are a wide range of tools
 - We will use **TensorFlow**