




# Modern CNNs

**Industrial AI Lab.**


**Prof. Seungchul Lee**

# ImageNet

- Human performance = 5.1 %

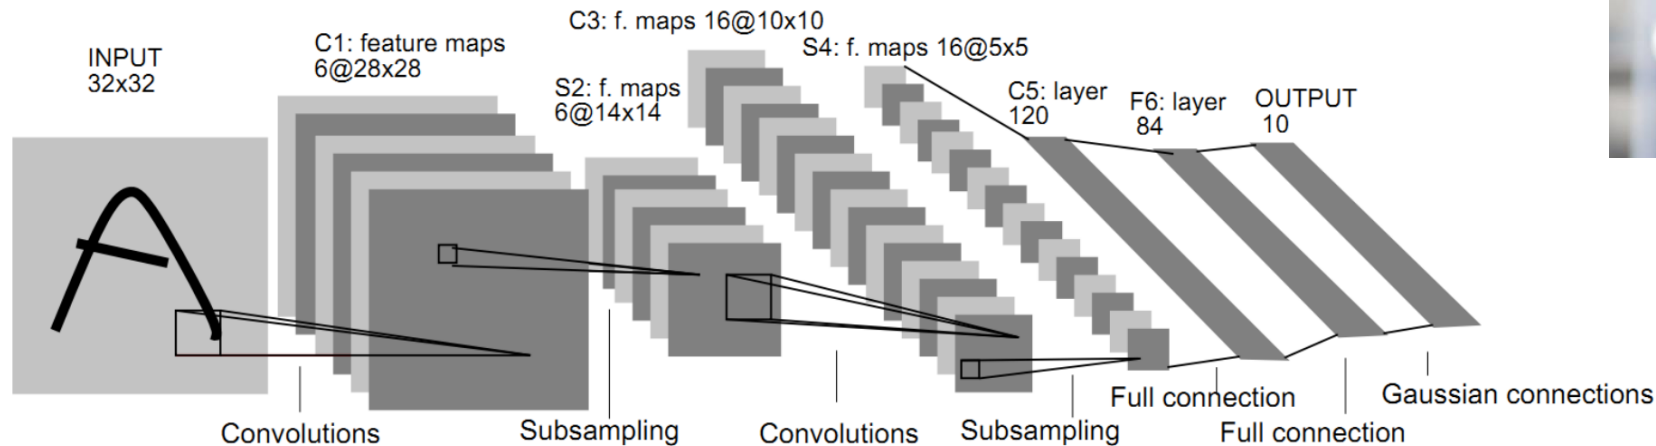


# ImageNet



# LeNet


- CNN = Convolutional Neural Networks = ConvNet
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition.
- All are still the basic components of modern ConvNets!



Yann LeCun



# AlexNet

- Simplified version of Krizhevsky, Alex, Sutskever, and Hinton. "Imagenet classification with deep convolutional neural networks." NIPS 2012
- LeNet-style backbone, plus:
  - ReLU [Nair & Hinton 2010]
    - “RevoLUtion of deep learning”\*
    - Accelerate training
  - Dropout [Hinton et al 2012]
    - In-network ensembling
    - Reduce overfitting
  - Data augmentation
    - Label-preserving transformation
    - Reduce overfitting




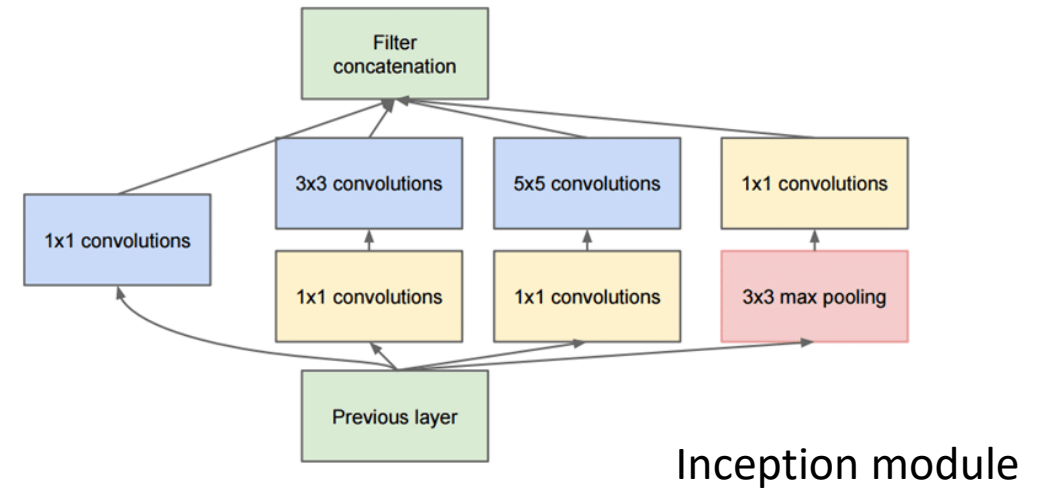
# VGG-16/19

- Simonyan, Karen, and Zisserman. "Very deep convolutional networks for large-scale image recognition." (2014)
- Simply "Very Deep"!
  - Modularized design
    - 3x3 Conv as the module
    - Stack the same module
    - Same computation for each module
  - Stage-wise training
    - VGG-11 → VGG-13 → VGG-16
    - We need a better initialization...



# GoogleNet/Inception


- Multiple branches
  - e.g., 1x1, 3x3, 5x5, pool
- Shortcuts
  - stand-alone 1x1, merged by concatenation
- Bottleneck
  - Reduce dim by 1x1 before expensive 3x3/5x5 conv



# ResNet (Deep Residual Learning)

- He, Kaiming, et al. “Deep residual learning for image recognition.” CVPR. 2016.
- Plane net


$H(x)$  is any desired mapping,  
hope the small subnet fit  $H(x)$





# ResNet (Deep Residual Learning)

- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Residual net
- Skip connection



$H(x)$  is any desired mapping,  
~~hope the small subnet fit  $H(x)$~~   
hope the small subnet fit  $F(x)$   
Let  $H(x) = F(x) + x$

- A direct connection between 2 non-consecutive layers
- No gradient vanishing




# Skip Connection

- A skip connection is a connection that bypasses at least one layer.
- Here, it is often used to transfer local information by concatenating or summing feature maps from the downsampling path with feature maps from the upsampling path.
  - Will see it at FCN later
  - Merging features from various resolution levels helps combining context information with spatial information.


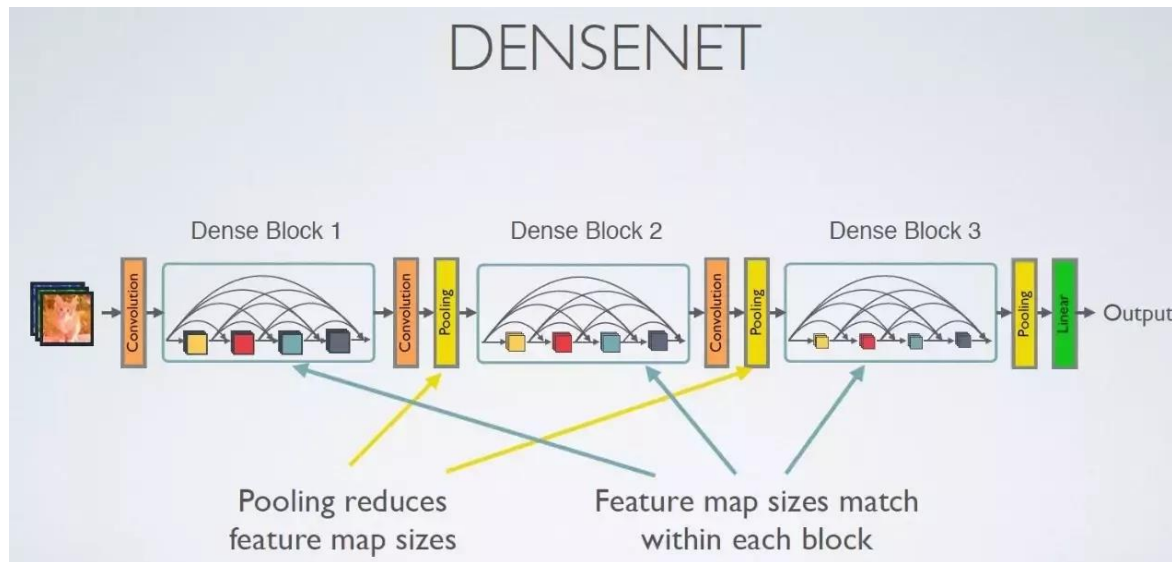
# Residual Net

```
def residual_net(x):  
    conv1 = tf.layers.conv2d(inputs = x,  
                              filters = 32,  
                              kernel_size = [3, 3],  
                              padding = "SAME",  
                              activation = tf.nn.relu)  
  
    conv2 = tf.layers.conv2d(inputs = conv1,  
                              filters = 32,  
                              kernel_size = [3, 3],  
                              padding = "SAME",  
                              activation = tf.nn.relu)  
  
→ maxp2 = tf.layers.max_pooling2d(inputs = x + conv2,  
                                   pool_size = [2, 2],  
                                   strides = 2)  
  
    flat = tf.layers.flatten(maxp2)  
    hidden = tf.layers.dense(inputs = flat,  
                              units = n_hidden,  
                              activation = tf.nn.relu)  
    output = tf.layers.dense(inputs = hidden,  
                              units = n_output)  
  
    return output
```



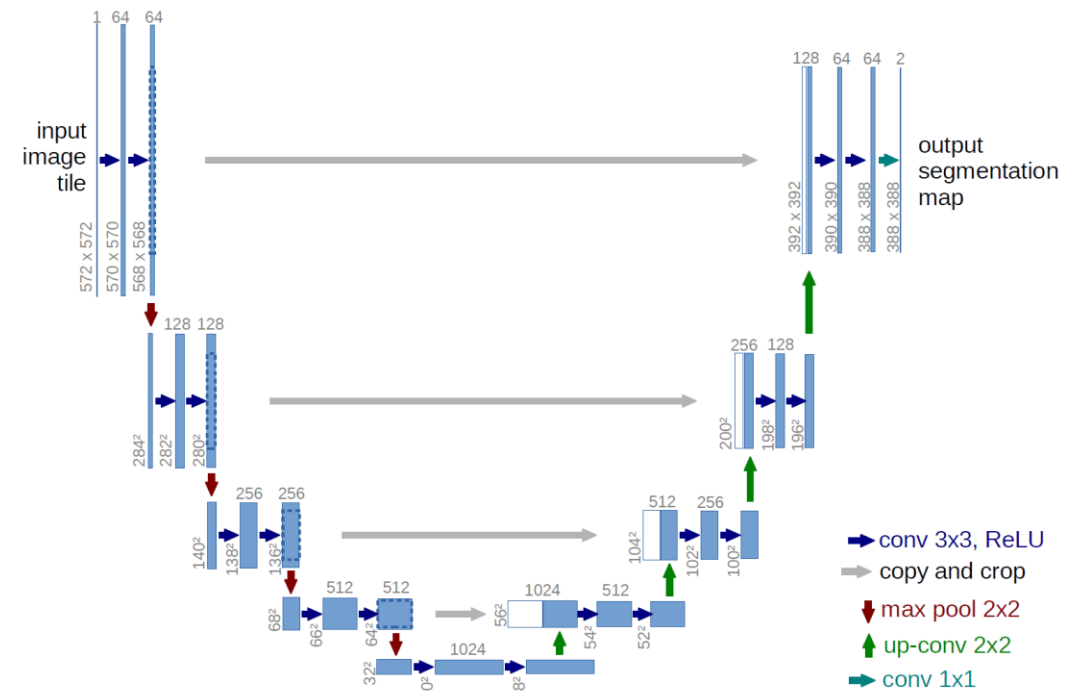
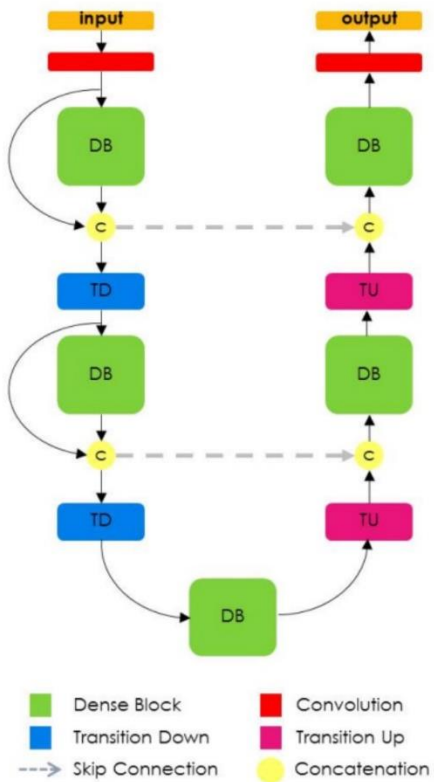
# DensNets

- Densely Connected Convolutional Networks




# U-Net

- The U-Net owes its name to its symmetric shape
  - better segmentation in medical imaging



# Pre-trained Models

- Training a model on ImageNet from scratch takes days or weeks.
- Many models trained on ImageNet and their weights are publicly available!
- Transfer learning
  - Use pre-trained weights, remove last layers to compute representations of images
  - **The network is used as a generic feature extractor**
  - Train a classification model from these features on a new classification task
  - Pre-trained models can extract more general image features that can help identify edges, textures, shapes, and object composition
  - Better than handcrafted feature extraction on natural images



# Image Classification with VGG16

- Target data

- 5 classes

```
Dict = ['Hat', 'Cube', 'Card', 'Torch', 'screw']
```

- Target data to VGG16

- Poor performance

1. mosquito\_net: 4.66%
2. toilet\_tissue: 4.00%
3. envelope: 2.29%
4. carton: 2.20%
5. photocopier: 1.86%

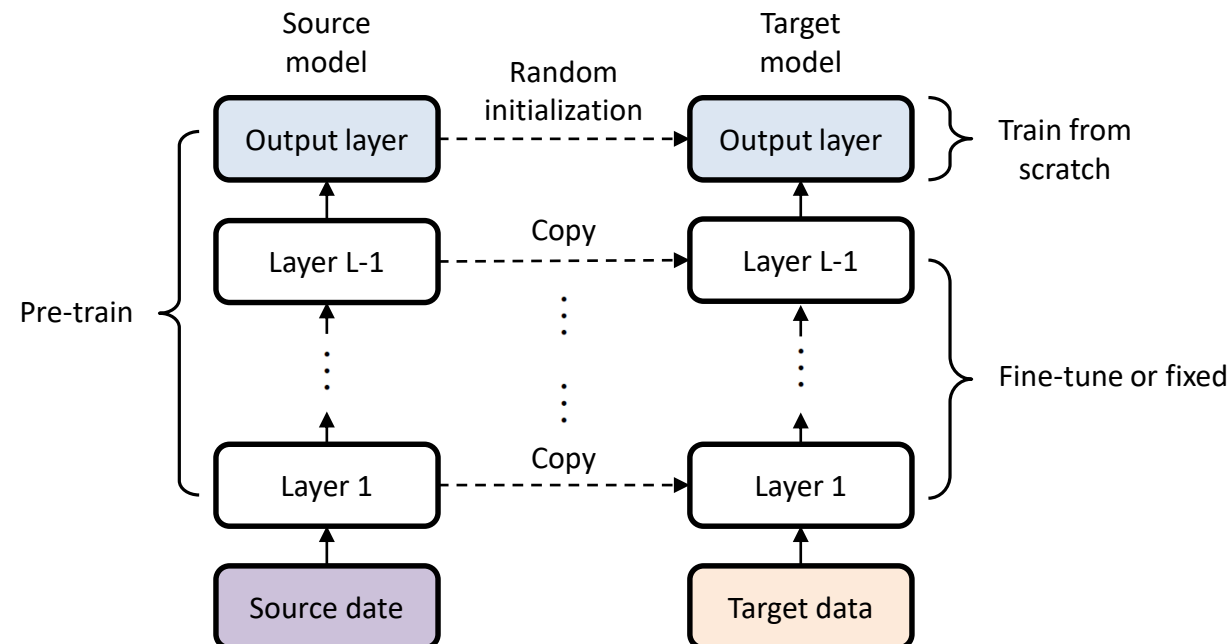
Label : Cube






# Transfer Learning

- We assume that these model parameters contain the knowledge learned from the source data set and that this knowledge will be equally applicable to the target data set.
- We will train the output layer from scratch, while the parameters of all remaining layers are fine tuned based on the parameters of the source model.
- Or initialize all weights from pre-trained model, then train them with target data



# Transfer Learning Structure



# Transfer Learning Implementation

```
vgg16_weights = model.get_weights()

weights = {
    'conv1_1' : tf.constant(vgg16_weights[0]),
    'conv1_2' : tf.constant(vgg16_weights[2]),

    'conv2_1' : tf.constant(vgg16_weights[4]),
    'conv2_2' : tf.constant(vgg16_weights[6]),

    'conv3_1' : tf.constant(vgg16_weights[8]),
    'conv3_2' : tf.constant(vgg16_weights[10]),
    'conv3_3' : tf.constant(vgg16_weights[12]),

    'conv4_1' : tf.constant(vgg16_weights[14]),
    'conv4_2' : tf.constant(vgg16_weights[16]),
    'conv4_3' : tf.constant(vgg16_weights[18]),

    'conv5_1' : tf.constant(vgg16_weights[20]),
    'conv5_2' : tf.constant(vgg16_weights[22]),
    'conv5_3' : tf.constant(vgg16_weights[24]),

    'fc1' : tf.constant(vgg16_weights[26]),
    'fc2' : tf.constant(vgg16_weights[28]),
    # train from scratch
    'out' : tf.Variable(tf.random_normal([4096, 5], stddev = 0.1))
}
```

```
biases = {
    'conv1_1' : tf.constant(vgg16_weights[1]),
    'conv1_2' : tf.constant(vgg16_weights[3]),

    'conv2_1' : tf.constant(vgg16_weights[5]),
    'conv2_2' : tf.constant(vgg16_weights[7]),

    'conv3_1' : tf.constant(vgg16_weights[9]),
    'conv3_2' : tf.constant(vgg16_weights[11]),
    'conv3_3' : tf.constant(vgg16_weights[13]),

    'conv4_1' : tf.constant(vgg16_weights[15]),
    'conv4_2' : tf.constant(vgg16_weights[17]),
    'conv4_3' : tf.constant(vgg16_weights[19]),

    'conv5_1' : tf.constant(vgg16_weights[21]),
    'conv5_2' : tf.constant(vgg16_weights[23]),
    'conv5_3' : tf.constant(vgg16_weights[25]),

    'fc1' : tf.constant(vgg16_weights[27]),
    'fc2' : tf.constant(vgg16_weights[29]),
    # train from scratch
    'out' : tf.Variable(tf.random_normal([5], stddev = 0.1))
}
```

# Testing

```
Dict = ['Hat', 'Cube', 'Card', 'Torch', 'screw']
```

Prediction : Cube

Probability : [0. 0.99 0.01 0. 0. ]

