



# Modern CNNs

**Industrial AI Lab.**  
**Prof. Seungchul Lee**

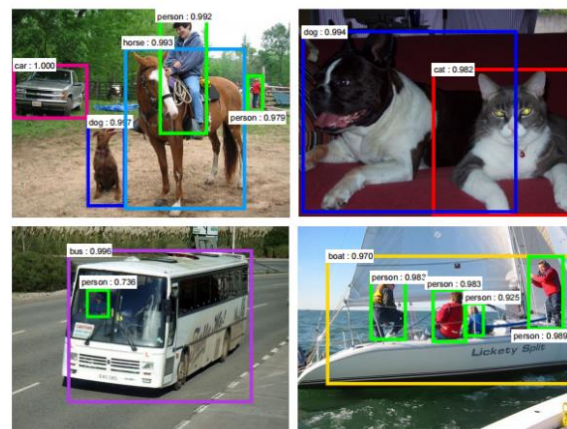
# CNNs for Computer Vision



[Krizhevsky 2012]



[Ciresan et al. 2013]



[Faster R-CNN - Ren 2015]




[NVIDIA dev blog]


# Modern CNN

# ImageNet

- Human performance = 5.1 %




# ImageNet



# LeNet


- CNN = Convolutional Neural Networks = ConvNet
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition.
- All are still the basic components of modern ConvNets!



Yann LeCun

# AlexNet



- Simplified version of Krizhevsky, Alex, Sutskever, and Hinton. "Imagenet classification with deep convolutional neural networks." NIPS 2012
- LeNet-style backbone, plus:
  - ReLU [Nair & Hinton 2010]
    - “RevoLUtion of deep learning”\*
    - Accelerate training; better grad prop (vs. tanh)
  - Dropout [Hinton et al 2012]
    - In-network ensembling
    - Reduce overfitting
  - Data augmentation
    - Label-preserving transformation
    - Reduce overfitting



# VGG-16/19

- Simonyan, Karen, and Zisserman. "Very deep convolutional networks for large-scale image recognition." (2014)

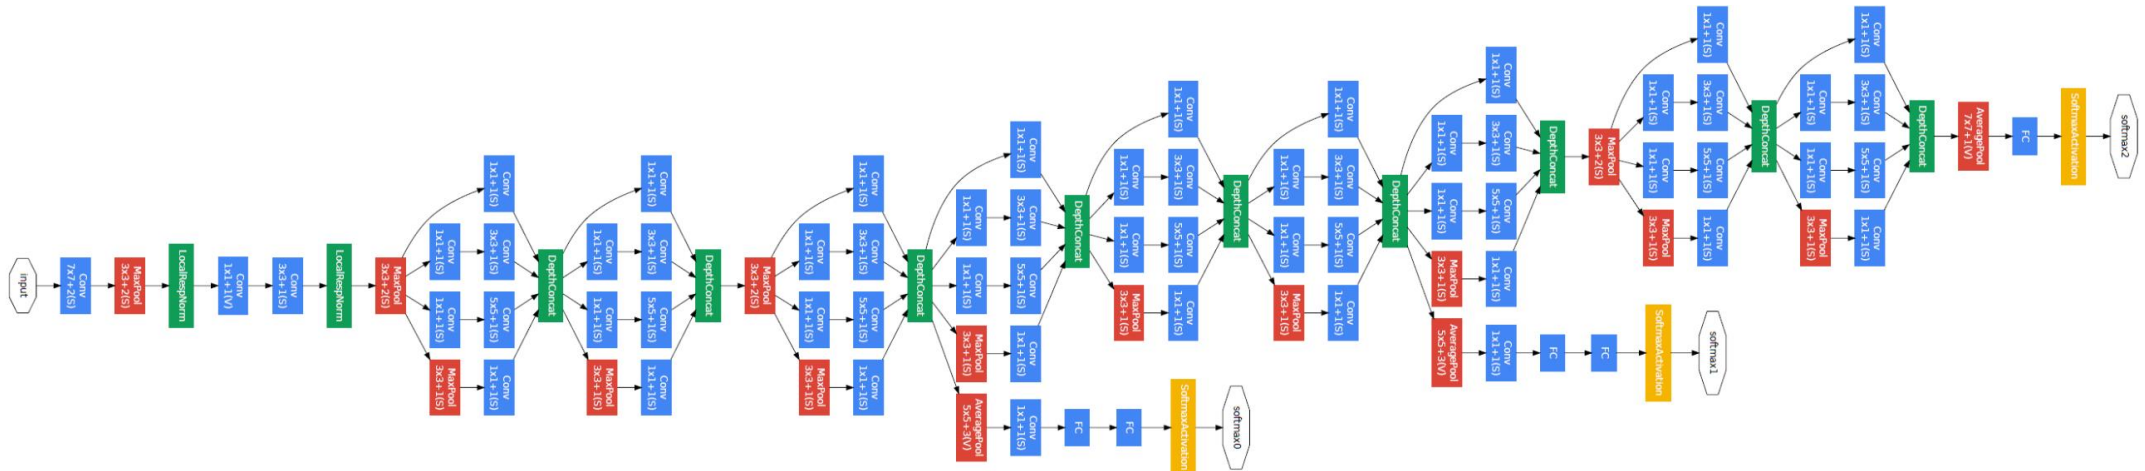
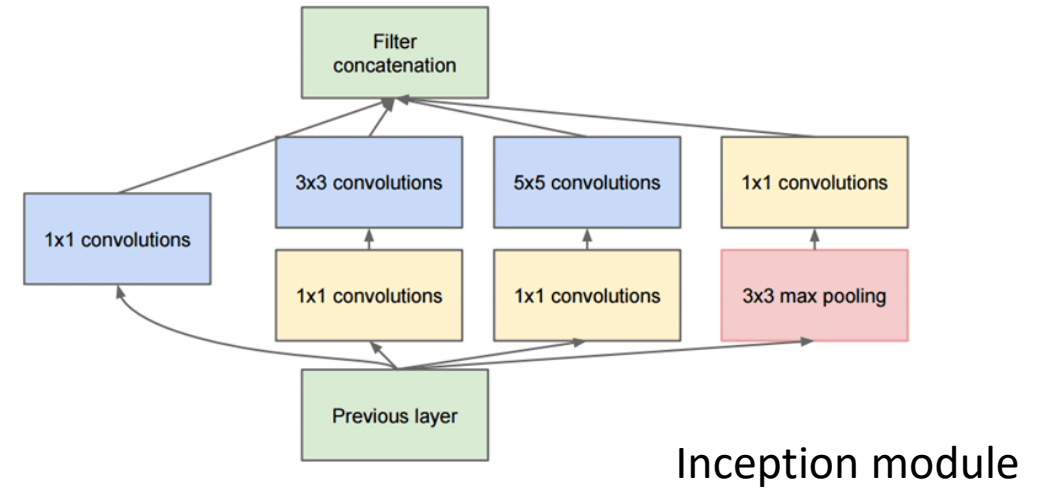
- Simply “Very Deep”!
  - Modularized design
    - 3x3 Conv as the module
    - Stack the same module
    - Same computation for each module
  - Stage-wise training
    - VGG-11 → VGG-13 → VGG-16
    - We need a better initialization...





# GoogleNet/Inception


- Multiple branches
  - e.g., 1x1, 3x3, 5x5, pool
- Shortcuts
  - stand-alone 1x1, merged by concat.
- Bottleneck
  - Reduce dim by 1x1 before expensive 3x3/5x5 conv



# ResNet (Deep Residual Learning)


- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Plane net

$H(x)$  is any desired mapping,  
hope the small subnet fit  $H(x)$



# ResNet (Deep Residual Learning)

- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Residual net
- Skip connection




$H(x)$  is any desired mapping,  
~~hope the small subnet fit  $H(x)$~~   
hope the small subnet fit  $F(x)$   
Let  $H(x) = F(x) + x$

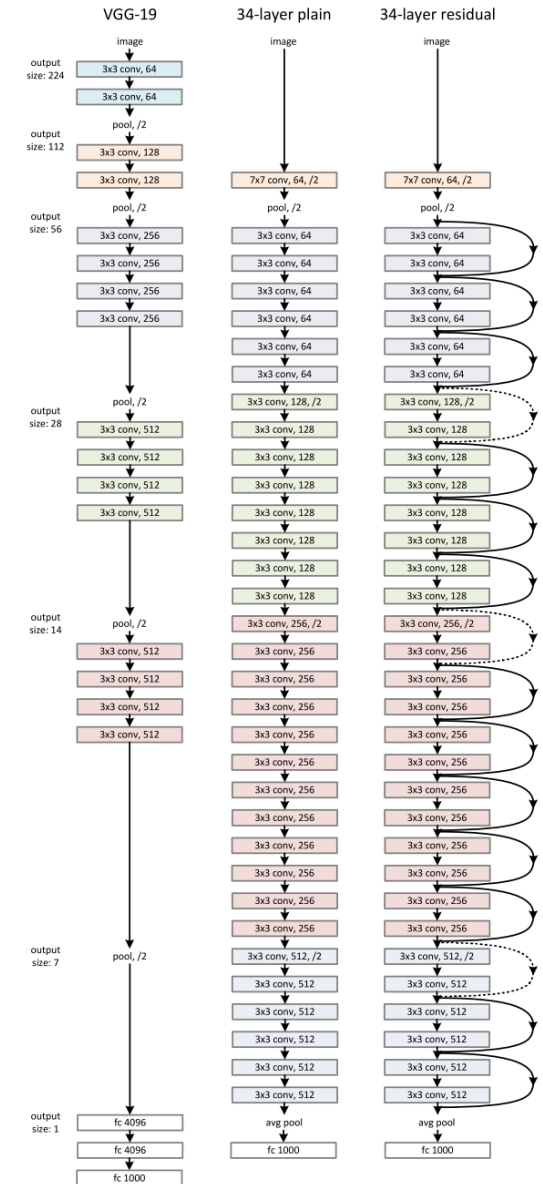
- A direct connection between 2 non-consecutive layers
- No vanishing gradient

# ResNet (Deep Residual Learning)

- Parameters are optimized to learn a residual, that is the difference between the value before the block and the one needed after.
- $F(x)$  is a residual mapping w.r.t. identity


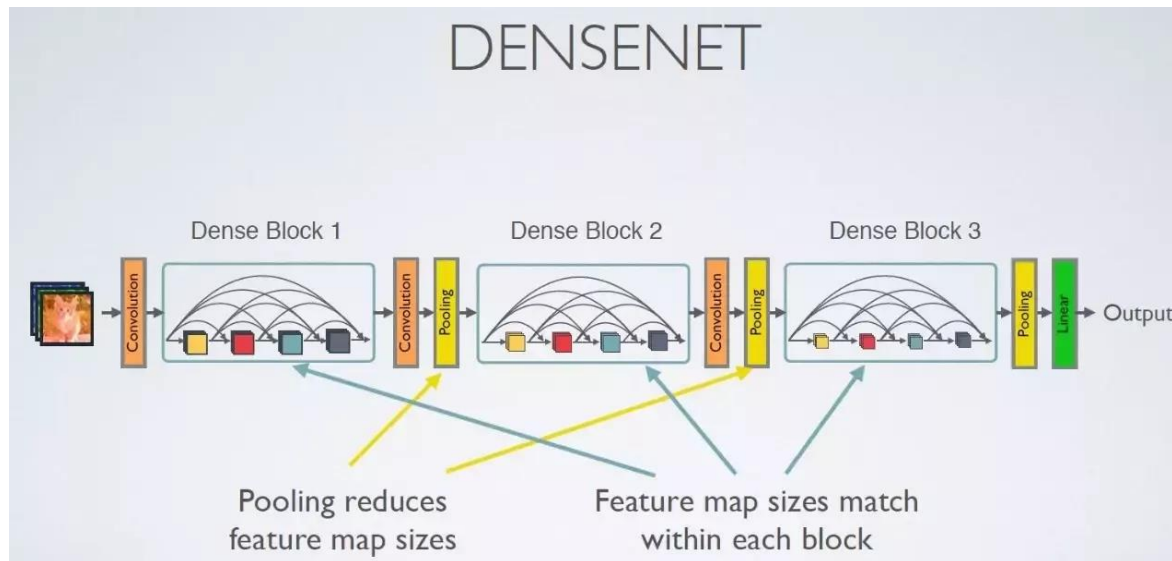


- If identity were optimal, easy to set weights as 0
- If optimal mapping is closer to identity, easier to find small fluctuations

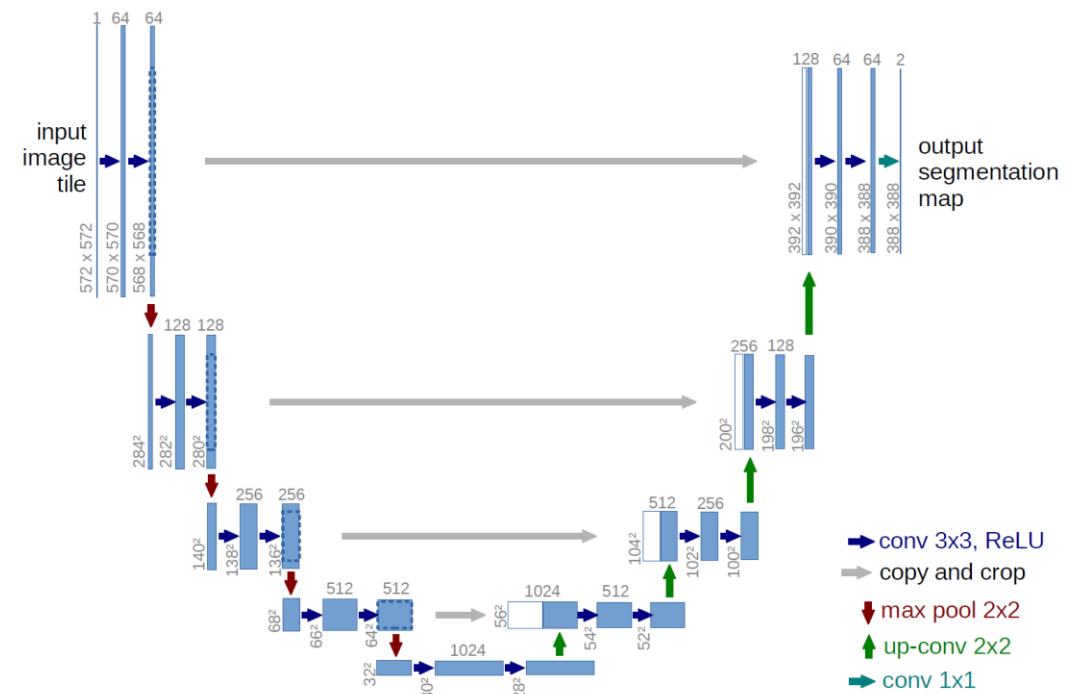
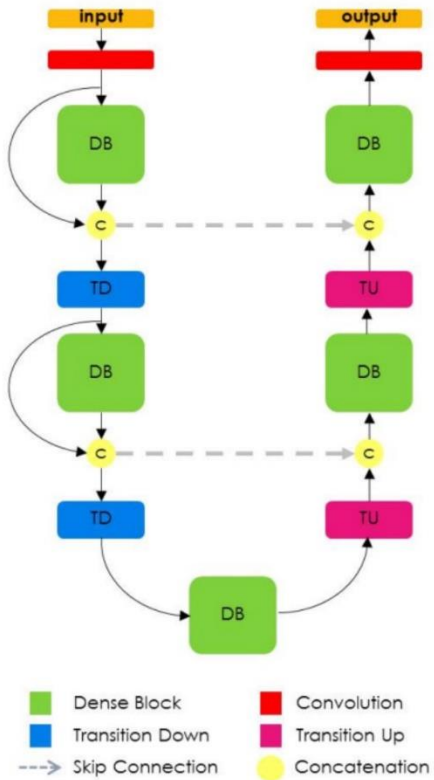


# DensNets

- Densely Connected Convolutional Networks



# U-Net



# U-Net

- The U-Net owes its name to its symmetric shape
- The U-Net architecture is built upon the Fully Convolutional Network and modified in a way that it yields better segmentation in medical imaging.
- Compared to FCN-8, the two main differences are
  - U-net is symmetric and
  - the skip connections between the downsampling path and the upsampling path apply a concatenation operator instead of a sum.
- These skip connections intend to provide local information to the global information while upsampling. Because of its symmetry, the network has a large number of feature maps in the upsampling path, which allows to transfer information.

# Pre-trained Models

- Training a model on ImageNet from scratch takes days or weeks.
- Many models trained on ImageNet and their weights are publicly available!
- Transfer learning
  - Use pre-trained weights, remove last layers to compute representations of images
  - Train a classification model from these features on a new classification task
  - The network is used as a generic feature extractor
  - Better than handcrafted feature extraction on natural images



# Pre-trained Models

- Training a model on ImageNet from scratch takes days or weeks.
- Many models trained on ImageNet and their weights are publicly available!
- Fine-tuning
  - Retraining the (some) parameters of the network (given enough data)
  - Truncate the last layer(s) of the pre-trained network
  - Freeze the remaining layers weights
  - Add a (linear) classifier on top and train it for a few epochs
  - Then fine-tune the whole network or the few deepest layers
  - Use a smaller learning rate when fine tuning