

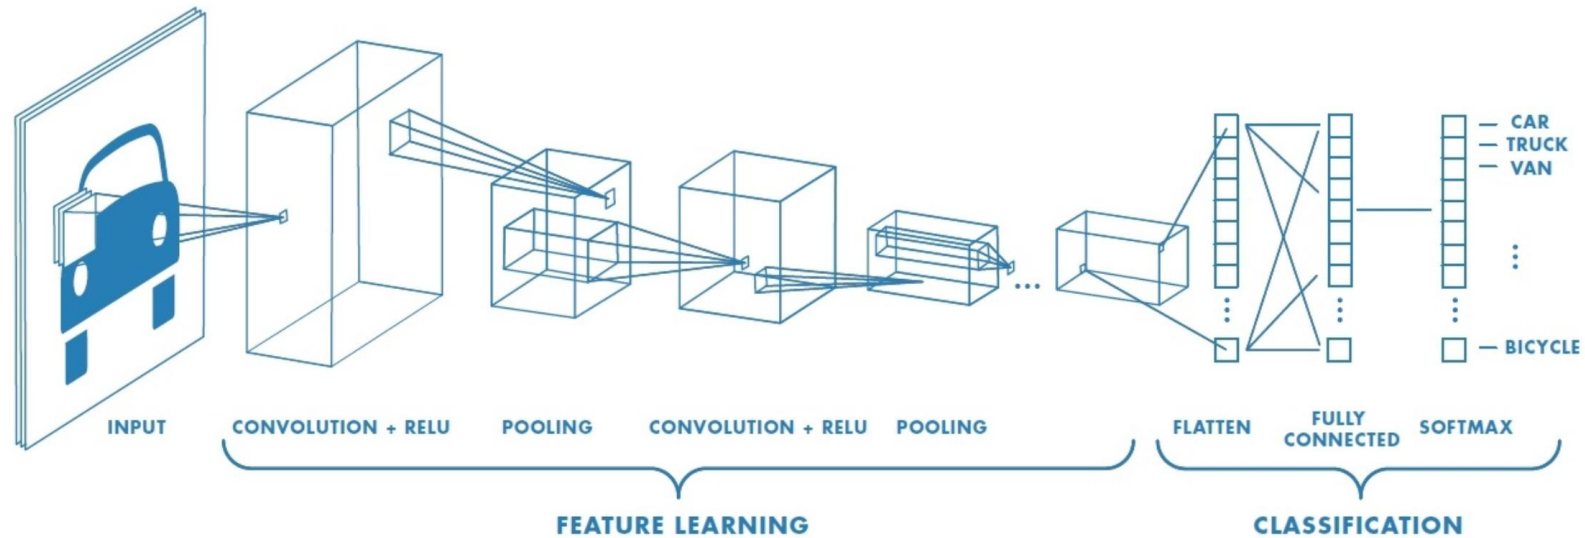


Beyond Convolutional Neural Networks (CNN)

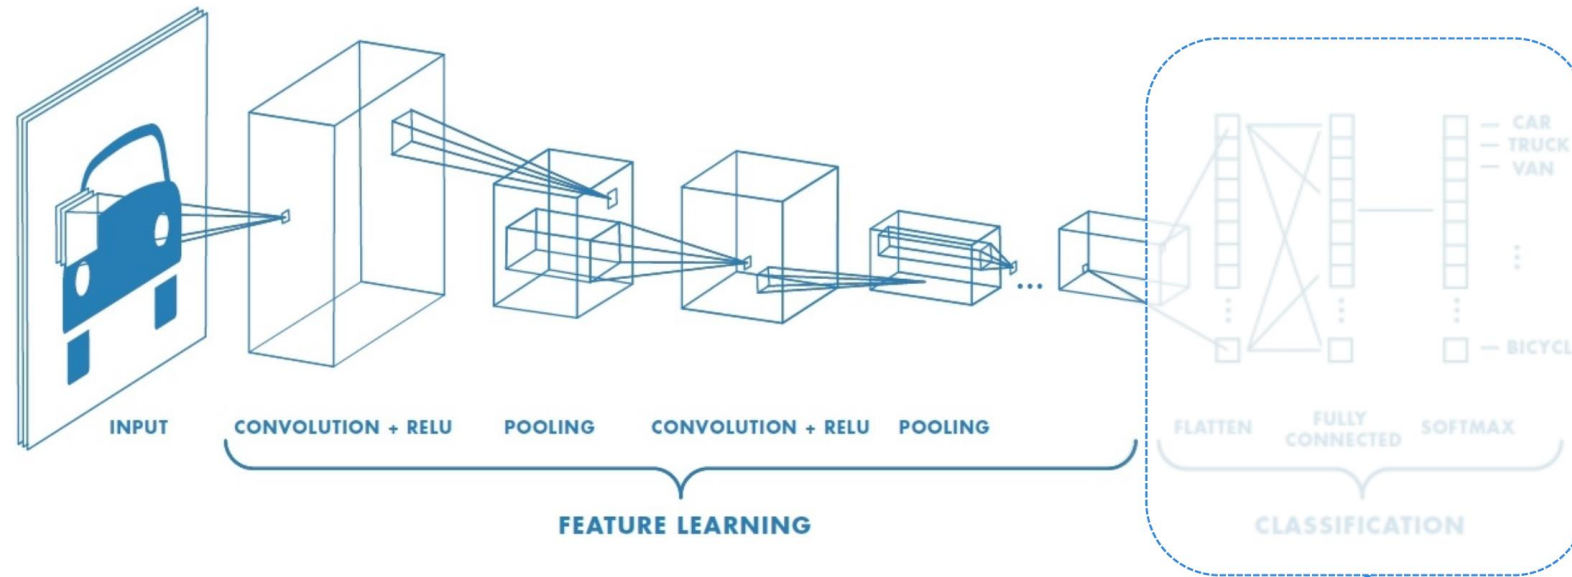
Industrial AI Lab.
Prof. Seungchul Lee

Beyond Classification

An Architecture for Many Applications




An Architecture for Many Applications



Different image classification domains
Object detection with R-CNNs
Segmentation with fully convolutional networks
Image captioning with RNNs


Beyond Classification

Semantic Segmentation



CAT

Object Detection



CAT, DOG, DUCK

Image Captioning



The cat is in the grass.

Instance Segmentation



CAT, DOG, DUCK


Classification + Localization



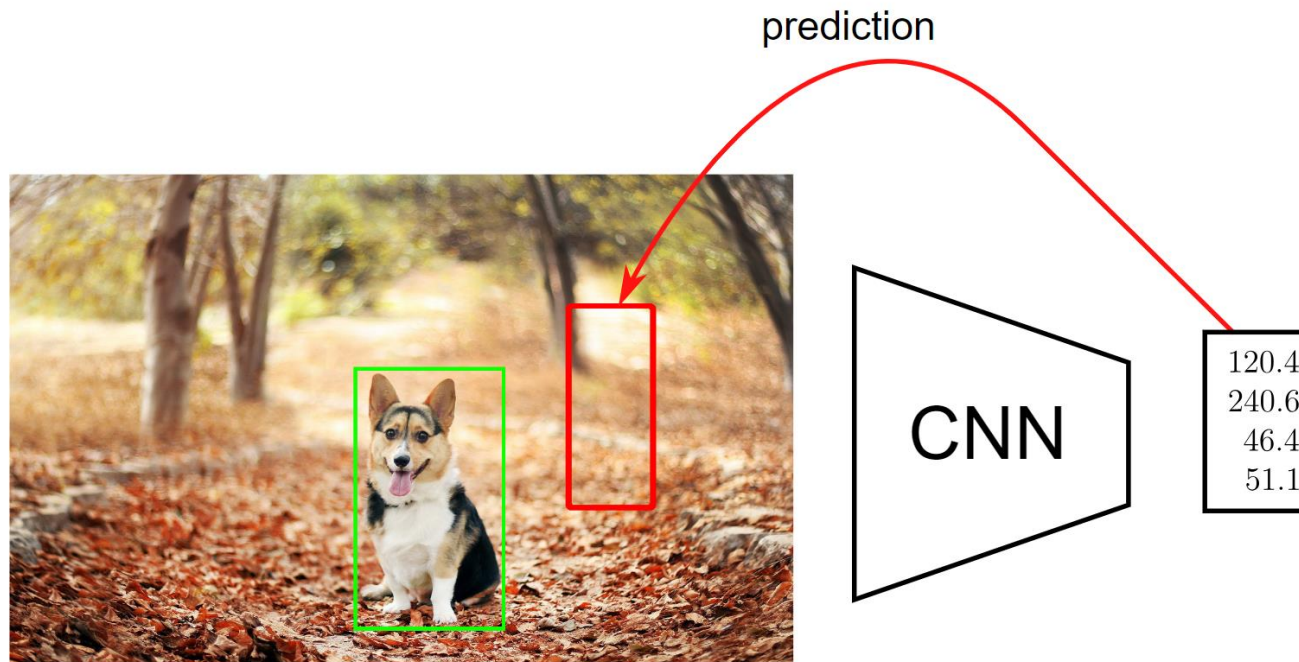
CAT

Localization

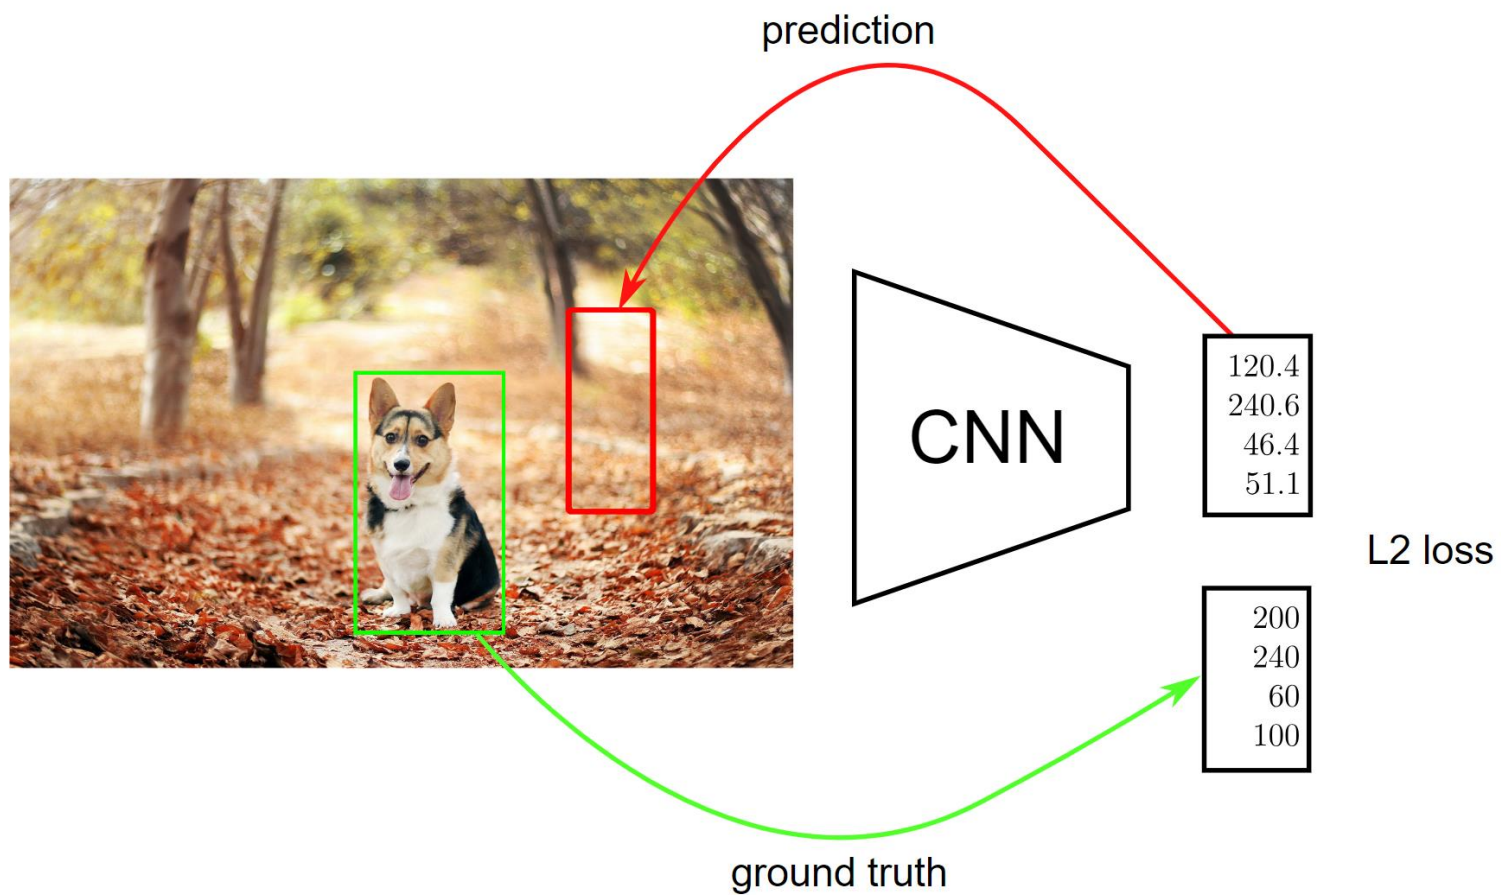
- Single object per image
- Predict coordinates of a bounding box (x, y, w, h)
- Evaluate via Intersection over Union (IoU)



Localization as Regression




Localization as Regression



Classification + Localization

- Use a pre-trained CNN on ImageNet (ex. ResNet)
- The "localization head" is trained separately with regression
- Possible end-to-end fine-tuning of both tasks

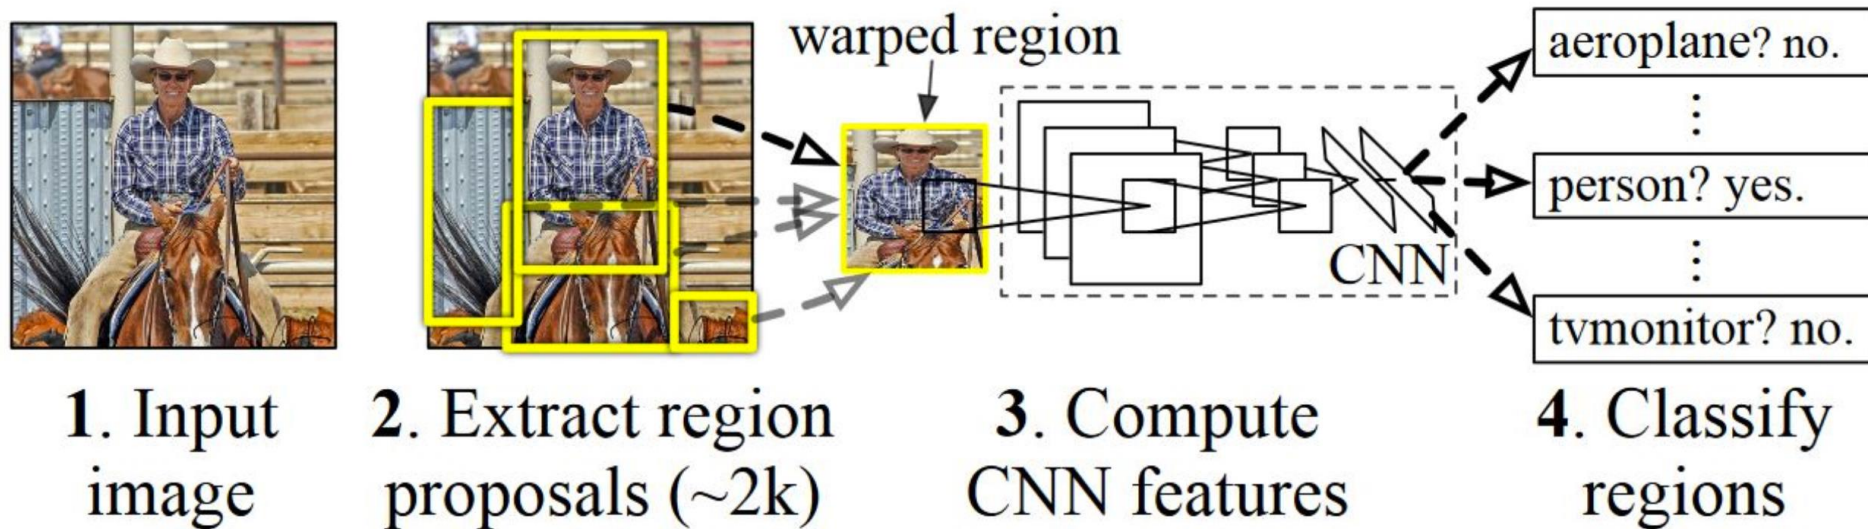


Object Detection

- We don't know in advance the number of objects in the image.
- Object detection relies on object proposal and object classification
 - Object proposal: find regions of interest (ROIs) in the image
 - Object classification: classify the object in these regions
- Two main families:
 - A grid in the image where each cell is a proposal (SSD, YOLO)
 - Region proposal (SPP, MultiBox, Faster RCNN, ...)

Object Detection with R-CNNs

- R-CNN:
 - Find regions that we think have objects.
 - Use CNN to classify



Segmentation

- Output a class map for each pixel (here: dog vs background)



- Instance segmentation: specify each object instance as well (two dogs have different instances)
- This can be done through object detection + segmentation

Semantic Segmentation: FCNs

- FCN: Fully Convolutional Network
 - Network designed with all convolutional layers, with down-sampling and up-sampling operations




Image Captioning using RNNs


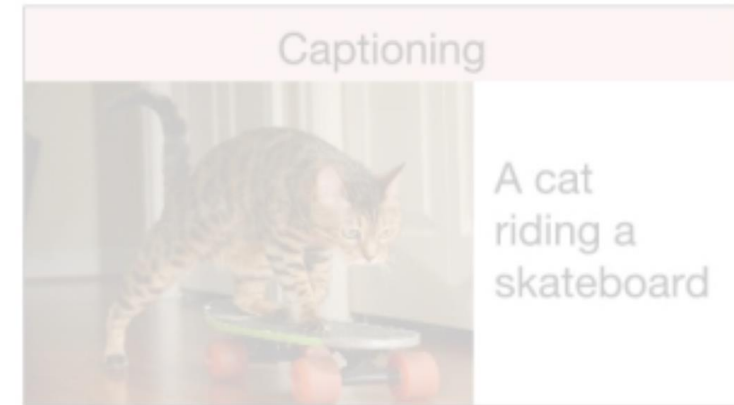



Image Captioning using RNNs

