




(Artificial) Neural Networks: From Perceptron to MLP


**Industrial AI Lab.
Prof. Seungchul Lee**

Perceptron: Example




$$\begin{aligned}\hat{y} &= g(\omega_0 + X^T \omega) \\ &= g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 3 \\ -2 \end{bmatrix}\right) \\ &= g(1 + 3x_1 - 2x_2)\end{aligned}$$


Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example




$$\hat{y} = g(1 + 3 \times (-1) - 2 \times 2) = g(-6) = -1$$


$$\hat{y} = g(1 + 3x_1 - 2x_2)$$




Perceptron: Example



$$\hat{y} = g(1 + 3x_1 - 2x_2)$$



Perceptron: Forward Propagation




$$\hat{y} = g(\omega_0 + X^T \omega)$$

$$= g\left(\omega_0 + \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}^T \begin{bmatrix} \omega_1 \\ \vdots \\ \omega_m \end{bmatrix}\right)$$

From Perceptron to MLP

Artificial Neural Networks: Perceptron


- Perceptron for $h(\theta)$ or $h(\omega)$
 - Neurons compute the weighted sum of their inputs
 - A neuron is activated or fired when the sum a is positive



- A step function is not differentiable
- One neuron is often not enough
 - One hyperplane

$$a = \omega_0 + \omega_1 x_1 + \omega_2 x_2$$

$$\hat{y} = g(a) = \begin{cases} 1 & a > 0 \\ 0 & \text{otherwise} \end{cases}$$






Here, a step function is illustrated instead of a sign function

XOR Problem

- Minsky-Papert Controversy on XOR
 - Not linearly separable
 - Limitation of perceptron


x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0



- Single neuron = **one linear classification boundary**




Artificial Neural Networks: MLP

- Multi-layer Perceptron (MLP) = Artificial Neural Networks (ANN)
 - Multi neurons = multiple linear classification boundaries




Artificial Neural Networks: Activation Function

- Differentiable nonlinear activation function




Common Activation Functions

Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$

 `tf.nn.sigmoid(z)`

Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$g'(z) = 1 - g(z)^2$$

 `tf.nn.tanh(z)`

Discuss later



Rectified Linear Unit (ReLU)




$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$

 `tf.nn.relu(z)`


Artificial Neural Networks

- In a compact representation



Artificial Neural Networks


- Multi-layer perceptron
 - Features of features
 - Mapping of mappings



- A single layer is not enough to be able to represent complex relationship between input and output
 \Rightarrow perceptron with many layers and units

Another Perspective: ANN as Kernel Learning


Nonlinear Classification



Neuron


- We can represent this “neuron” as follows:

$$f(x) = \sigma(w \cdot x + b)$$



XOR Problem


- The main weakness of linear predictors is their lack of capacity. For classification, the populations have to be linearly separable.



“xor”

Nonlinear Mapping


- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.


$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$



Nonlinear Mapping

- The XOR example can be solved by pre-processing the data to make the two populations linearly separable.

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$




Kernel

- Often we want to capture nonlinear patterns in the data
 - nonlinear regression: input and output relationship may not be linear
 - nonlinear classification: classes may not be separable by a linear boundary
- Linear models (e.g. linear regression, linear SVM) are not just rich enough
 - by mapping data to higher dimensions where it exhibits linear patterns
 - apply the linear model in the new input feature space
 - mapping = changing the feature representation
- Kernels: make linear model work in nonlinear settings

Kernel + Neuron


- Nonlinear mapping + neuron

$$\phi : (x_u, x_v) \rightarrow (x_u, x_v, x_u x_v)$$




Neuron + Neuron

- Nonlinear mapping can be represented by another neurons





- Nonlinear Kernel
 - Nonlinear activation functions

- Nonlinear mapping can be represented by another neurons
- We can generalize an MLP




Summary

- Universal function approximator
- Universal function classifier
- Parameterized




Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons





Deep Artificial Neural Networks

- Complex/Nonlinear universal function approximator
 - Linearly connected networks
 - Simple nonlinear neurons



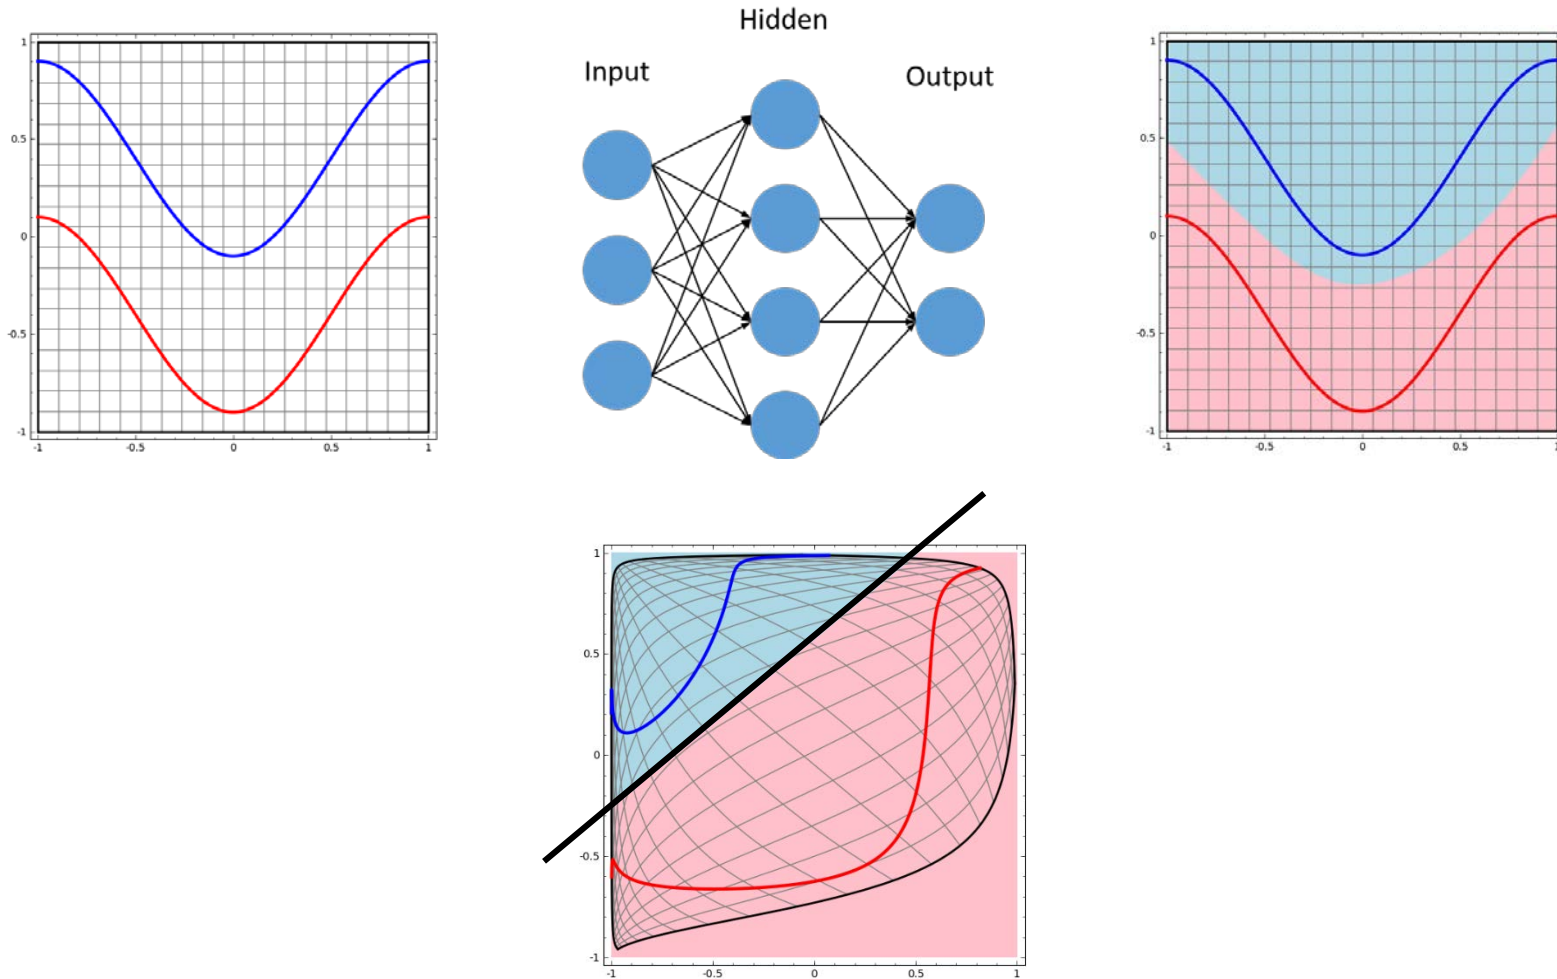
Example: Linear Classifier

- Perceptron tries to separate the two classes of data by dividing them with a line




Example: Neural Networks


- The hidden layer learns a representation so that the data gets linearly separable




Nonlinearly Distributed Data




Multi Layers




Multi Layers




Nonlinearly Distributed Data



Multi Layers



Nonlinearly Distributed Data



Multi Layers

