




# Modern CNNs

**Industrial AI Lab.**


**Prof. Seungchul Lee**

# ImageNet

- Human performance = 5.1 %

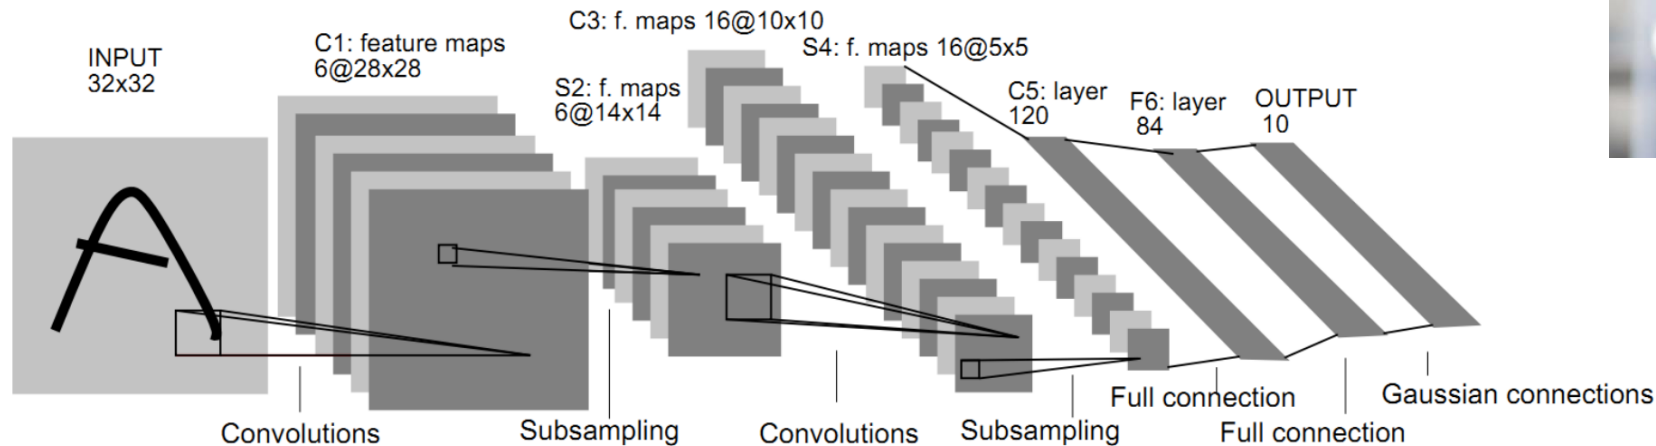


# ImageNet



# LeNet


- CNN = Convolutional Neural Networks = ConvNet
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. (1998). Gradient-based learning applied to document recognition.
- All are still the basic components of modern ConvNets!



Yann LeCun



# AlexNet

- Simplified version of Krizhevsky, Alex, Sutskever, and Hinton. "Imagenet classification with deep convolutional neural networks." NIPS 2012
- LeNet-style backbone, plus:
  - ReLU [Nair & Hinton 2010]
    - “RevoLUtion of deep learning”\*
    - Accelerate training; better grad prop (vs. tanh)
  - Dropout [Hinton et al 2012]
    - In-network ensembling
    - Reduce overfitting
  - Data augmentation
    - Label-preserving transformation
    - Reduce overfitting



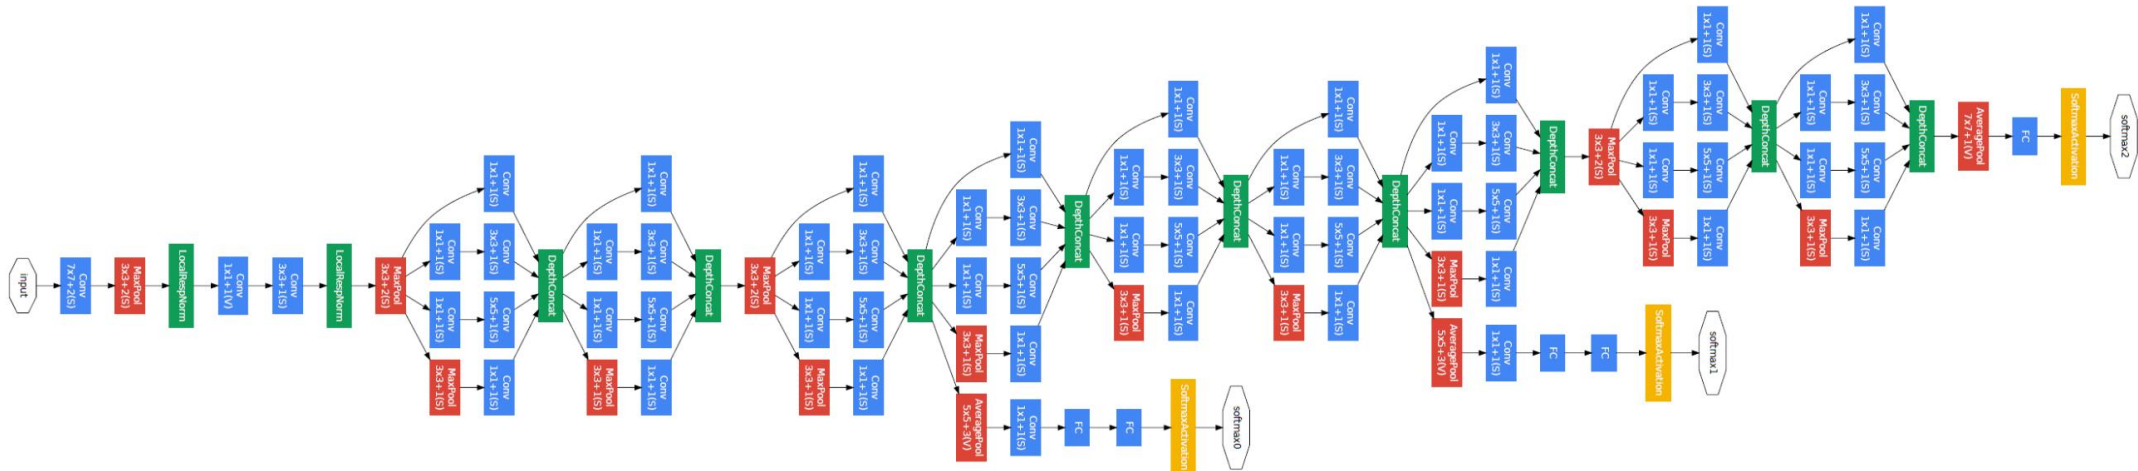

# VGG-16/19

- Simonyan, Karen, and Zisserman. "Very deep convolutional networks for large-scale image recognition." (2014)
- Simply "Very Deep"!
  - Modularized design
    - 3x3 Conv as the module
    - Stack the same module
    - Same computation for each module
  - Stage-wise training
    - VGG-11 → VGG-13 → VGG-16
    - We need a better initialization...



# GoogleNet/Inception


- Multiple branches
  - e.g., 1x1, 3x3, 5x5, pool
- Shortcuts
  - stand-alone 1x1, merged by concat.
- Bottleneck
  - Reduce dim by 1x1 before expensive 3x3/5x5 conv



# ResNet (Deep Residual Learning)

- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Plane net


$H(x)$  is any desired mapping,  
hope the small subnet fit  $H(x)$





# ResNet (Deep Residual Learning)

- He, Kaiming, et al. "Deep residual learning for image recognition." CVPR. 2016.
- Residual net
- Skip connection




$H(x)$  is any desired mapping,  
~~hope the small subnet fit  $H(x)$~~   
hope the small subnet fit  $F(x)$   
Let  $H(x) = F(x) + x$


- A direct connection between 2 non-consecutive layers
- No vanishing gradient

# ResNet (Deep Residual Learning)

- Parameters are optimized to learn a residual, that is the difference between the value before the block and the one needed after.
- $F(x)$  is a residual mapping w.r.t. identity



- If identity were optimal, easy to set weights as 0
- If optimal mapping is closer to identity, easier to find small fluctuations




# Skip Connection

- A skip connection is a connection that bypasses at least one layer.
- Here, it is often used to transfer local information by concatenating or summing feature maps from the downsampling path with feature maps from the upsampling path.
- Merging features from various resolution levels helps combining context information with spatial information.


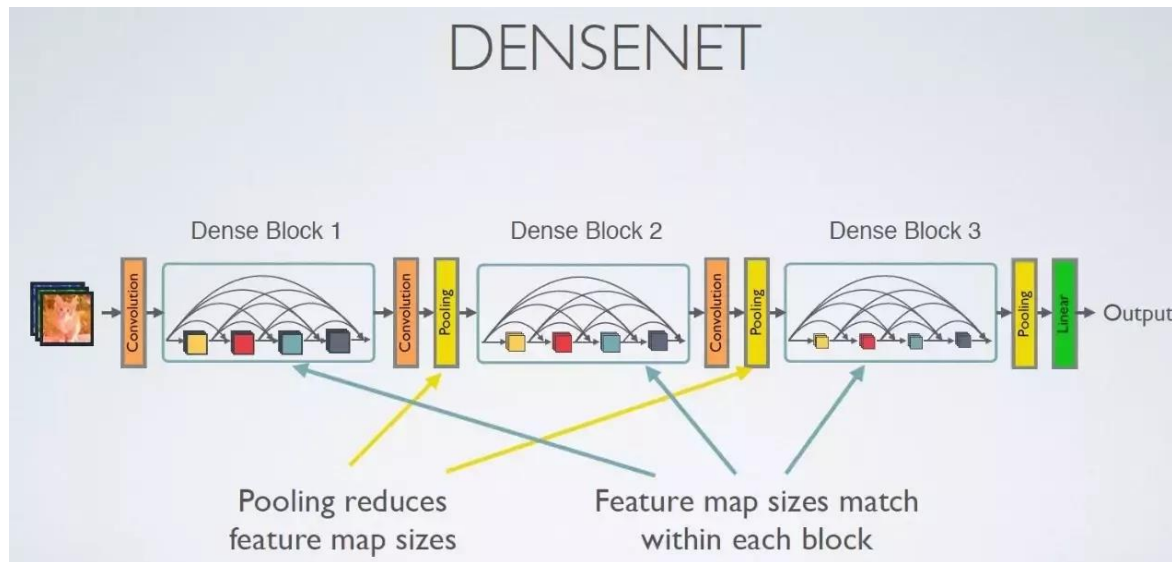
# Residual Net

```
def residual_net(x):  
    conv1 = tf.layers.conv2d(inputs = x,  
                              filters = 32,  
                              kernel_size = [3, 3],  
                              padding = "SAME",  
                              activation = tf.nn.relu)  
  
    conv2 = tf.layers.conv2d(inputs = conv1,  
                              filters = 32,  
                              kernel_size = [3, 3],  
                              padding = "SAME",  
                              activation = tf.nn.relu)  
  
→ maxp2 = tf.layers.max_pooling2d(inputs = x + conv2,  
                                   pool_size = [2, 2],  
                                   strides = 2)  
  
    flat = tf.layers.flatten(maxp2)  
    hidden = tf.layers.dense(inputs = flat,  
                              units = n_hidden,  
                              activation = tf.nn.relu)  
    output = tf.layers.dense(inputs = hidden,  
                              units = n_output)  
  
    return output
```

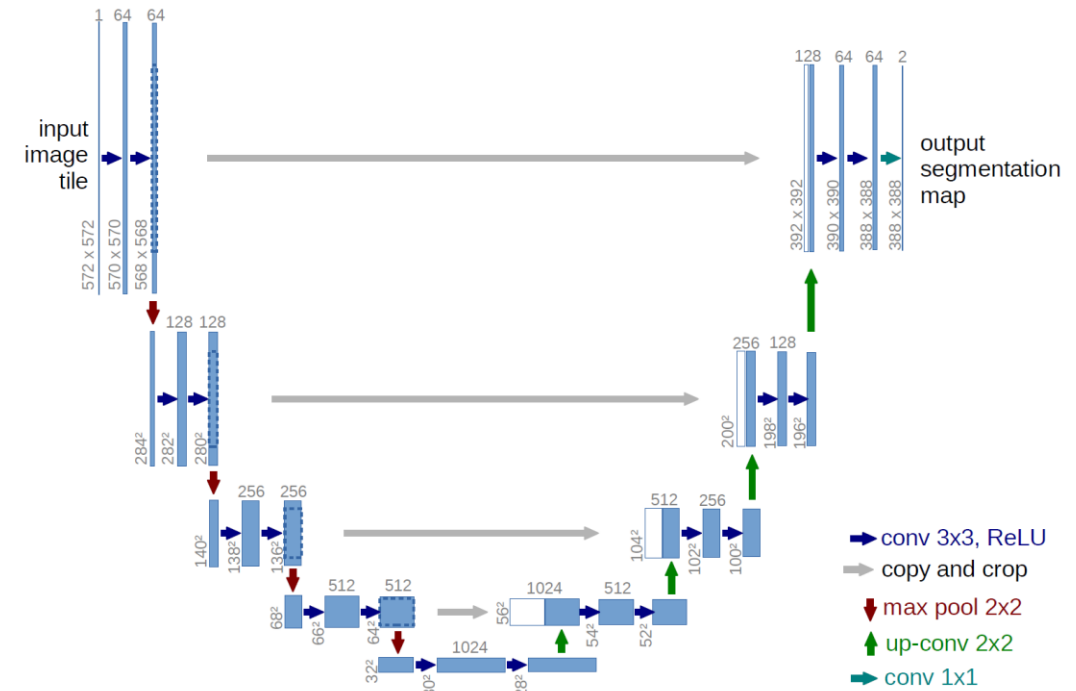
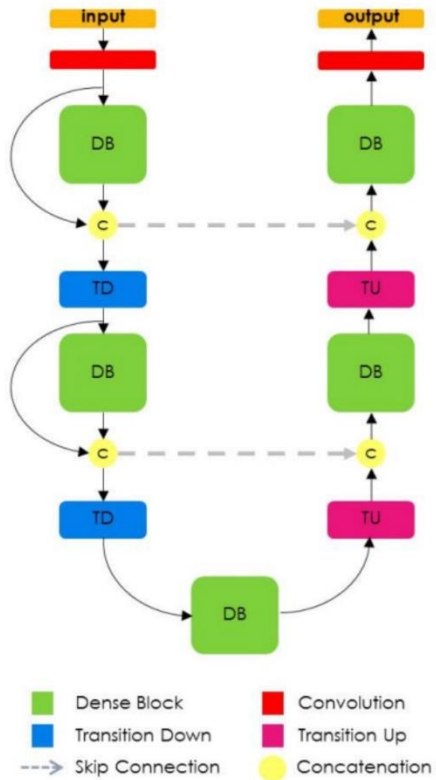


# DensNets

- Densely Connected Convolutional Networks



# U-Net



# U-Net

- The U-Net owes its name to its symmetric shape
- The U-Net architecture is built upon the Fully Convolutional Network and modified in a way that it yields better segmentation in medical imaging.
- Compared to FCN-8, the two main differences are
  - U-net is symmetric and
  - the skip connections between the downsampling path and the upsampling path apply a concatenation operator instead of a sum.
- These skip connections intend to provide local information to the global information while upsampling. Because of its symmetry, the network has a large number of feature maps in the upsampling path, which allows to transfer information.