Recurrent Neural Network (RNN)

Industrial AI Lab.

Prof. Seungchul Lee

Recurrent NN (RNN)

- Hidden state extraction and transformation

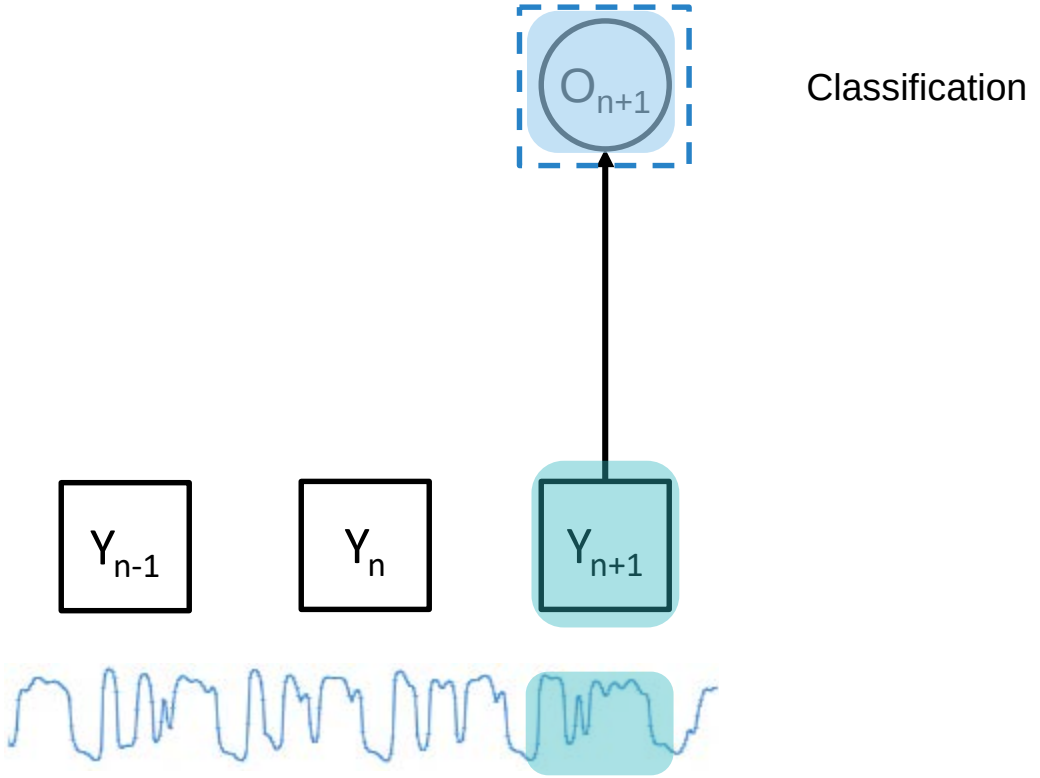

Recurrent NN (RNN)

- Hidden state extraction and transformation

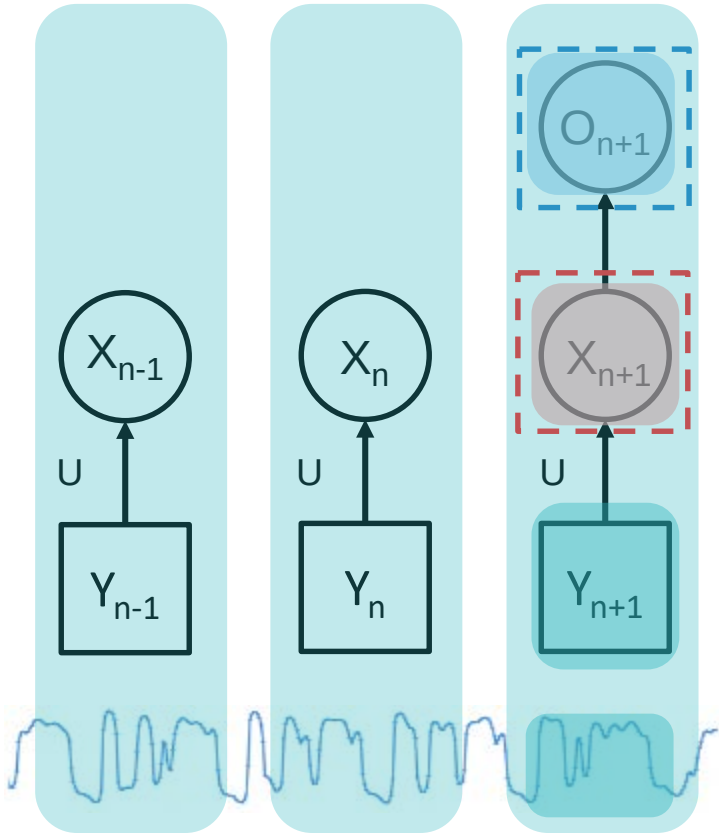

Classification based on states

Learned latent state

Recurrent NN (RNN)

- Hidden state extraction and transformation
- Good for sequential data (dynamic behavior)

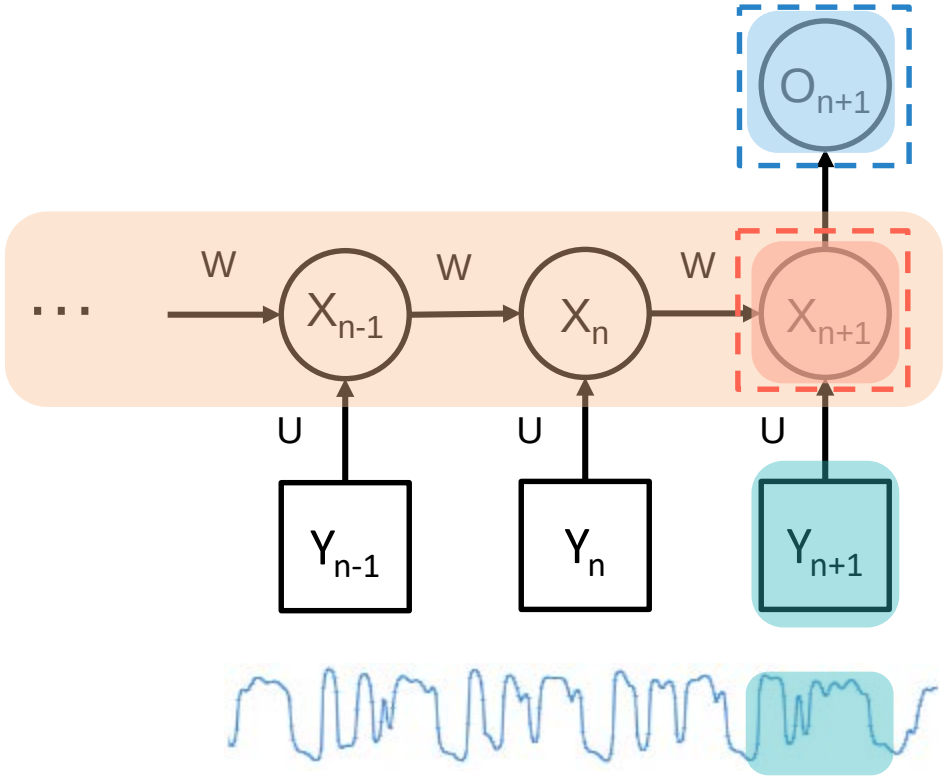

Classification based on states

Learned latent state and its dynamics

Recurrent NN

- Recurrence
  - Consider the classical form of a dynamical system:

$$s^{(t)} = f(s^{(t-1)}; \theta)$$

  - This is recurrent because the definition of s at time t refers back to the same definition at time $t-1$

- Hidden state representation

- Learn both from sequential data

Representation Shortcut

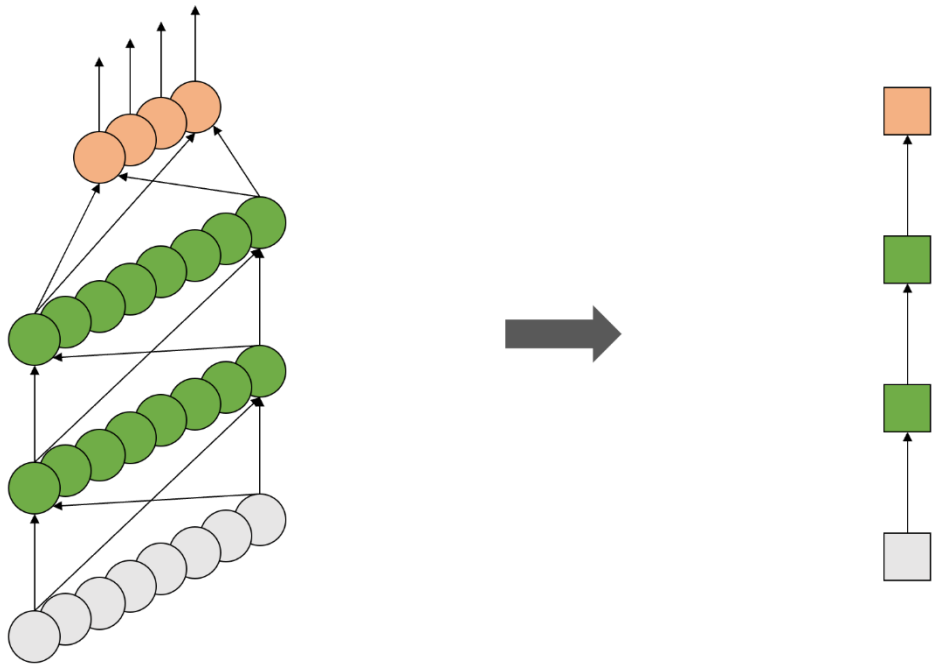

- Input at each time is a vector
- Each layer has many neurons
  - Output layer too may have many neurons
- But will represent everything simple boxes
  - Each box actually represents an entire layer with many units

Representation Shortcut

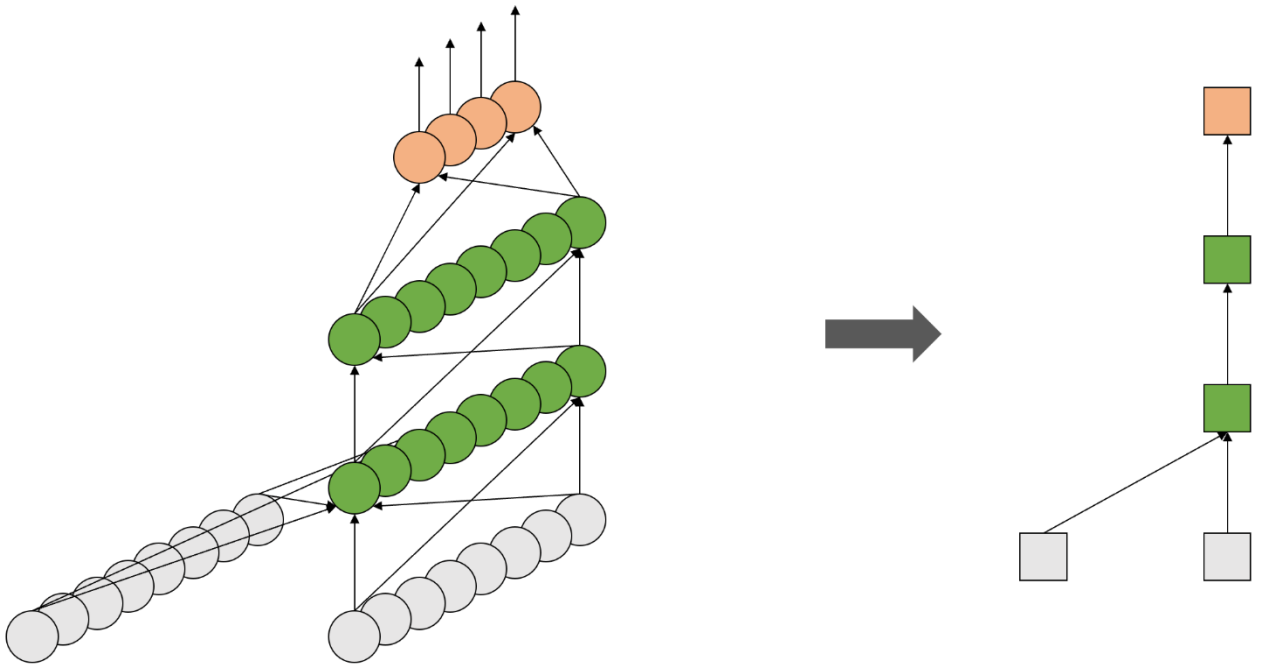

- Input at each time is a vector
- Each layer has many neurons
  - Output layer too may have many neurons
- But will represent everything simple boxes
  - Each box actually represents an entire layer with many units

Representation Shortcut

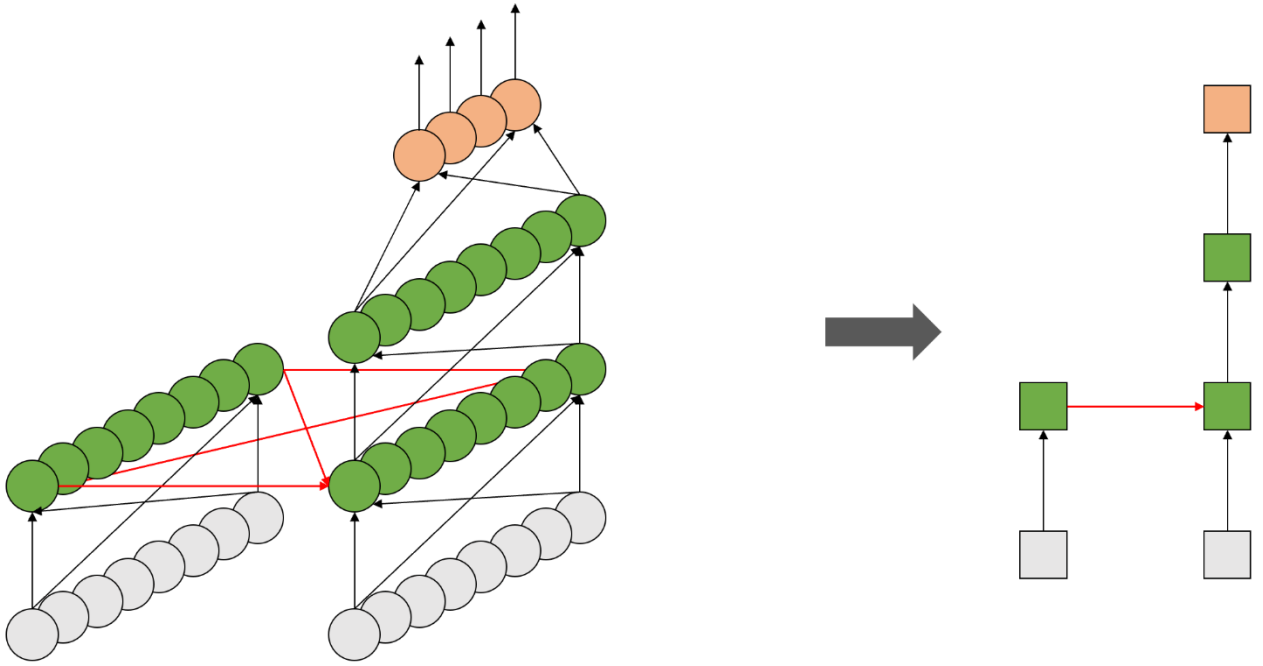

- Input at each time is a vector
- Each layer has many neurons
  - Output layer too may have many neurons
- But will represent everything simple boxes
  - Each box actually represents an entire layer with many units

Sliding Predictor

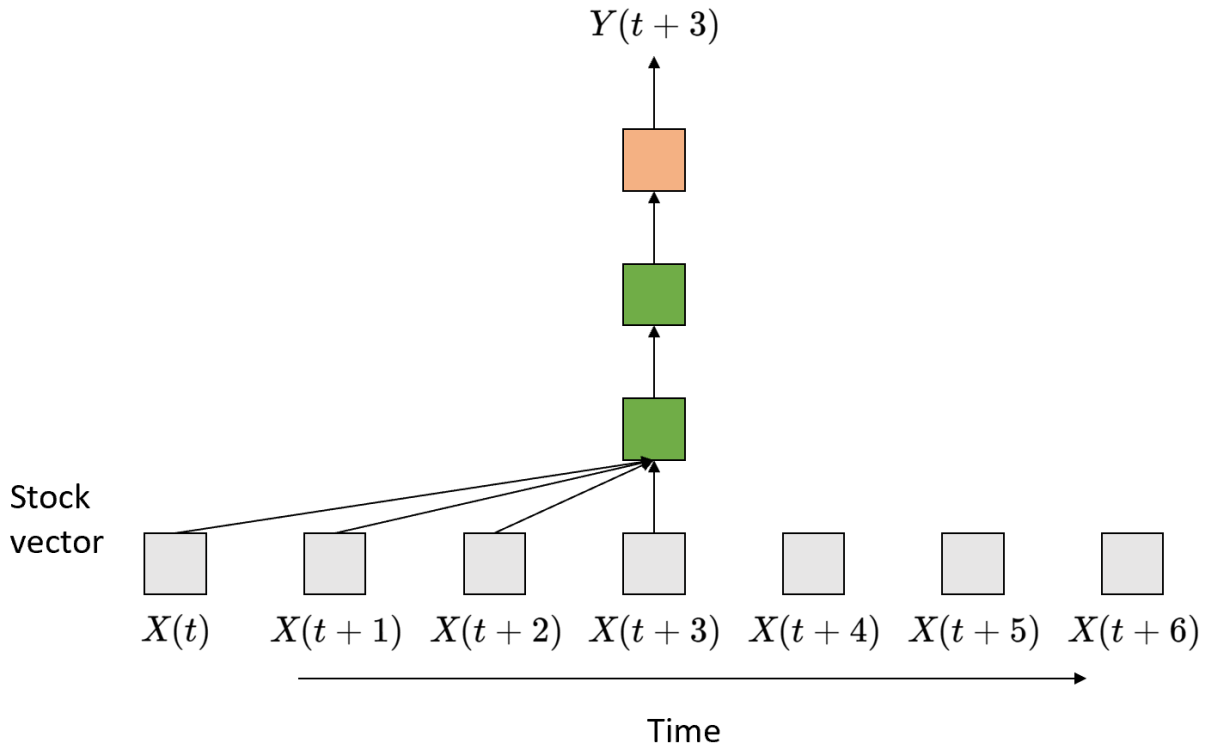

- The sliding predictor
  - Look at the last few days
  - This is just a convolutional neural net applied to sequential data

Sliding Predictor

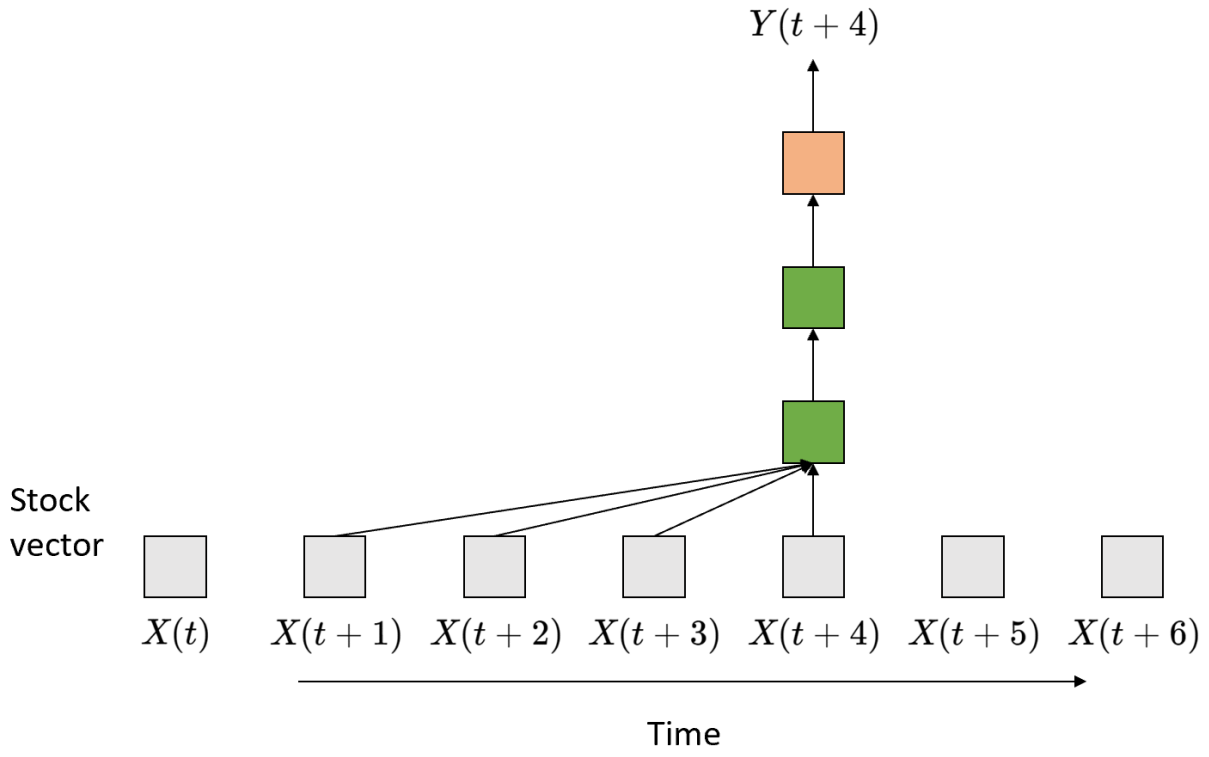

- The sliding predictor
  - Look at the last few days
  - This is just a convolutional neural net applied to sequential data

Sliding Predictor

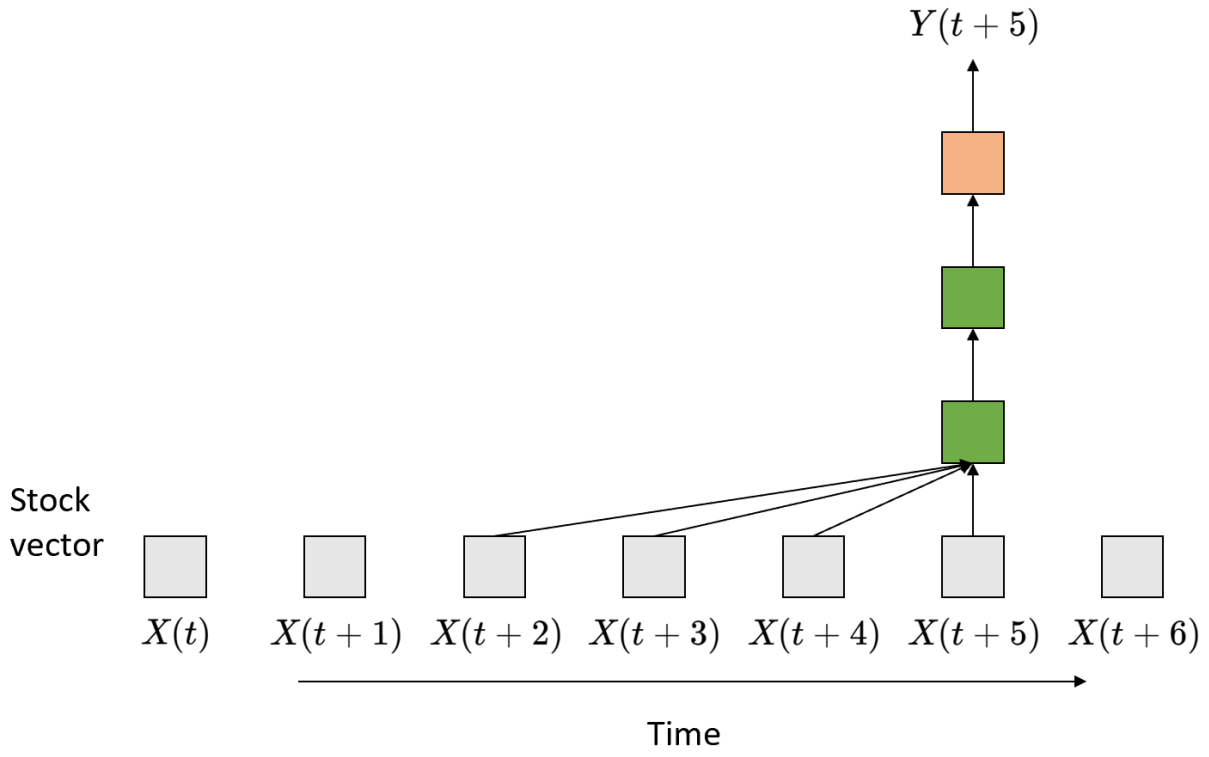

- The sliding predictor
  - Look at the last few days
  - This is just a convolutional neural net applied to sequential data

Sliding Predictor

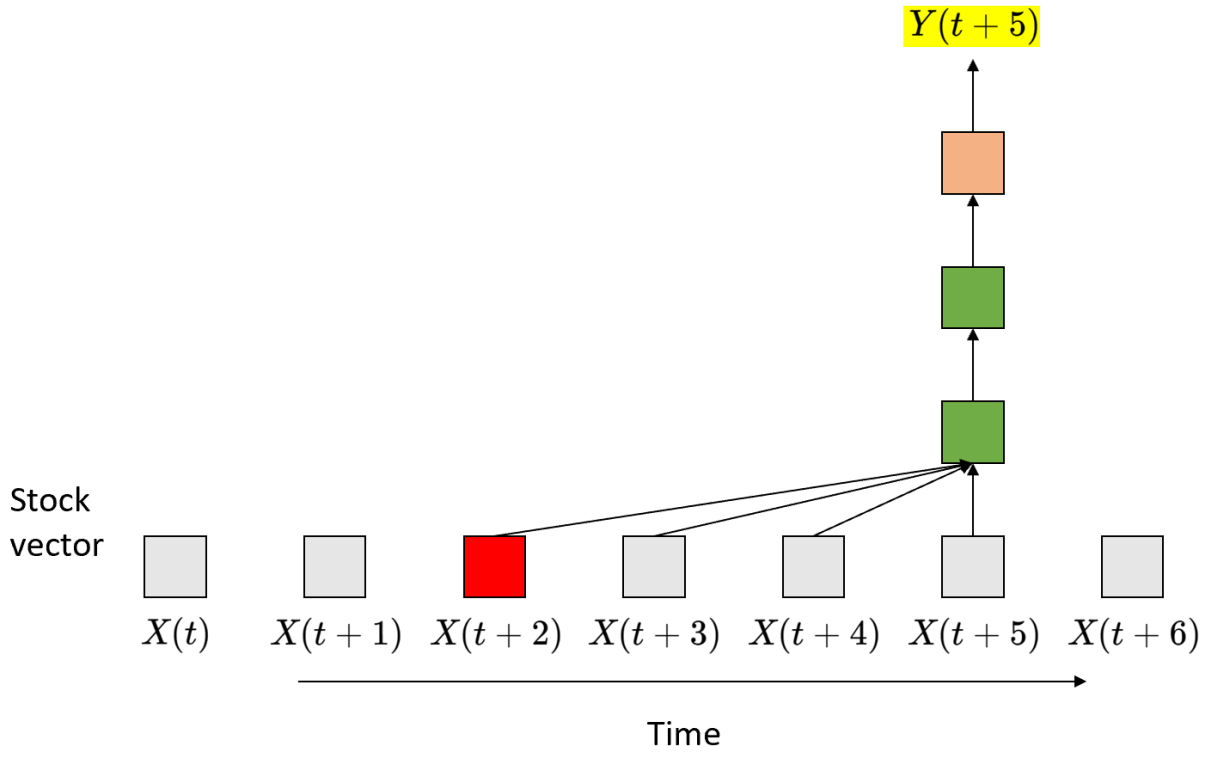

- The sliding predictor
  - Look at the last few days
  - This is just a convolutional neural net applied to sequential data

Finite-Response Model

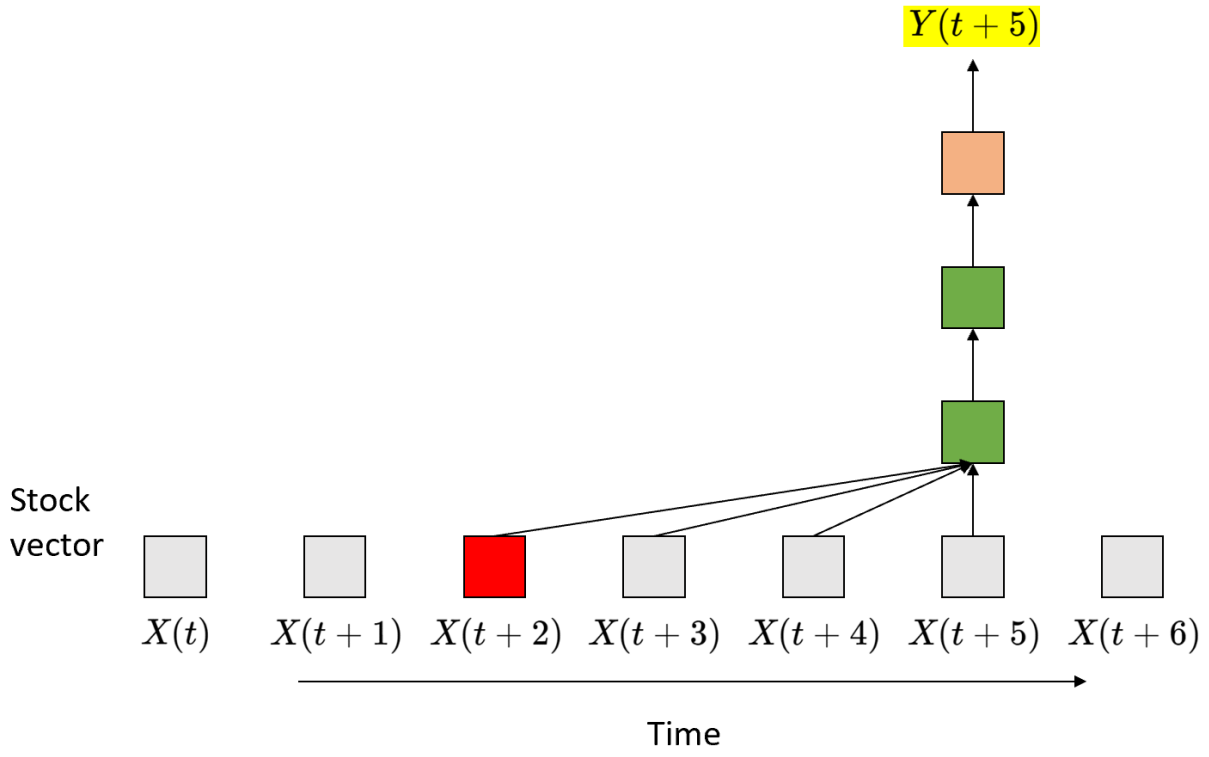

- This is a finite response system
  - Something that happens today only affects the output of the system for N days into the future
  - N is the width of the system

Finite-Response Model

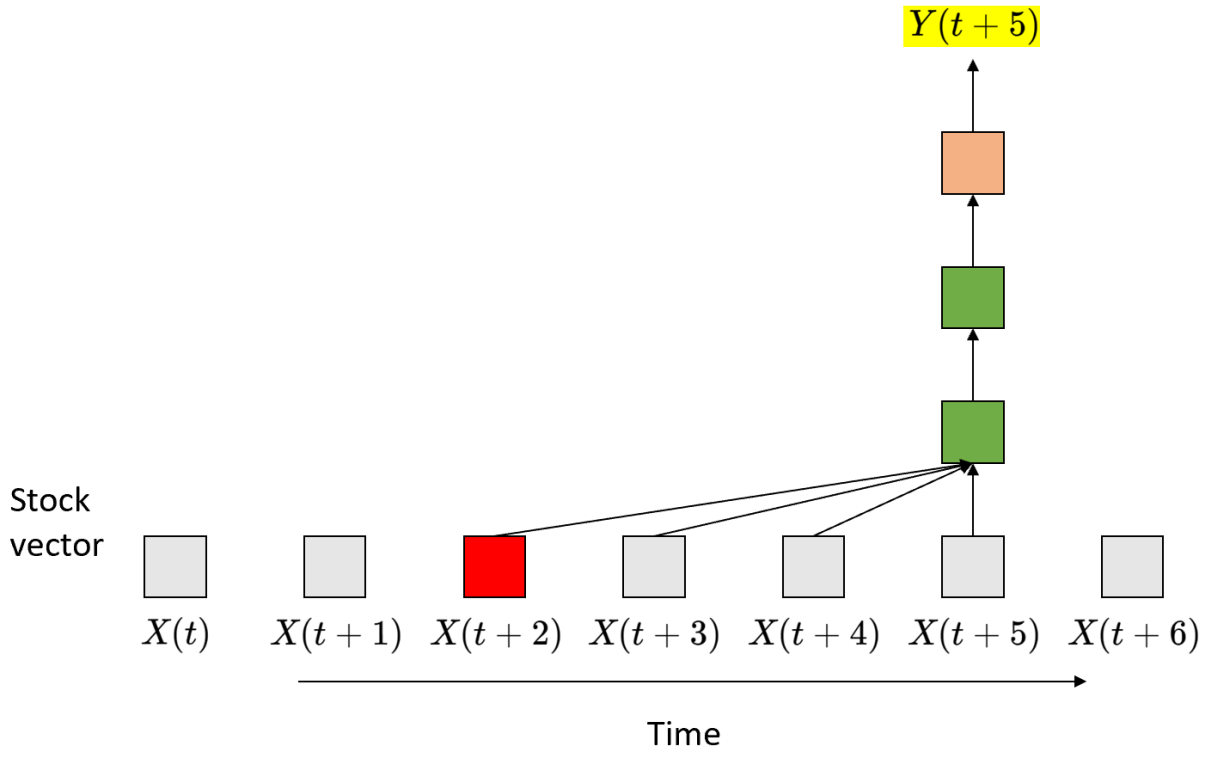

- Problem: Increasing the “history” makes the network more complex

$$Y_t = f(X_t, X_{t-1}, \ldots, X_{t-N})$$

In Theory, We Want Infinite Memory

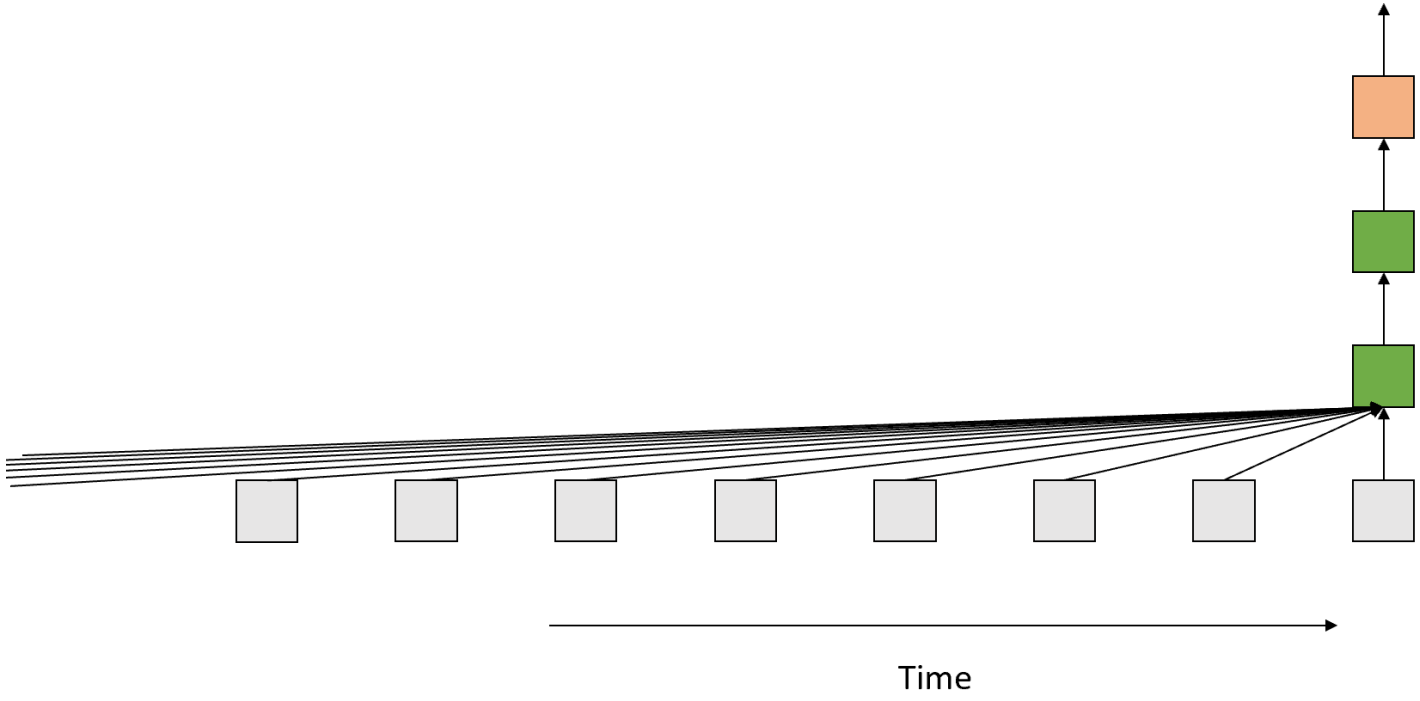

- Required: Infinite response systems
  - What happens today can continue to affect the output forever
  - Possibly with weaker and weaker influence

$$Y_t = f\left(X_t, X_{t-1}, \cdots, X_{t-\infty}\right)$$

Infinite Response Systems

$$Y_t = f(X_t, X_{t-1}, \cdots, X_{t-\infty}) \quad \Longrightarrow \quad Y_t = f(X_t, Y_{t-1})$$

- Recursive
  - Required: Define initial output: Y_{t-1} for $t = 0$
  - An input at X_0 at $t = 0$ produces Y_0
  - Y_0 produces Y_1 which produces Y_2 and so on until Y_∞ even if $X_1, \cdots, X_\infty$ are 0
  - Nonlinear autoregressive

- Output contains information about the entire past

Autoregression

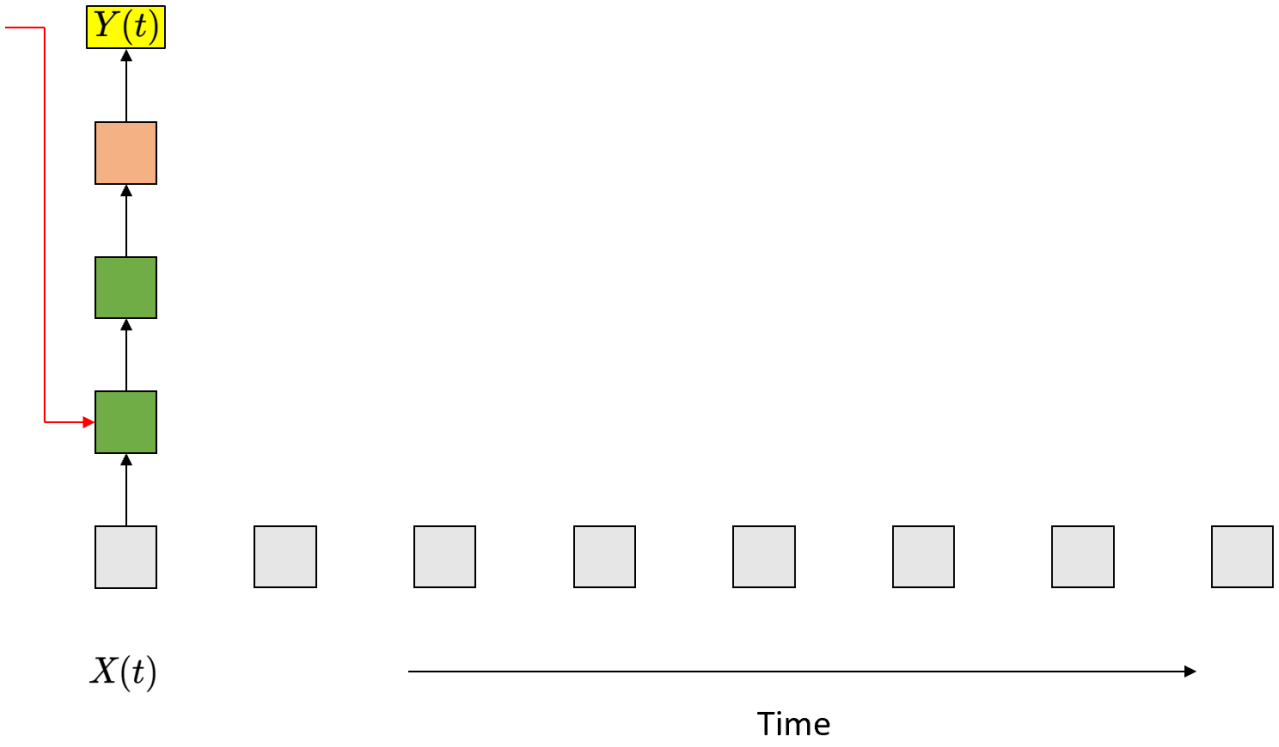

- An autoregressive net with recursion from the output

Autoregression

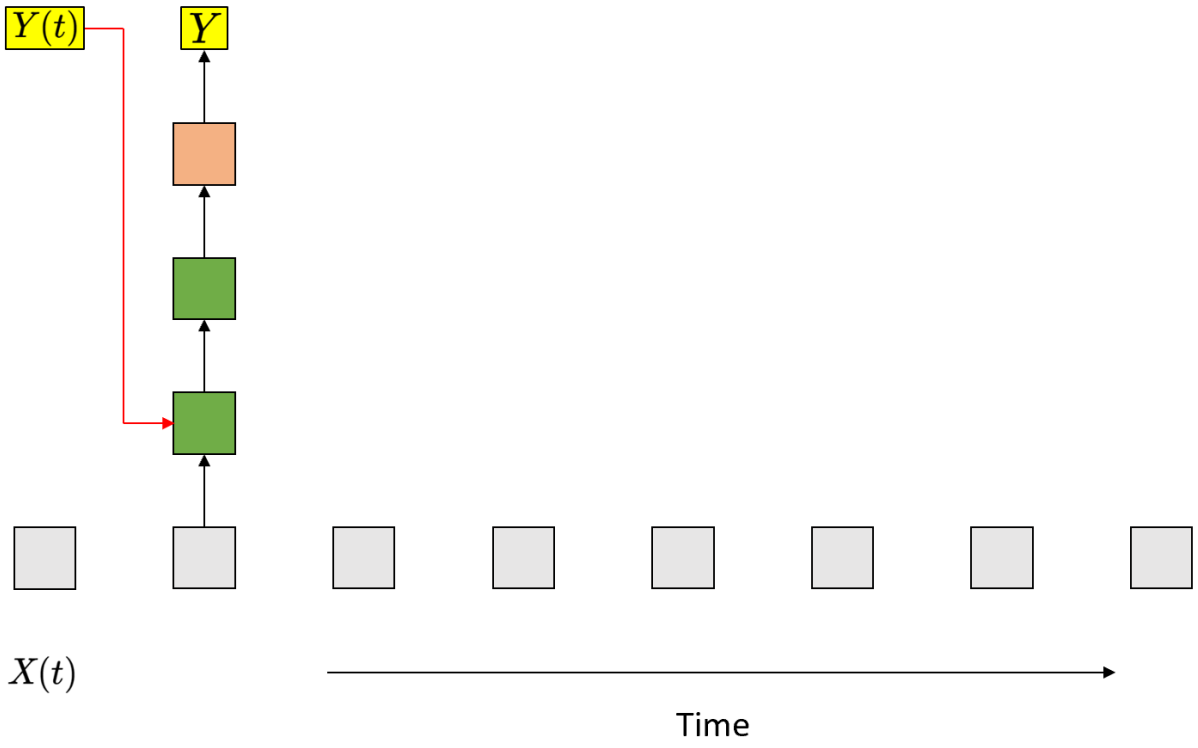

- An autoregressive net with recursion from the output

Autoregression

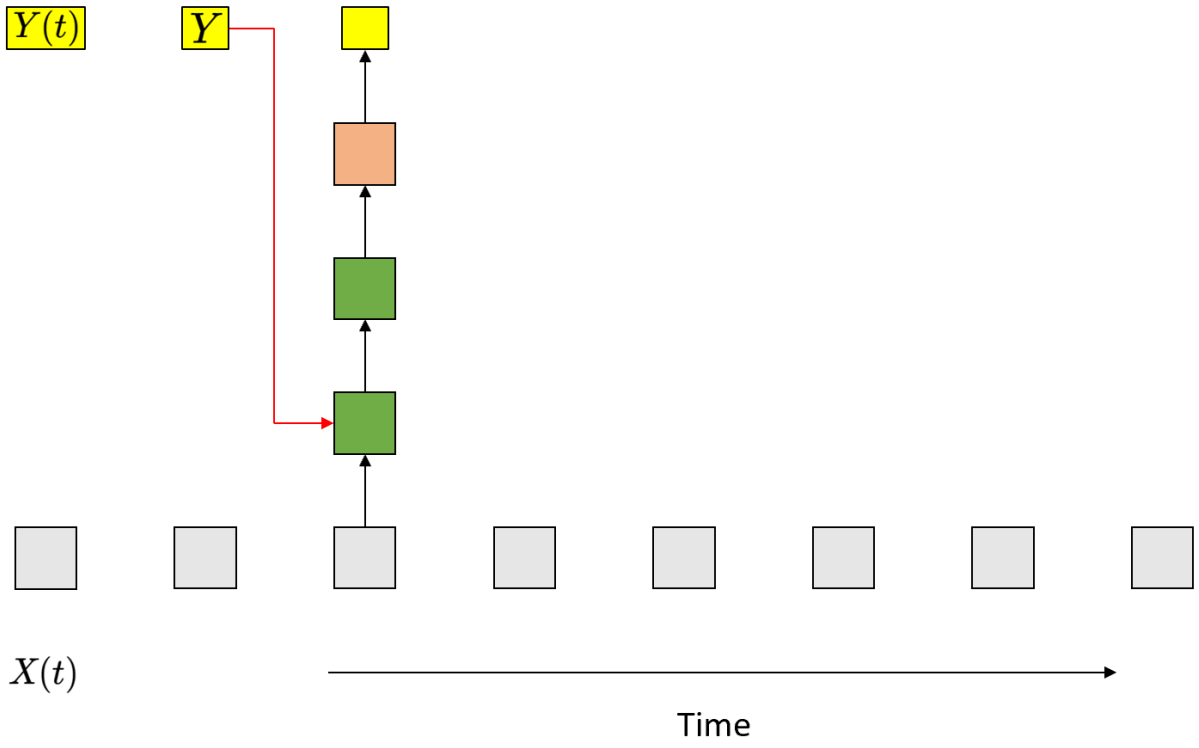

- An autoregressive net with recursion from the output

Autoregression

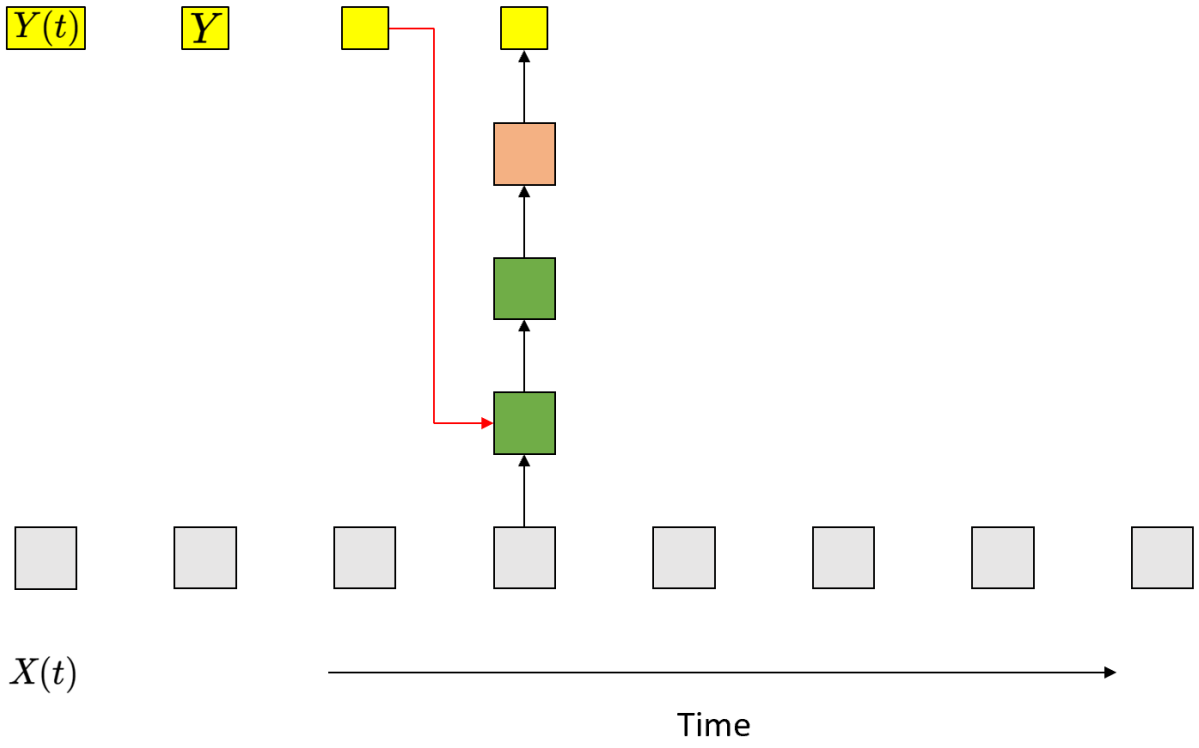

- An autoregressive net with recursion from the output

Autoregression

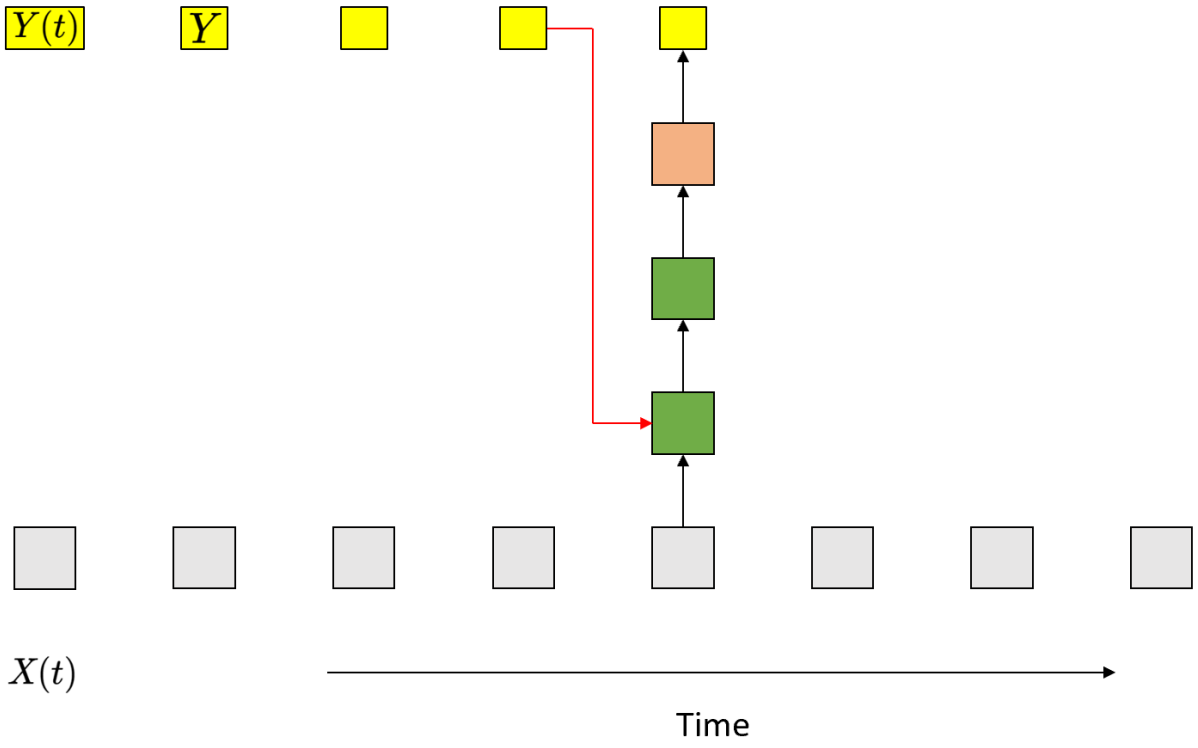

- An autoregressive net with recursion from the output

Autoregression

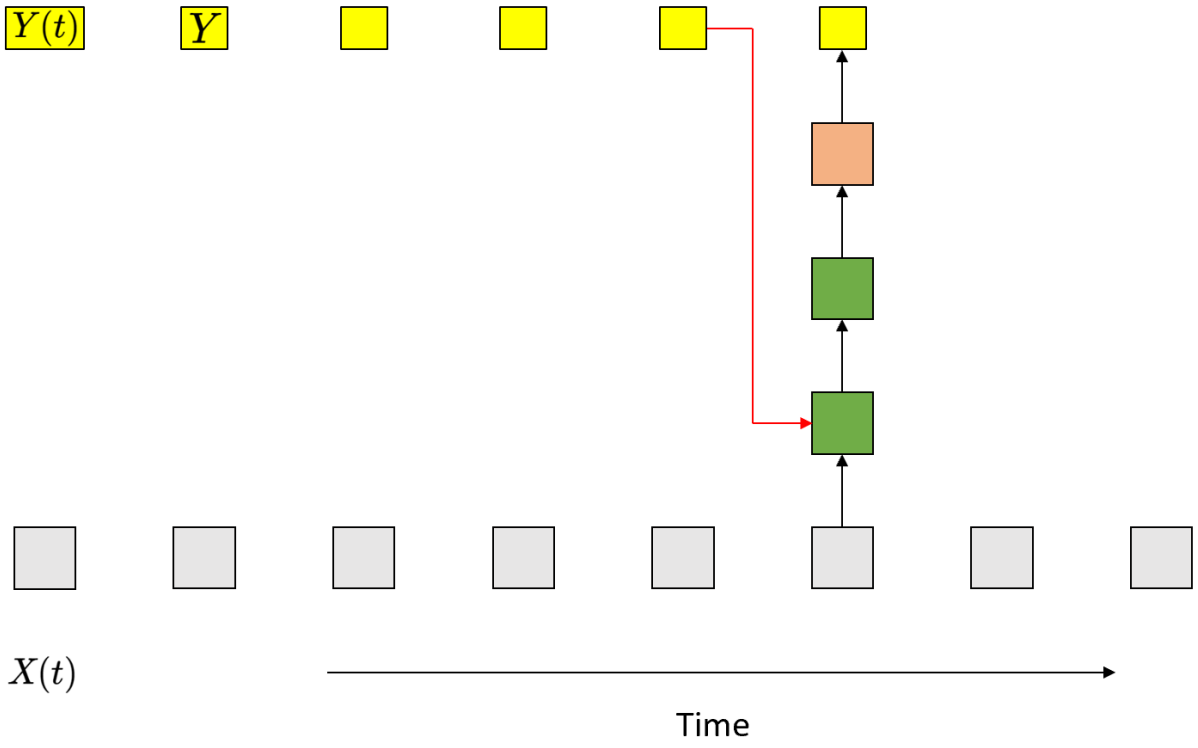

- An autoregressive net with recursion from the output

Autoregression

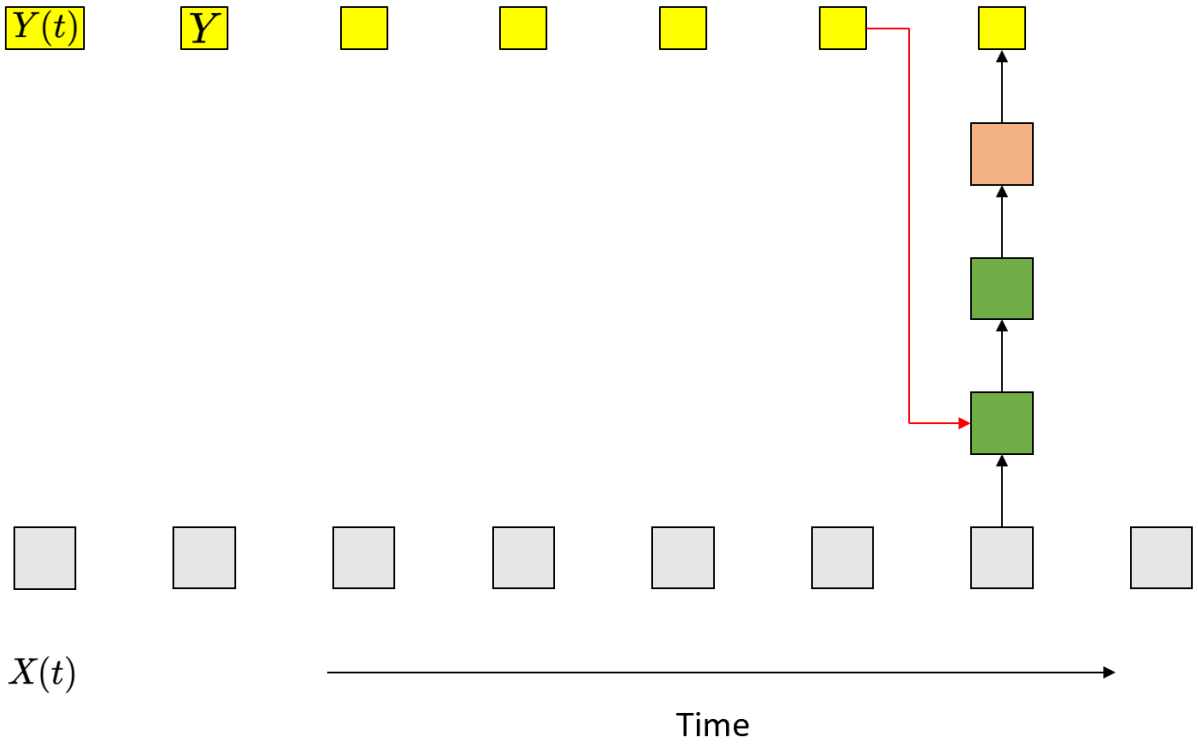

- An autoregressive net with recursion from the output

Autoregression

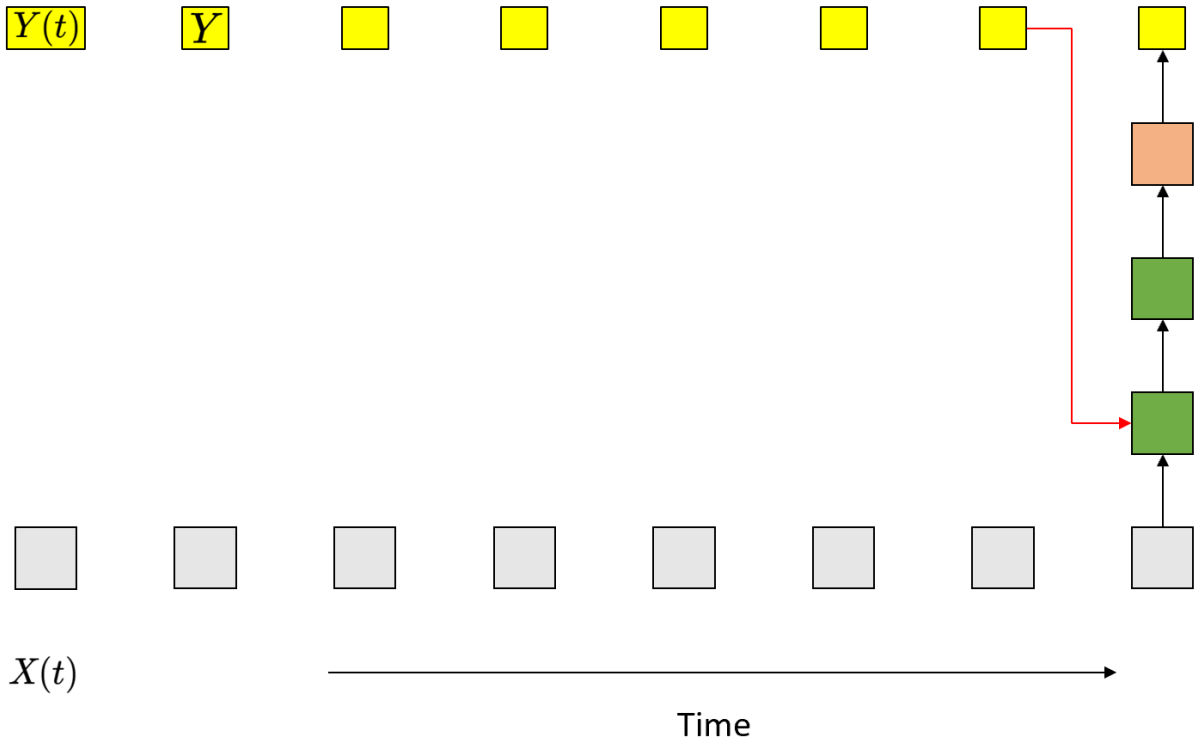

- An autoregressive net with recursion from the output

More Complete Representation

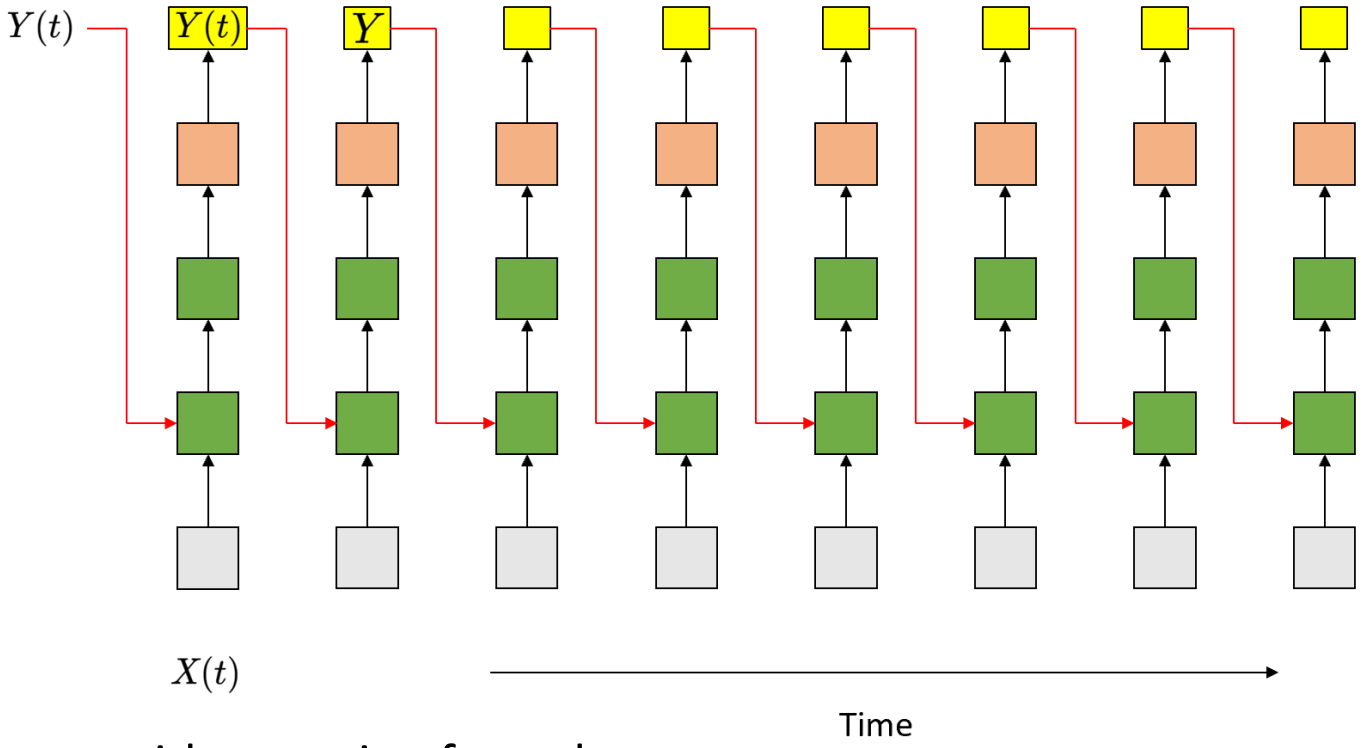

- An autoregressive net with recursion from the output
- Showing all computations
- All columns are identical
- An input at $t = 0$ affects outputs forever

An Alternate Model for Infinite Response Systems

- the state-space model

$$h_t = f(x_t, h_{t-1})$$
$$y_t = g(h_t)$$

- h_t is the state of the network
- Need to define initial state h_{-1}

- This is a recurrent neural network
- State summarizes information about the entire past

Single Hidden Layer RNN (Simplest State-Space Model)

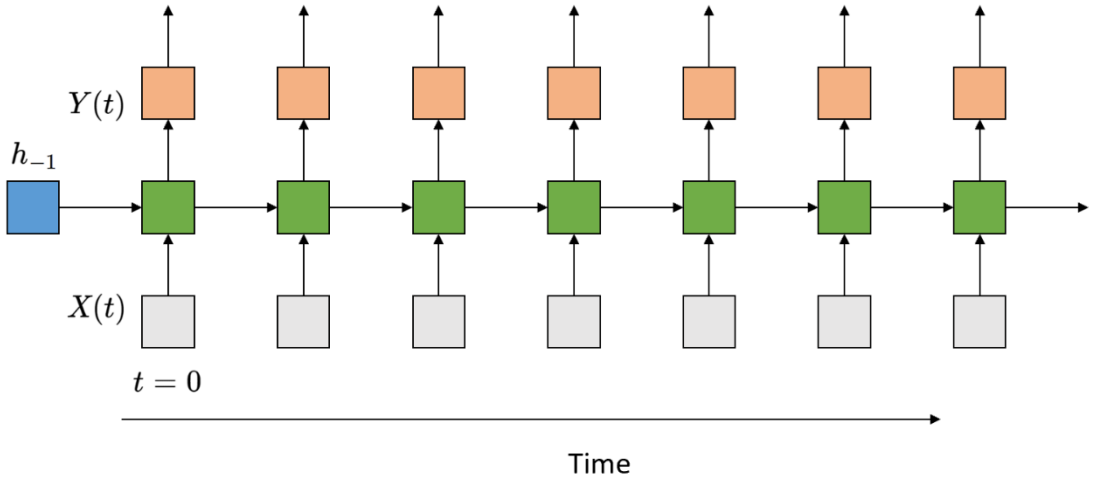

- The state (green) at any time is determined by the input at that time, and the state at the previous time
- All columns are identical
- An input at $t = 0$ affects outputs forever
- Also known as a recurrent neural net

Multiple Recurrent Layer RNN

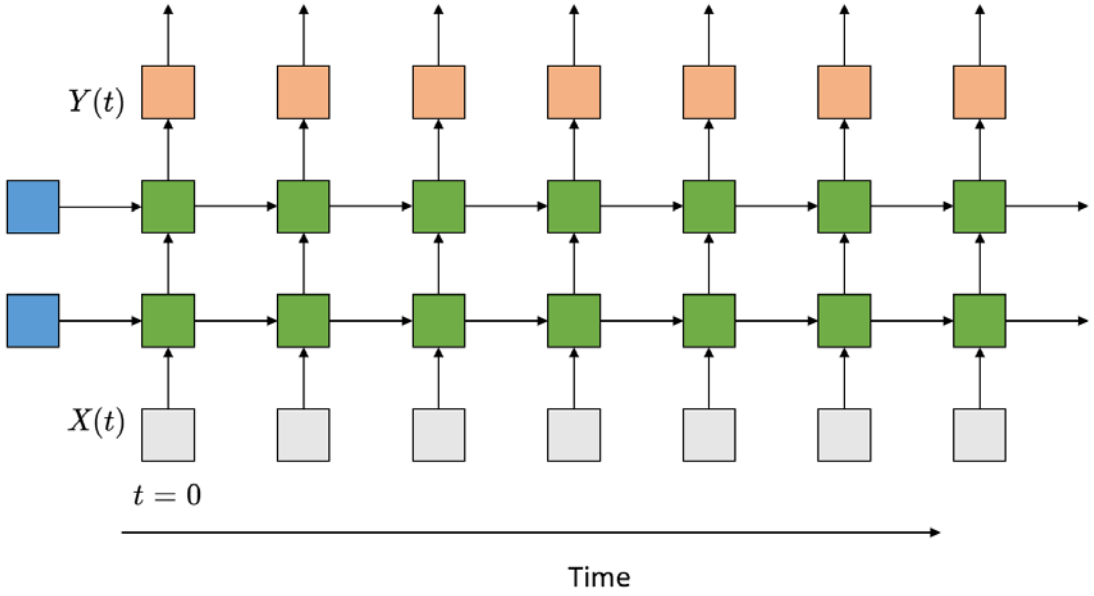

- The state (green) at any time is determined by the input at that time, and the state at the previous time
- All columns are identical
- An input at $t = 0$ affects outputs forever
- Also known as a recurrent neural net

Recurrent Neural Network

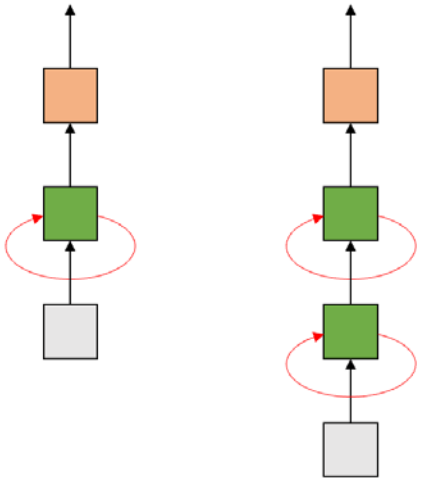

- Simplified models often drawn
- The loops imply recurrence

The Detailed Version of RNN

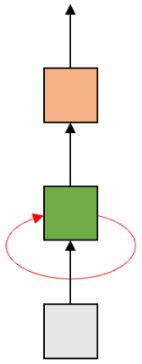

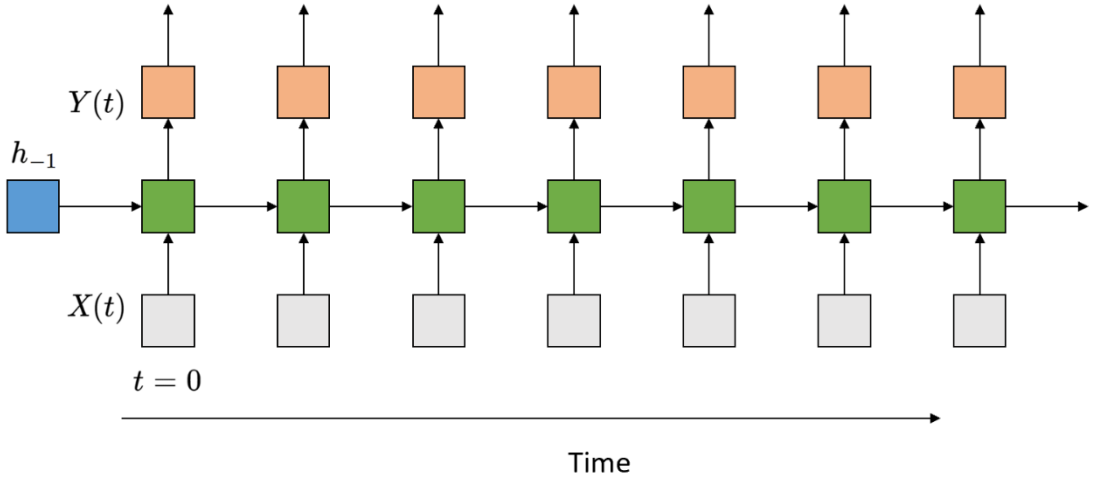

The Detailed Version of RNN

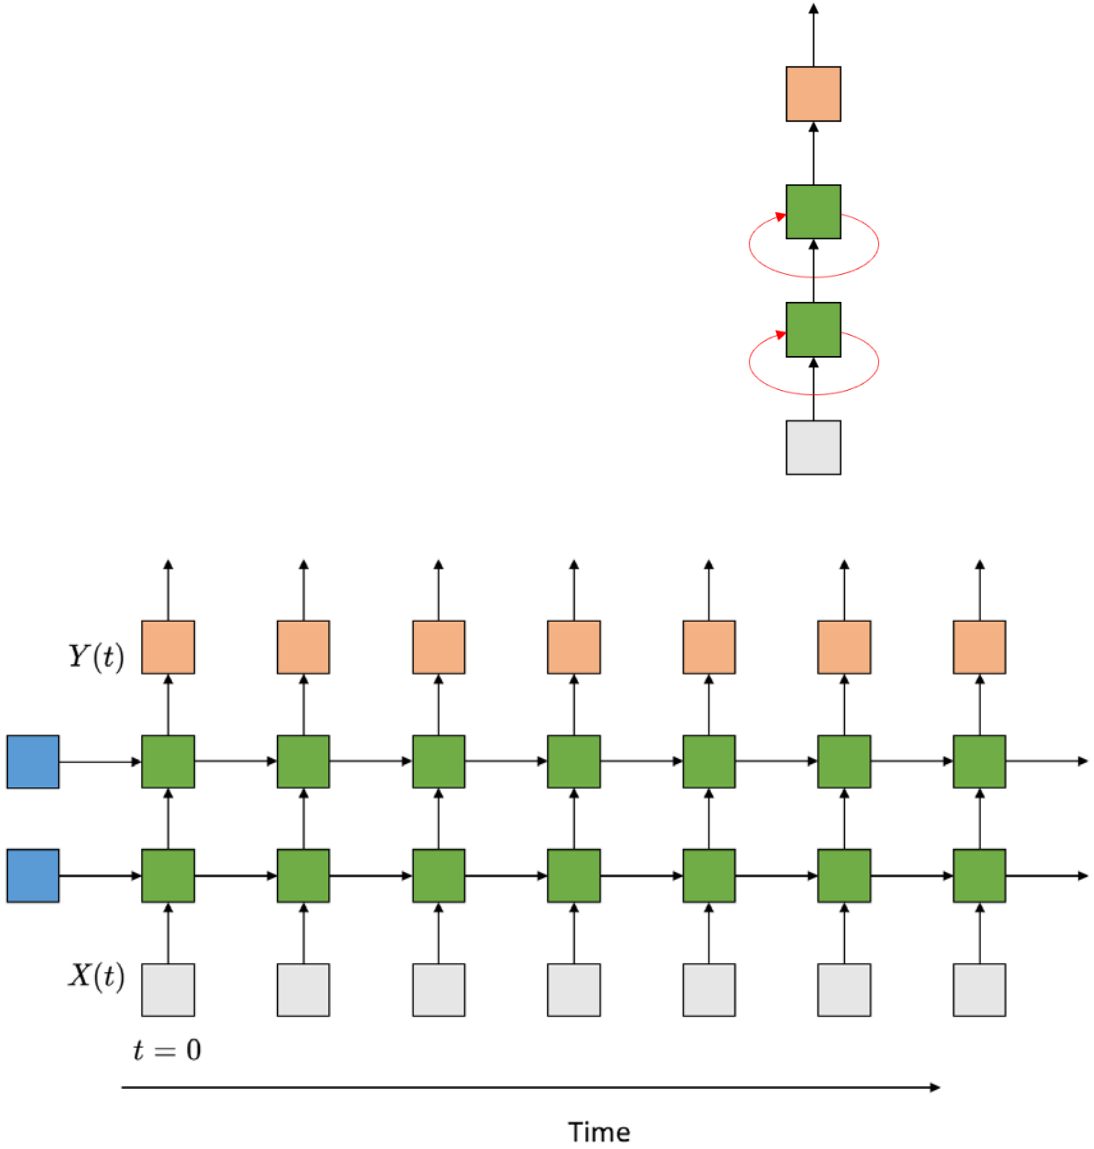

RNN Applications

- Machine translation
- Speech recognition
- Text-to-speech
- Image captioning
- Video analysis/understanding

“Vanilla” Neural Network

one to one

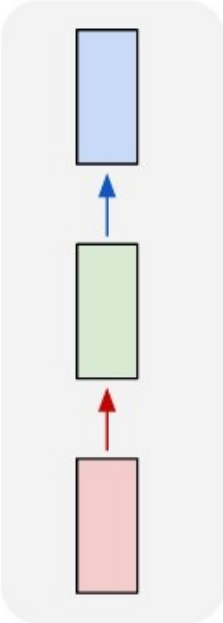

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Recurrent Neural Network: Process Sequences

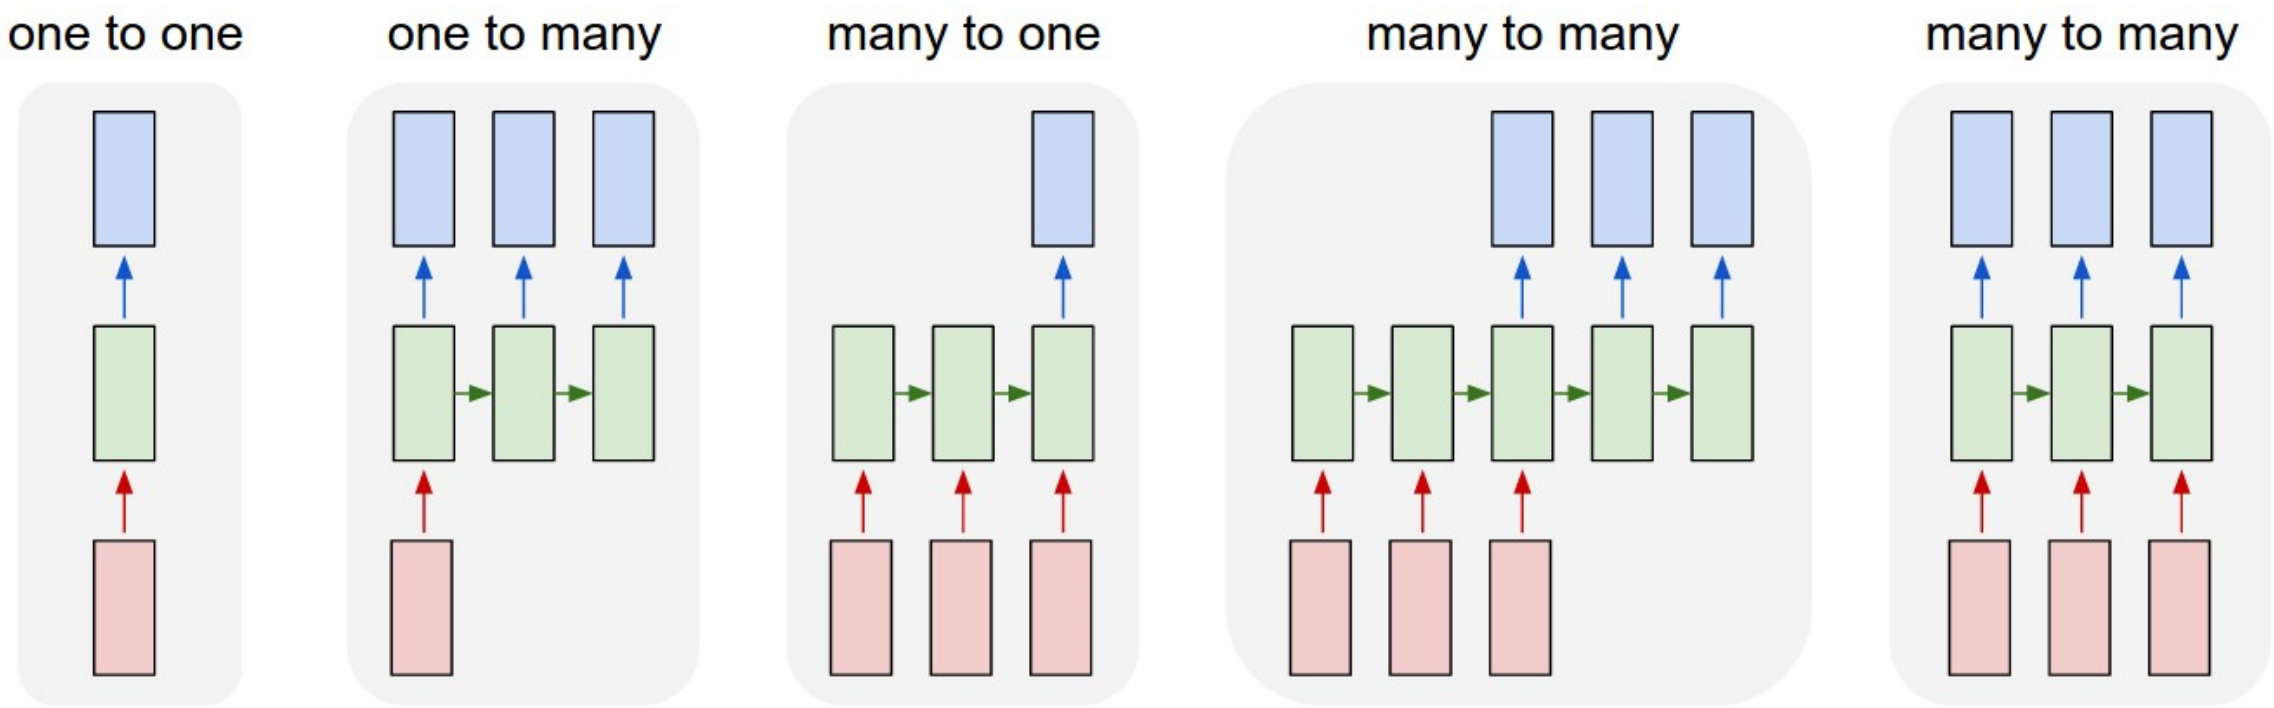

e.g. Image Captioning
image → sequence of words

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Recurrent Neural Network: Process Sequences

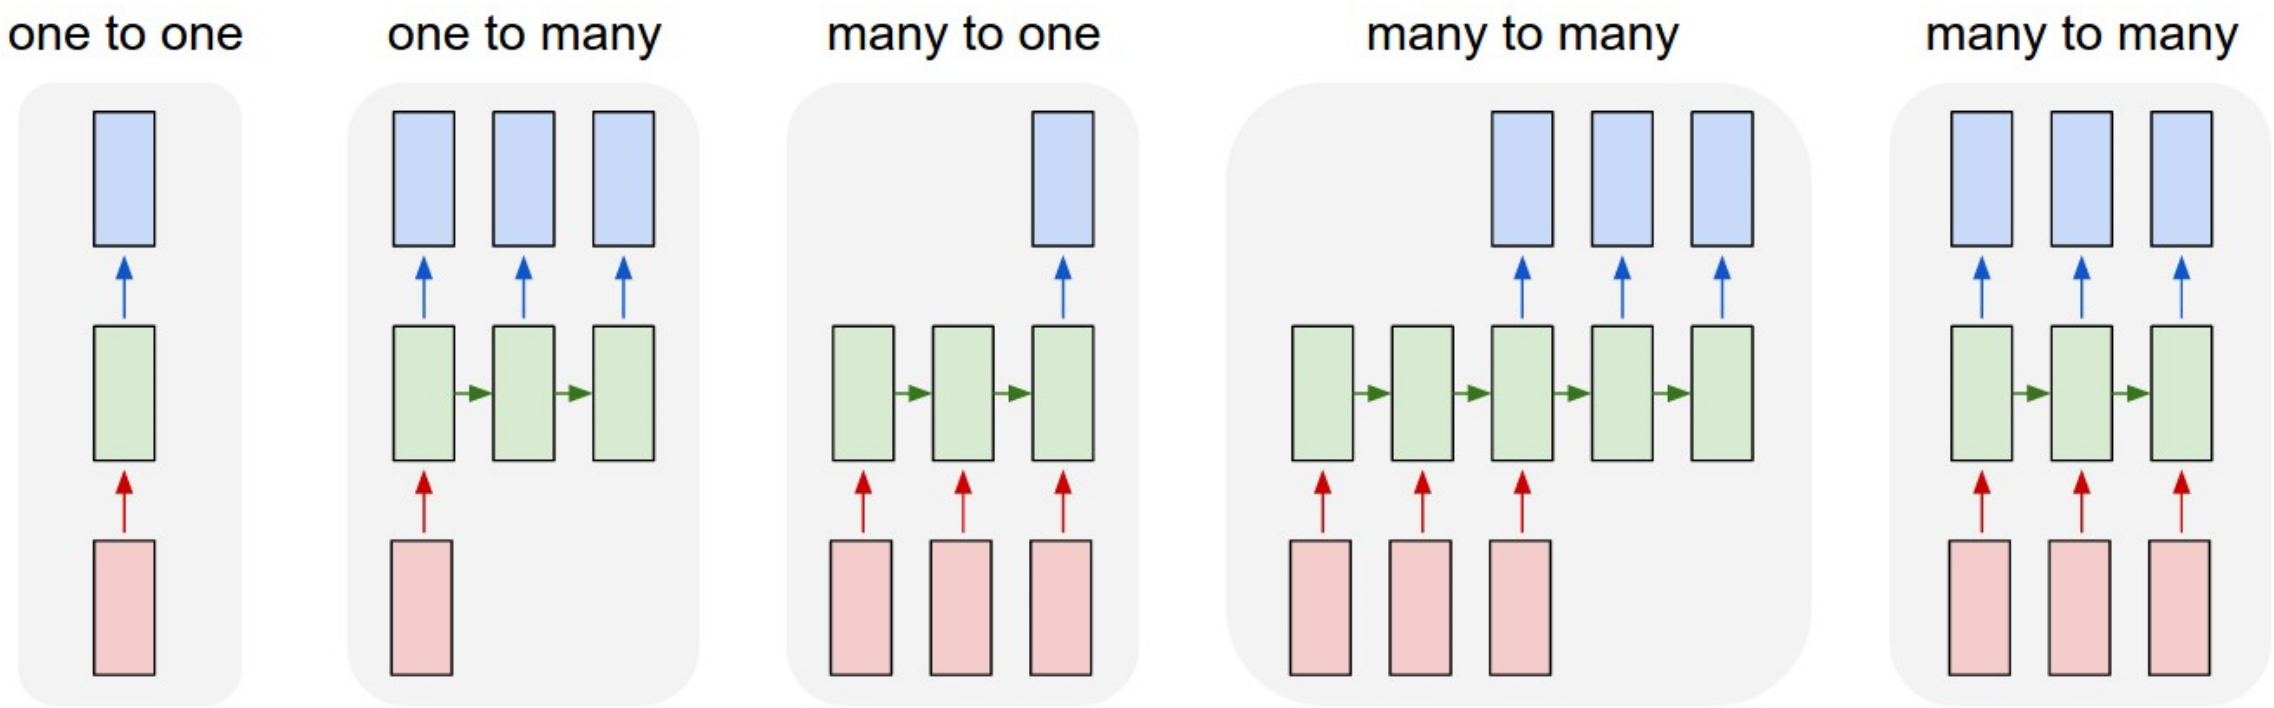

e.g. Sentiment Classification
sequence of words → sentiment

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Recurrent Neural Network: Process Sequences

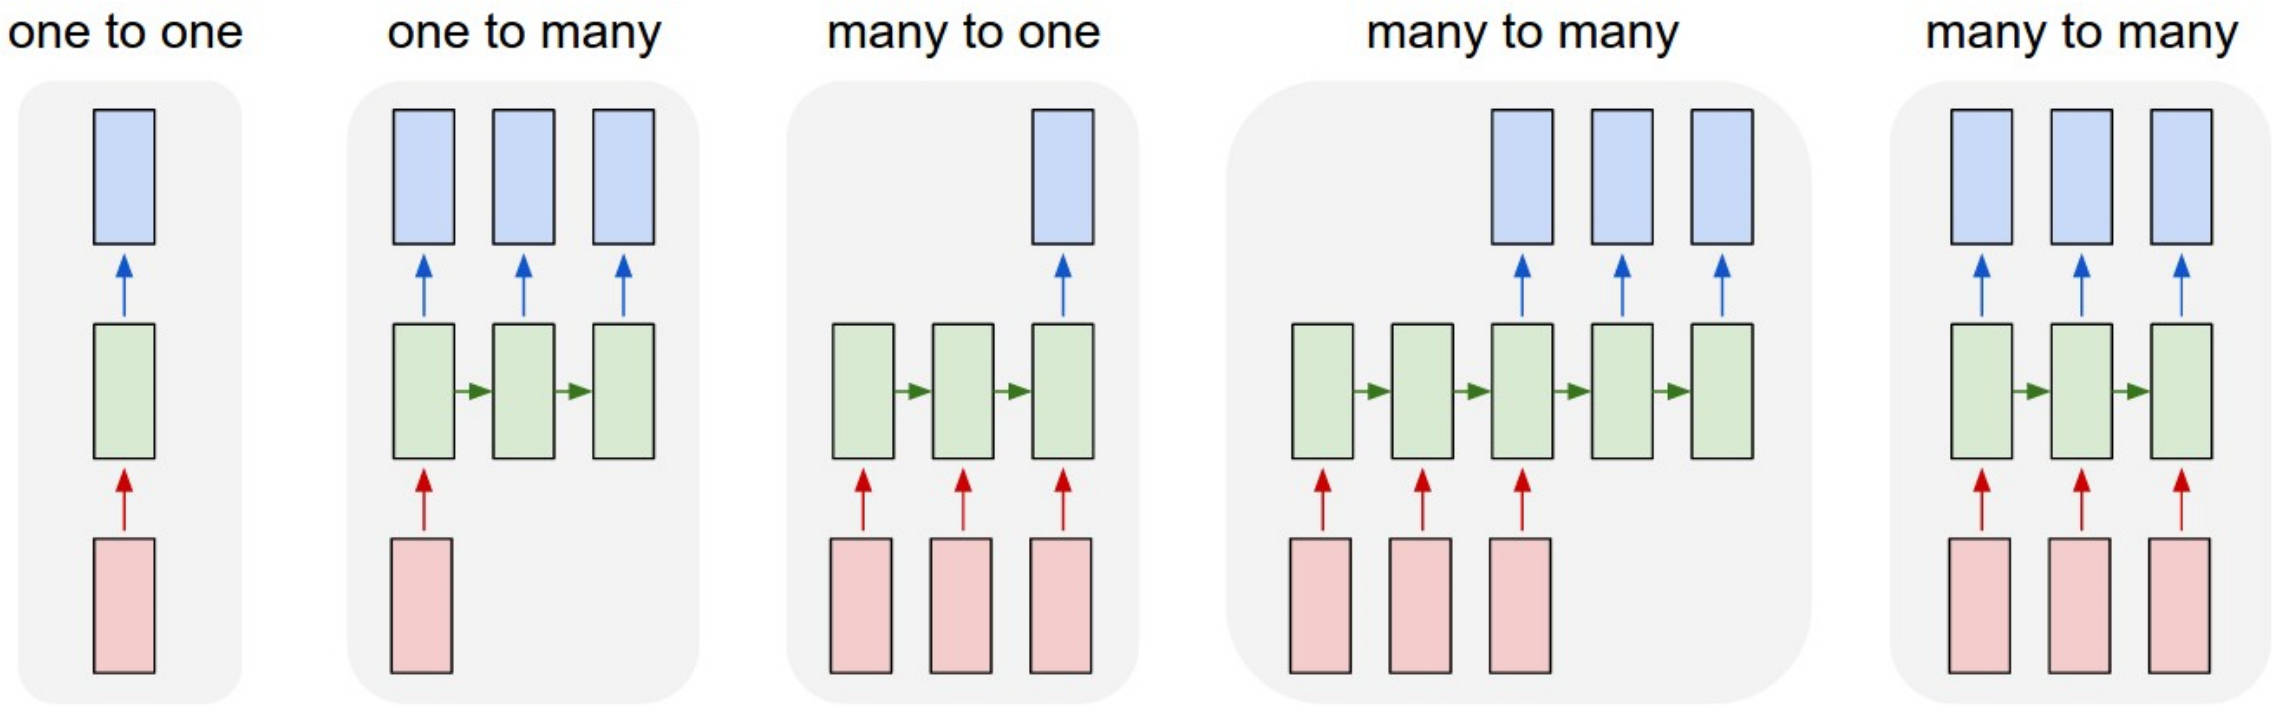

e.g. Machine Translation
Seq. of words → seq. of words

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Recurrent Neural Network: Process Sequences

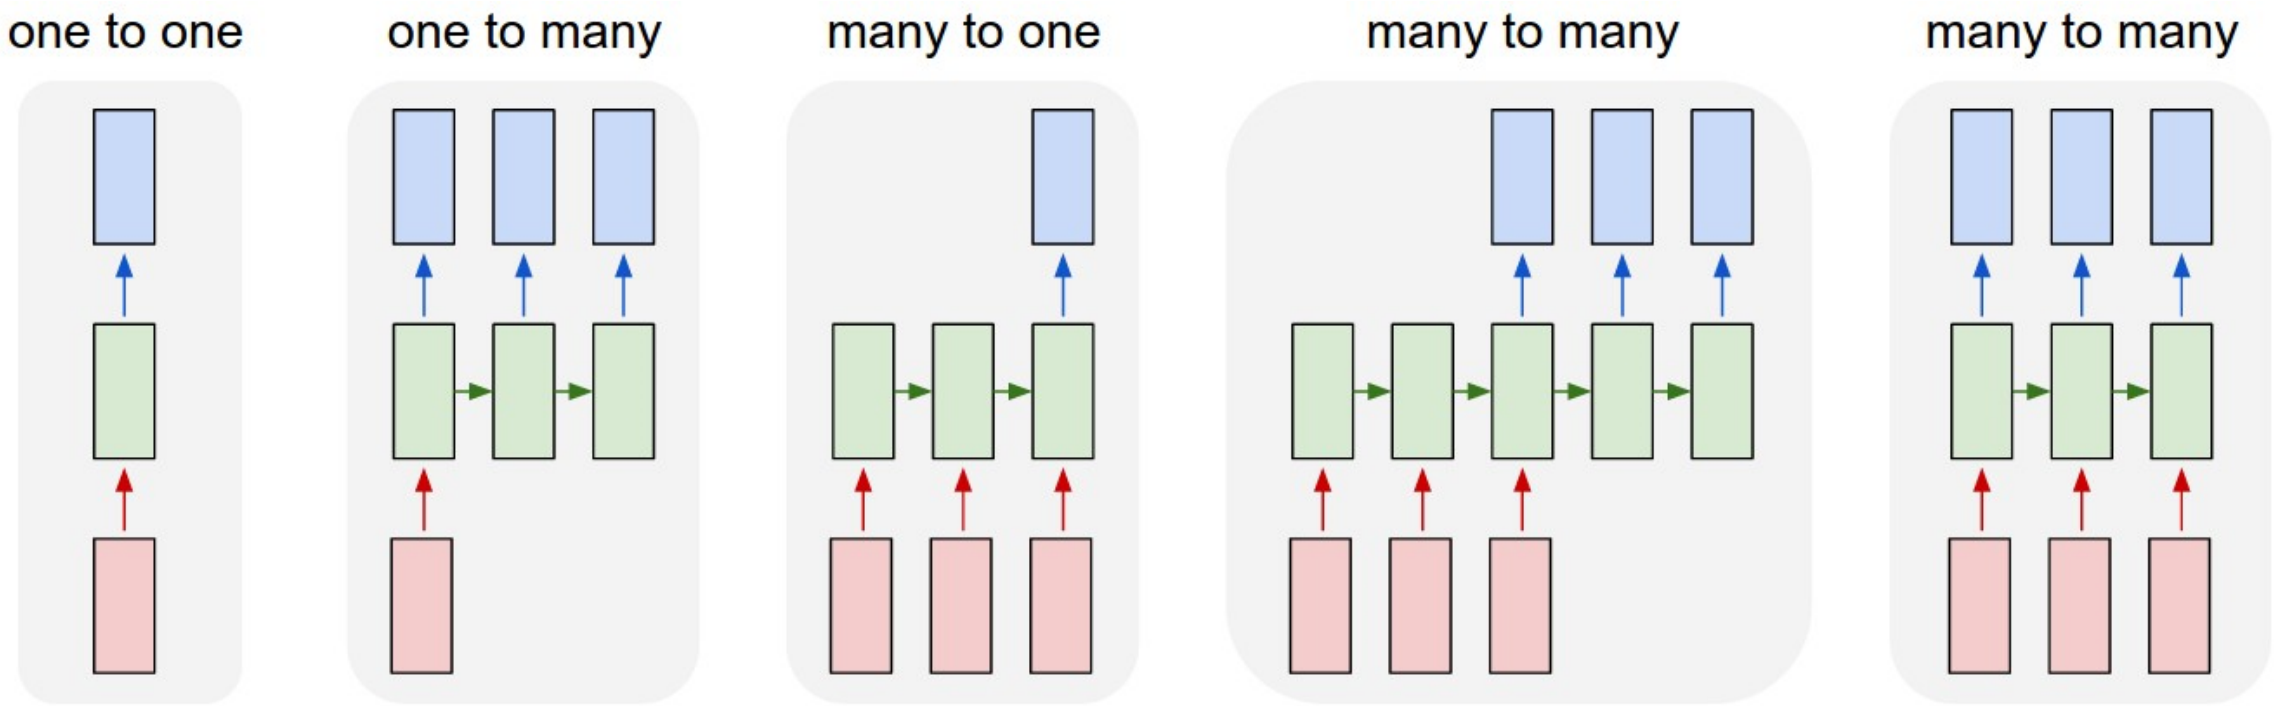

e.g. Video classification on frame level

http://karpathy.github.io/2015/05/21/rnn-effectiveness/

Recurrent Neural Networks

Unrolling the Recurrence

$$x_1 \qquad x_2 \qquad x_3 \qquad \cdots \qquad x_\tau$$

Unrolling the Recurrence

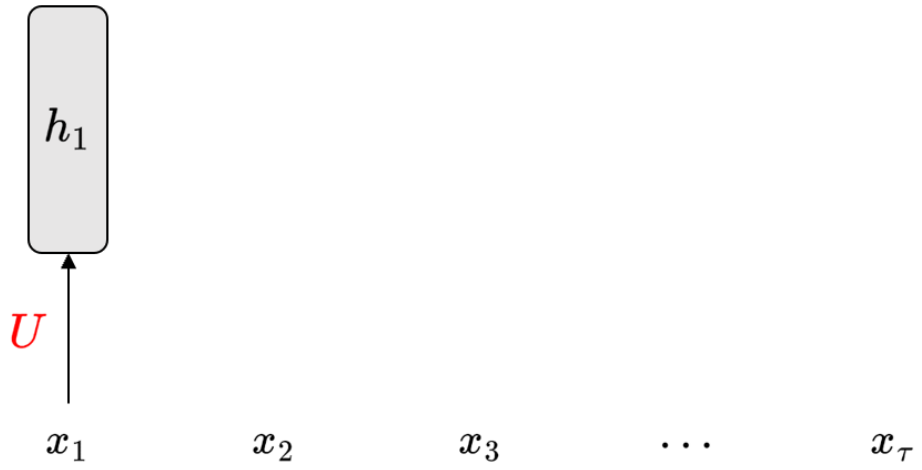

Unrolling the Recurrence

Unrolling the Recurrence

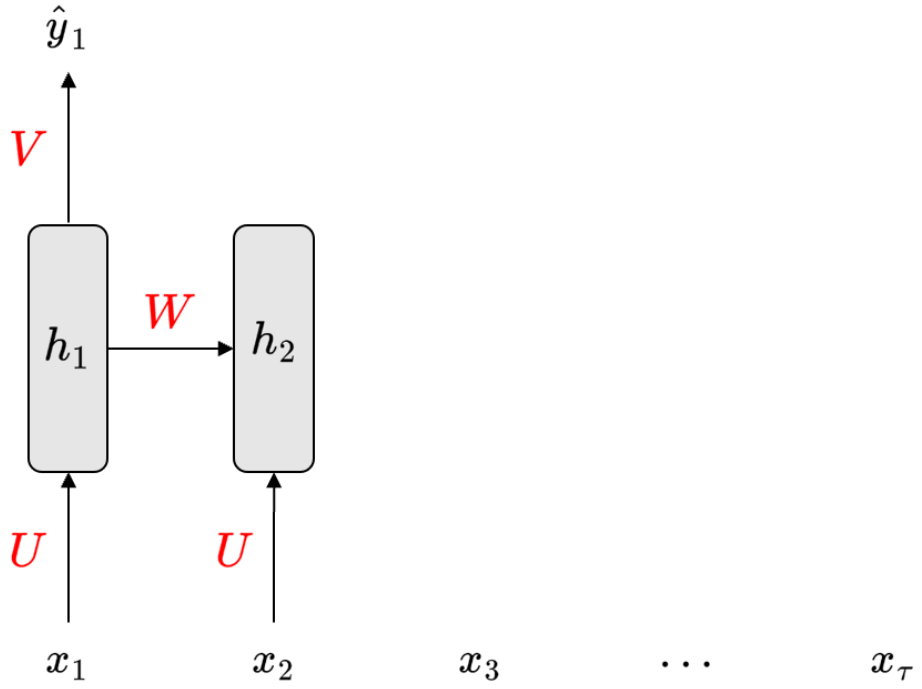

Unrolling the Recurrence

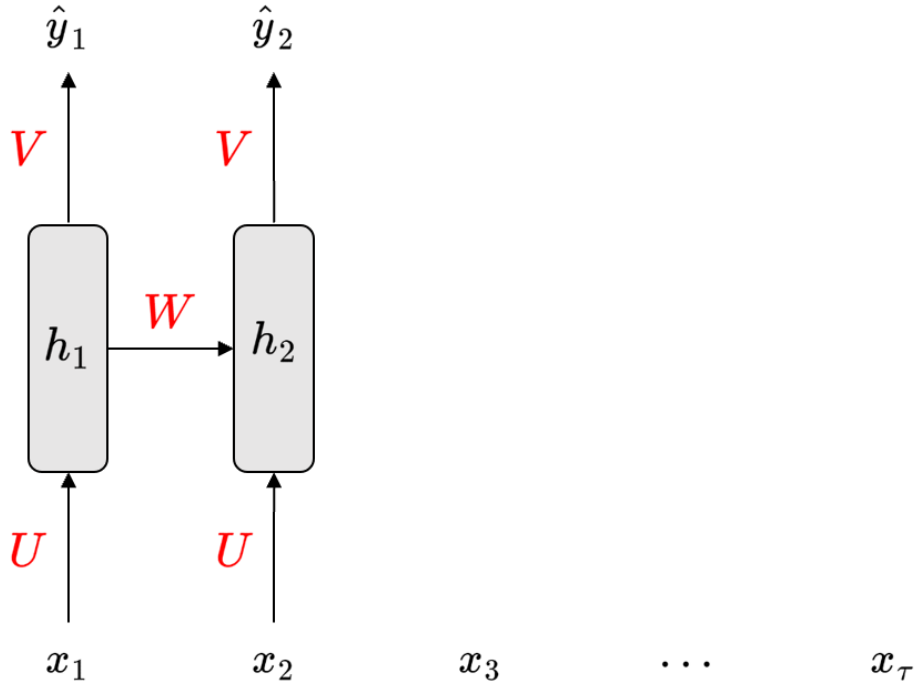

Unrolling the Recurrence

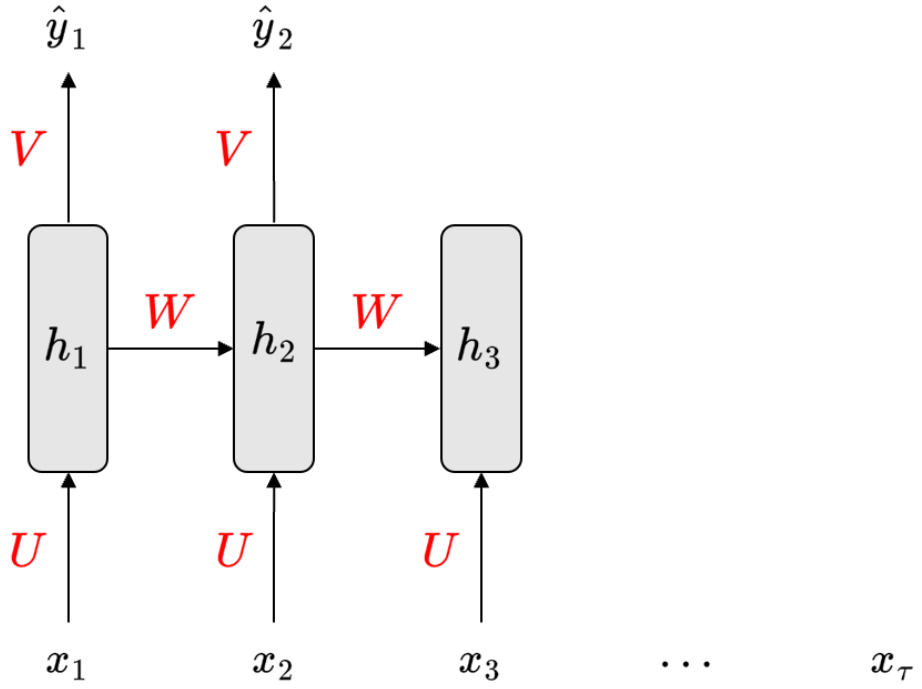

Unrolling the Recurrence

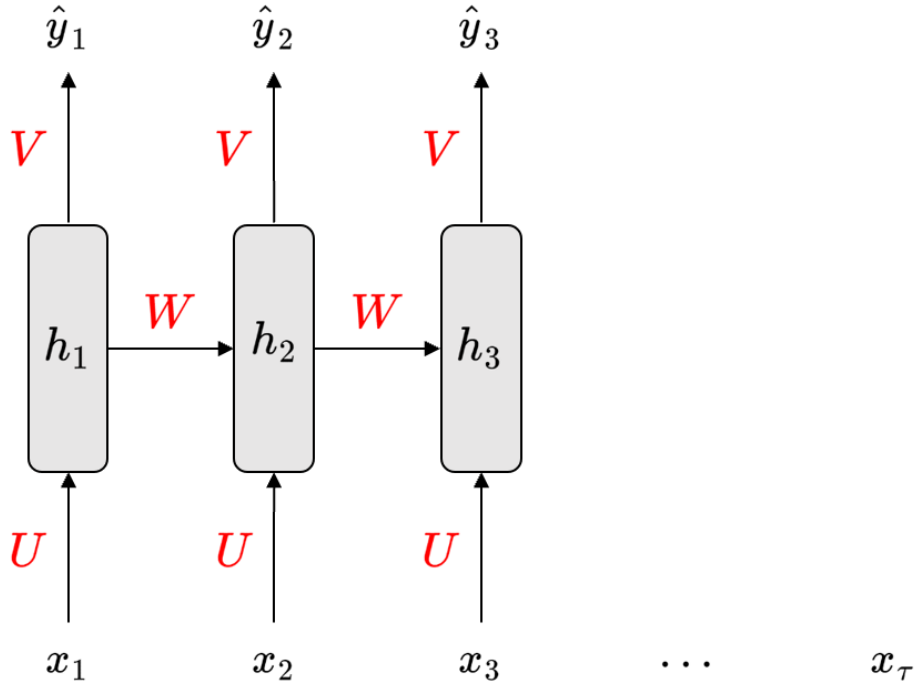

Unrolling the Recurrence

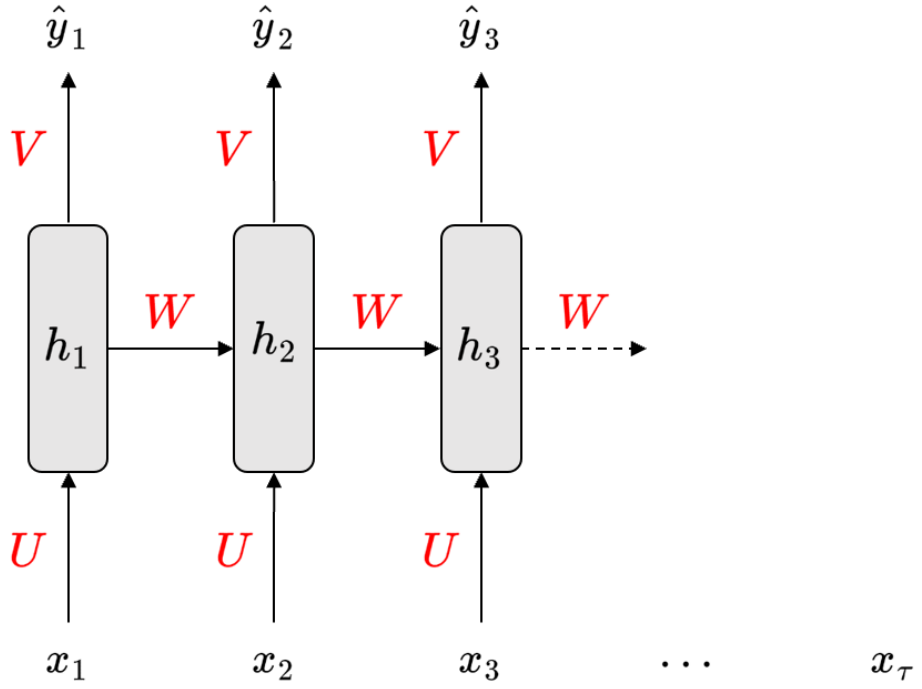

Unrolling the Recurrence

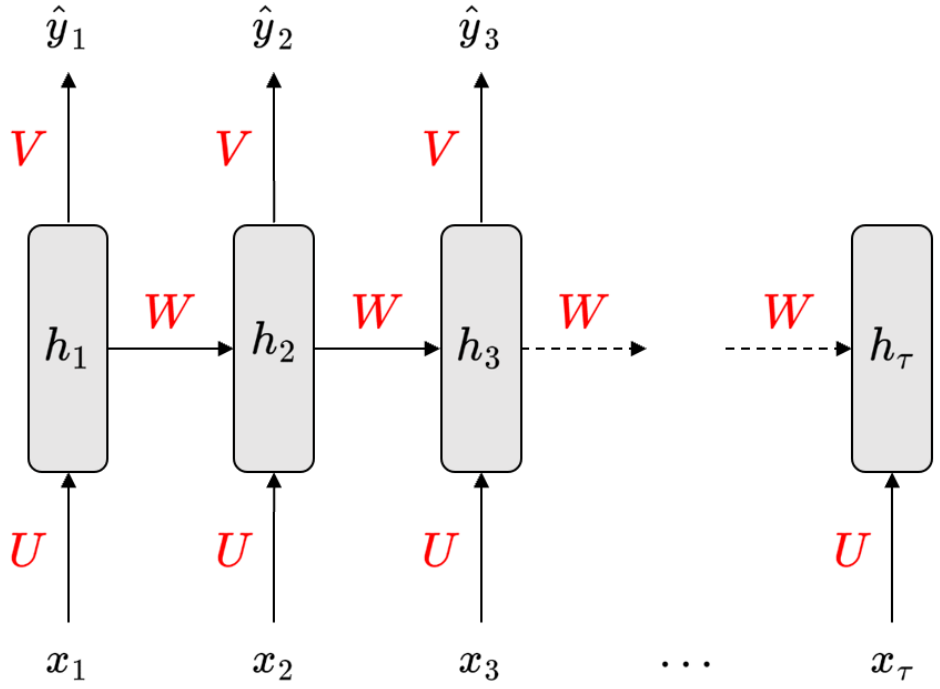

Unrolling the Recurrence

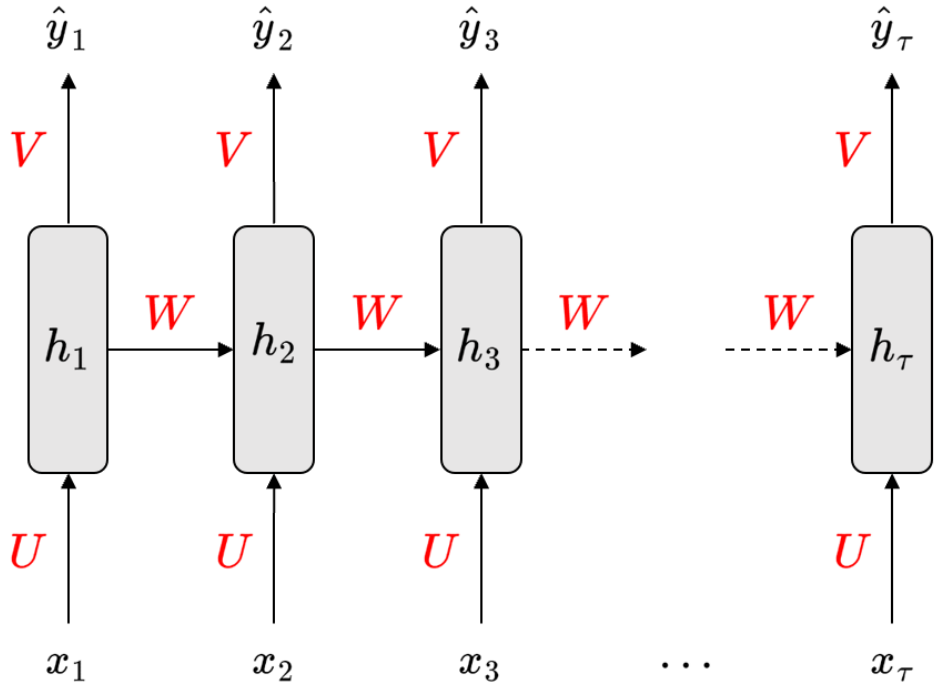

Recurrent Connections

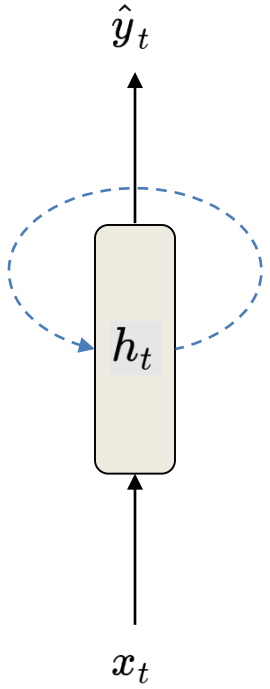

Recurrent Connections

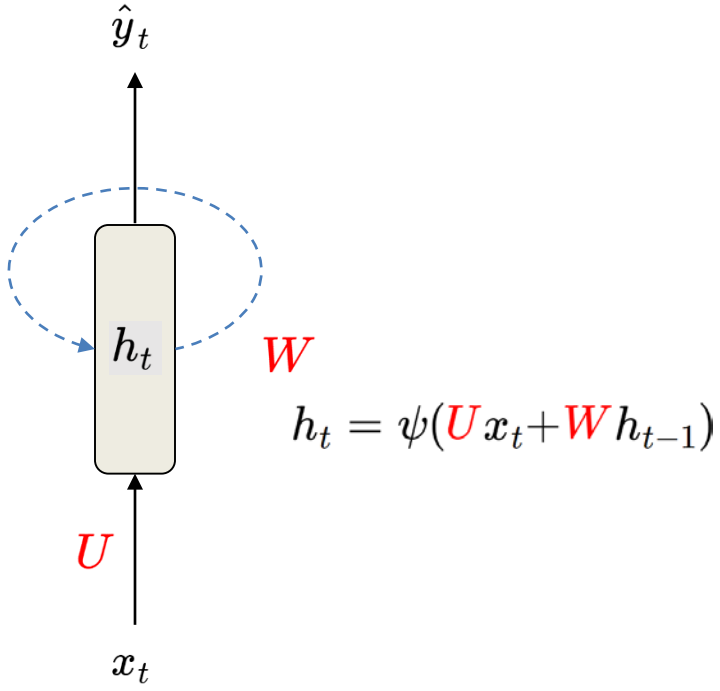

Feedforward Propagation

- This is a RNN where the input and output sequences are of the same length
- Feedforward operation proceeds from left to right
- Update Equations:

$$\mathbf{a}_t = b + W\mathbf{h}_{t-1} + U\mathbf{x}_t$$
$$\mathbf{h}_t = \tanh \mathbf{a}_t$$
$$\mathbf{o}_t = c + V\mathbf{h}_t$$
$$\hat{\mathbf{y}}_t = \text{softmax}(\mathbf{o}_t)$$

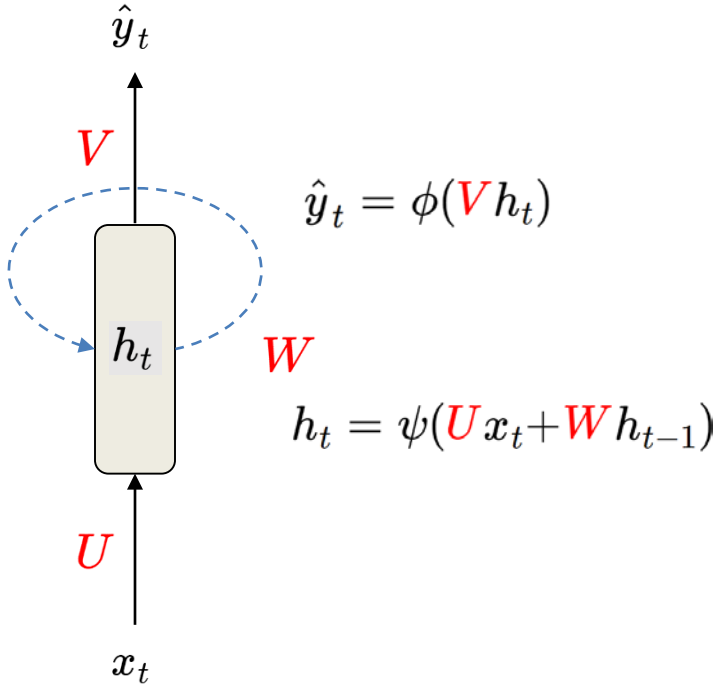

How to Train RNN

Backward Propagation

- Loss would just be the sum of losses over time steps
- Treat the recurrent network as a usual multilayer network and apply backpropagation on the unrolled network
- This is called Backpropagation through time (BPTT)

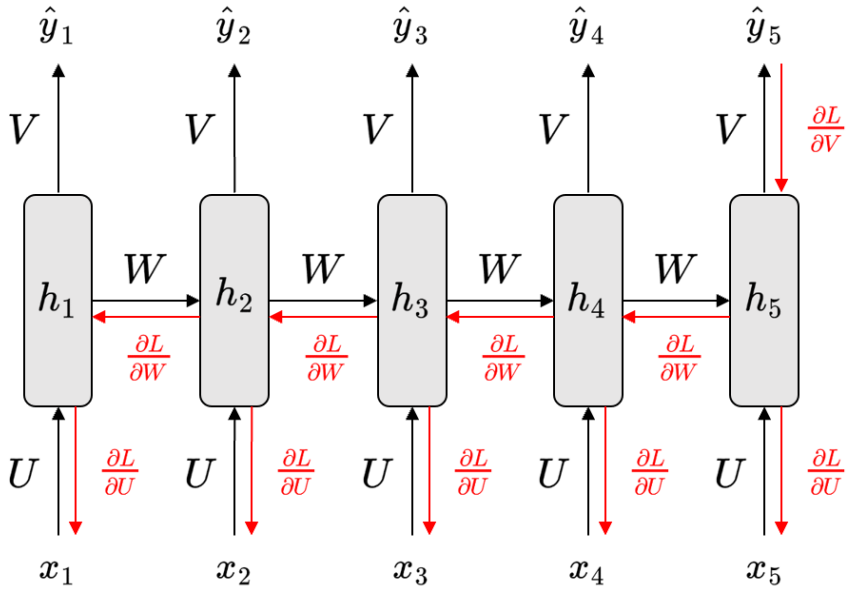

Long Short-Term Memory (LSTM)

Long Short-Term Memory (LSTM)

- Long-Term Dependencies
  - Gradients propagated over many stages tend to either **vanish** or **explode**
  - Difficulty with long-term dependencies arises from the exponentially smaller weights given to long-term interactions
  - Introduce a memory state that runs through only linear operators
  - Use gating units to control the updates of the state

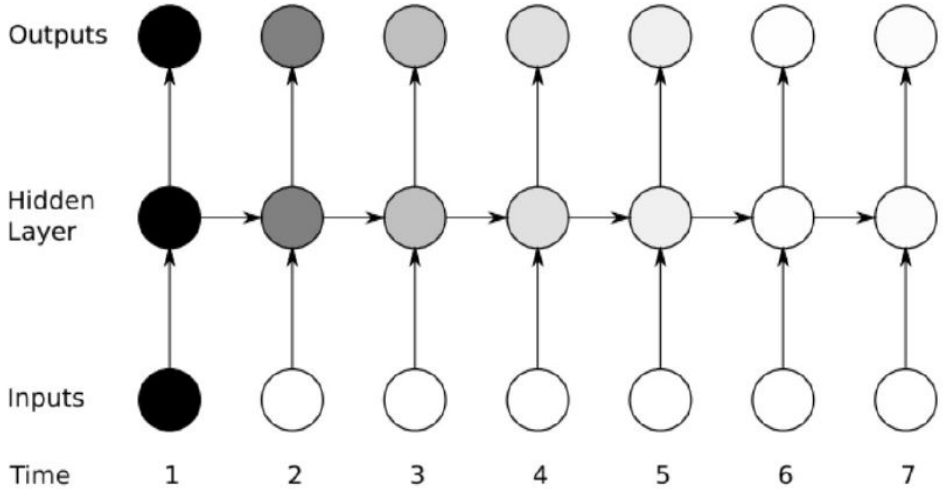

Example

- “I grew up in France… I speak fluent *French*.”

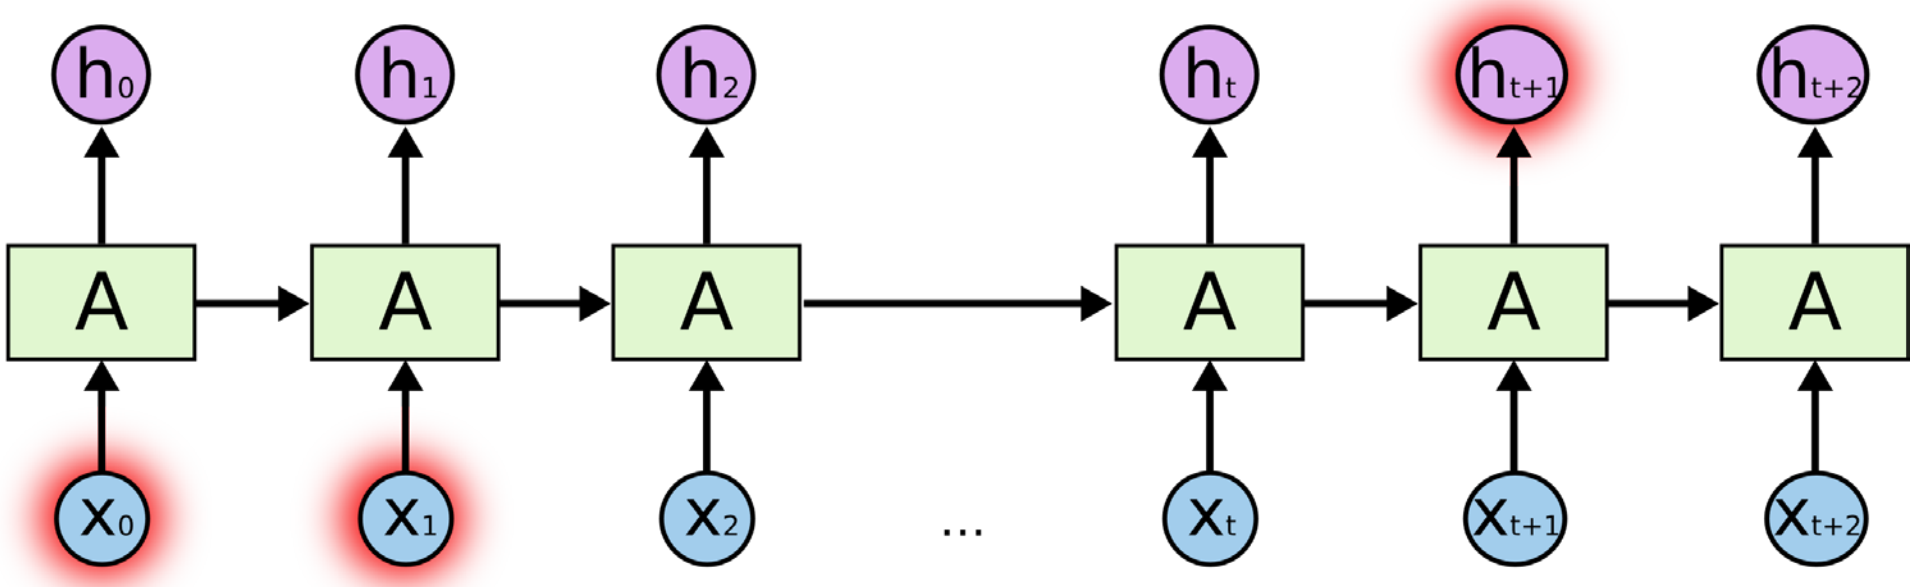

Long Short-Term Memory (LSTM)

- Consists of a memory cell and a set of gating units
  - Memory cell is the context that carries over
  - Forget gate controls erase operation
  - Input gate controls write operation
  - Output gate controls the read operation

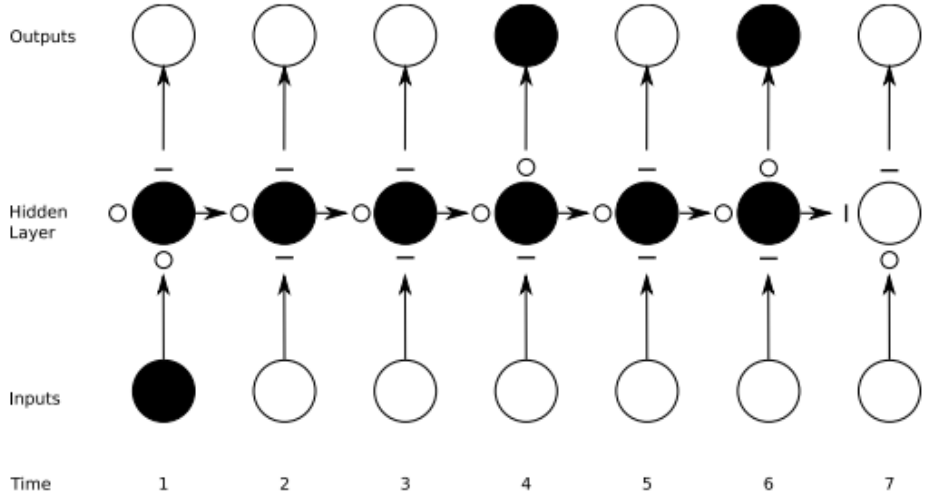

Long Short-Term Memory (LSTM)

- Consists of a memory cell and a set of gating units
  - Memory cell is the context that carries over
  - Forget gate controls erase operation
  - Input gate controls write operation
  - Output gate controls the read operation

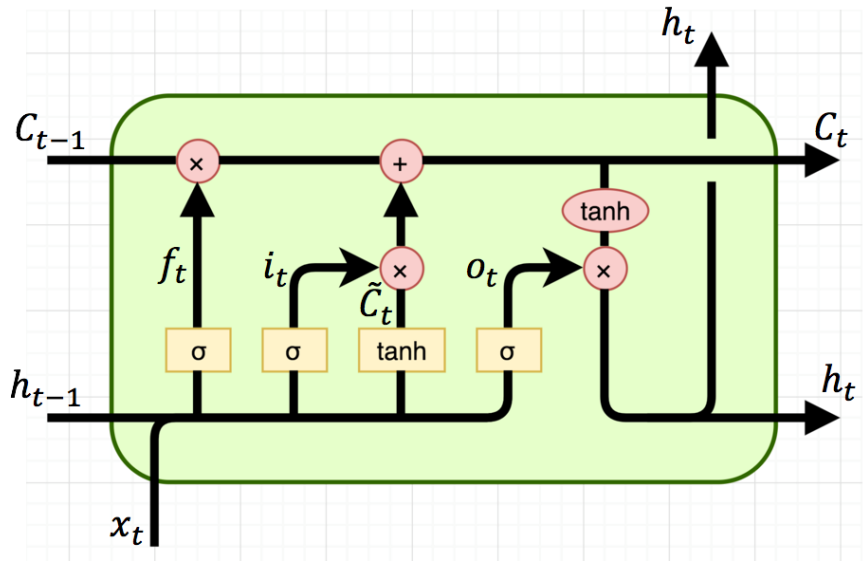

$$i_t = \sigma\left(x_t U^i + h_{t-1} W^i\right)$$
$$f_t = \sigma\left(x_t U^f + h_{t-1} W^f\right)$$
$$o_t = \sigma\left(x_t U^o + h_{t-1} W^o\right)$$
$$\tilde{C}_t = \tanh\left(x_t U^g + h_{t-1} W^g\right)$$
$$C_t = \sigma\left(f_t * C_{t-1} + i_t * \tilde{C}_t\right)$$
$$h_t = \tanh(C_t) * o_t$$

Element-by-Element

Element-by-Element Addition

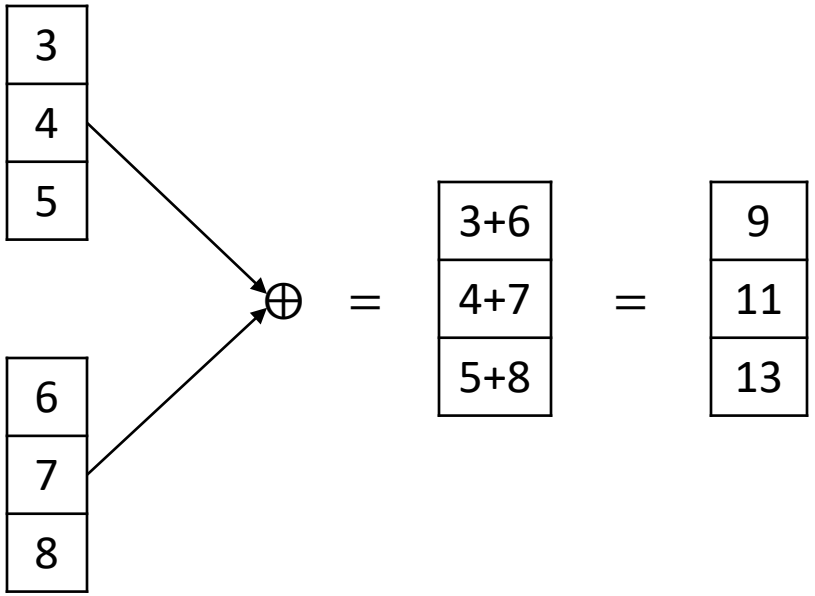

Element-by-Element Addition

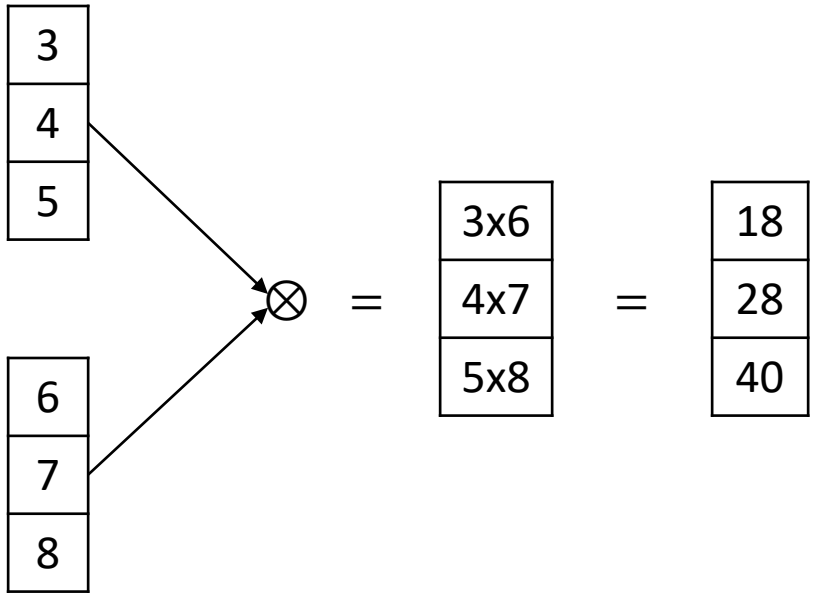

Gating

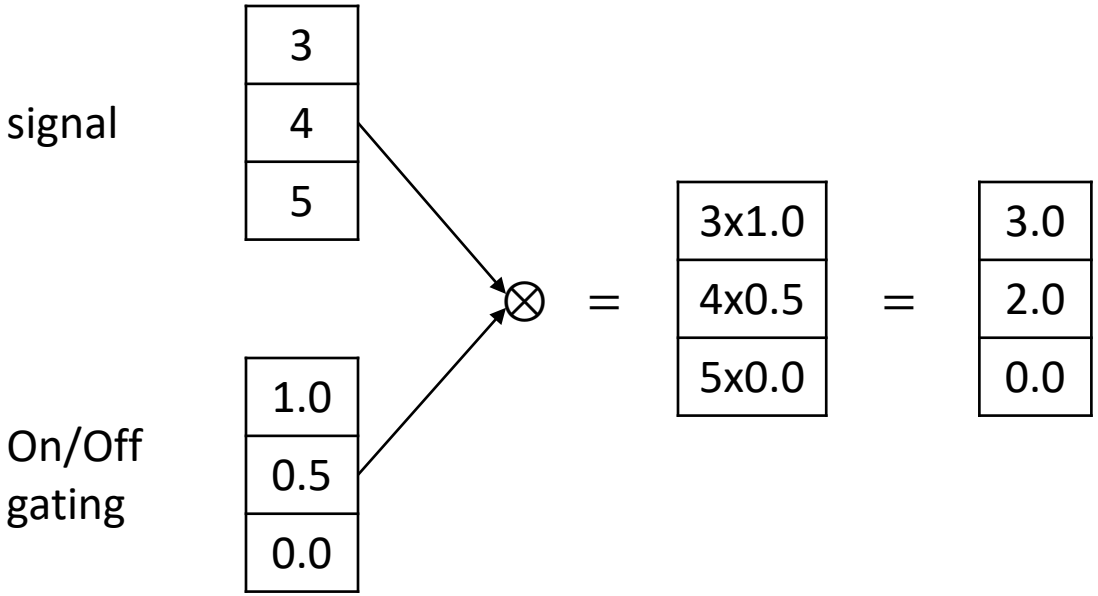

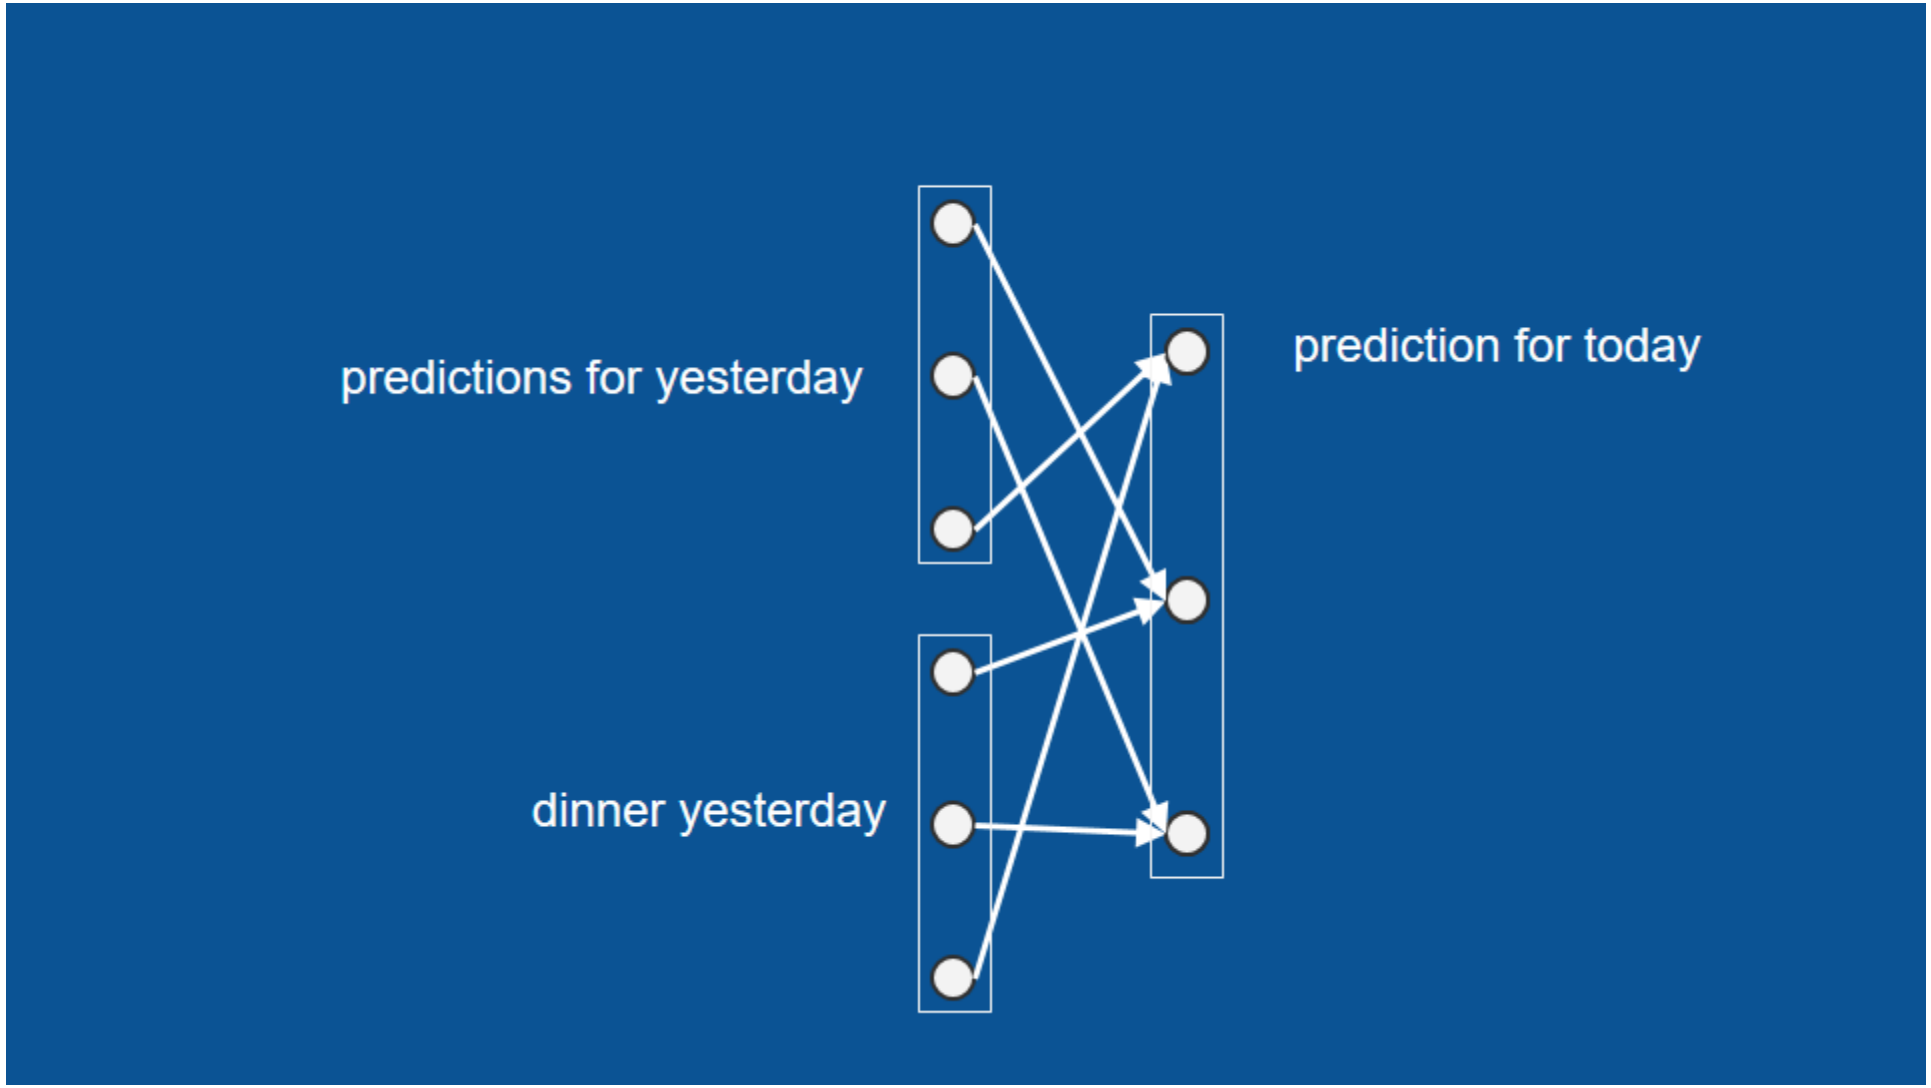

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

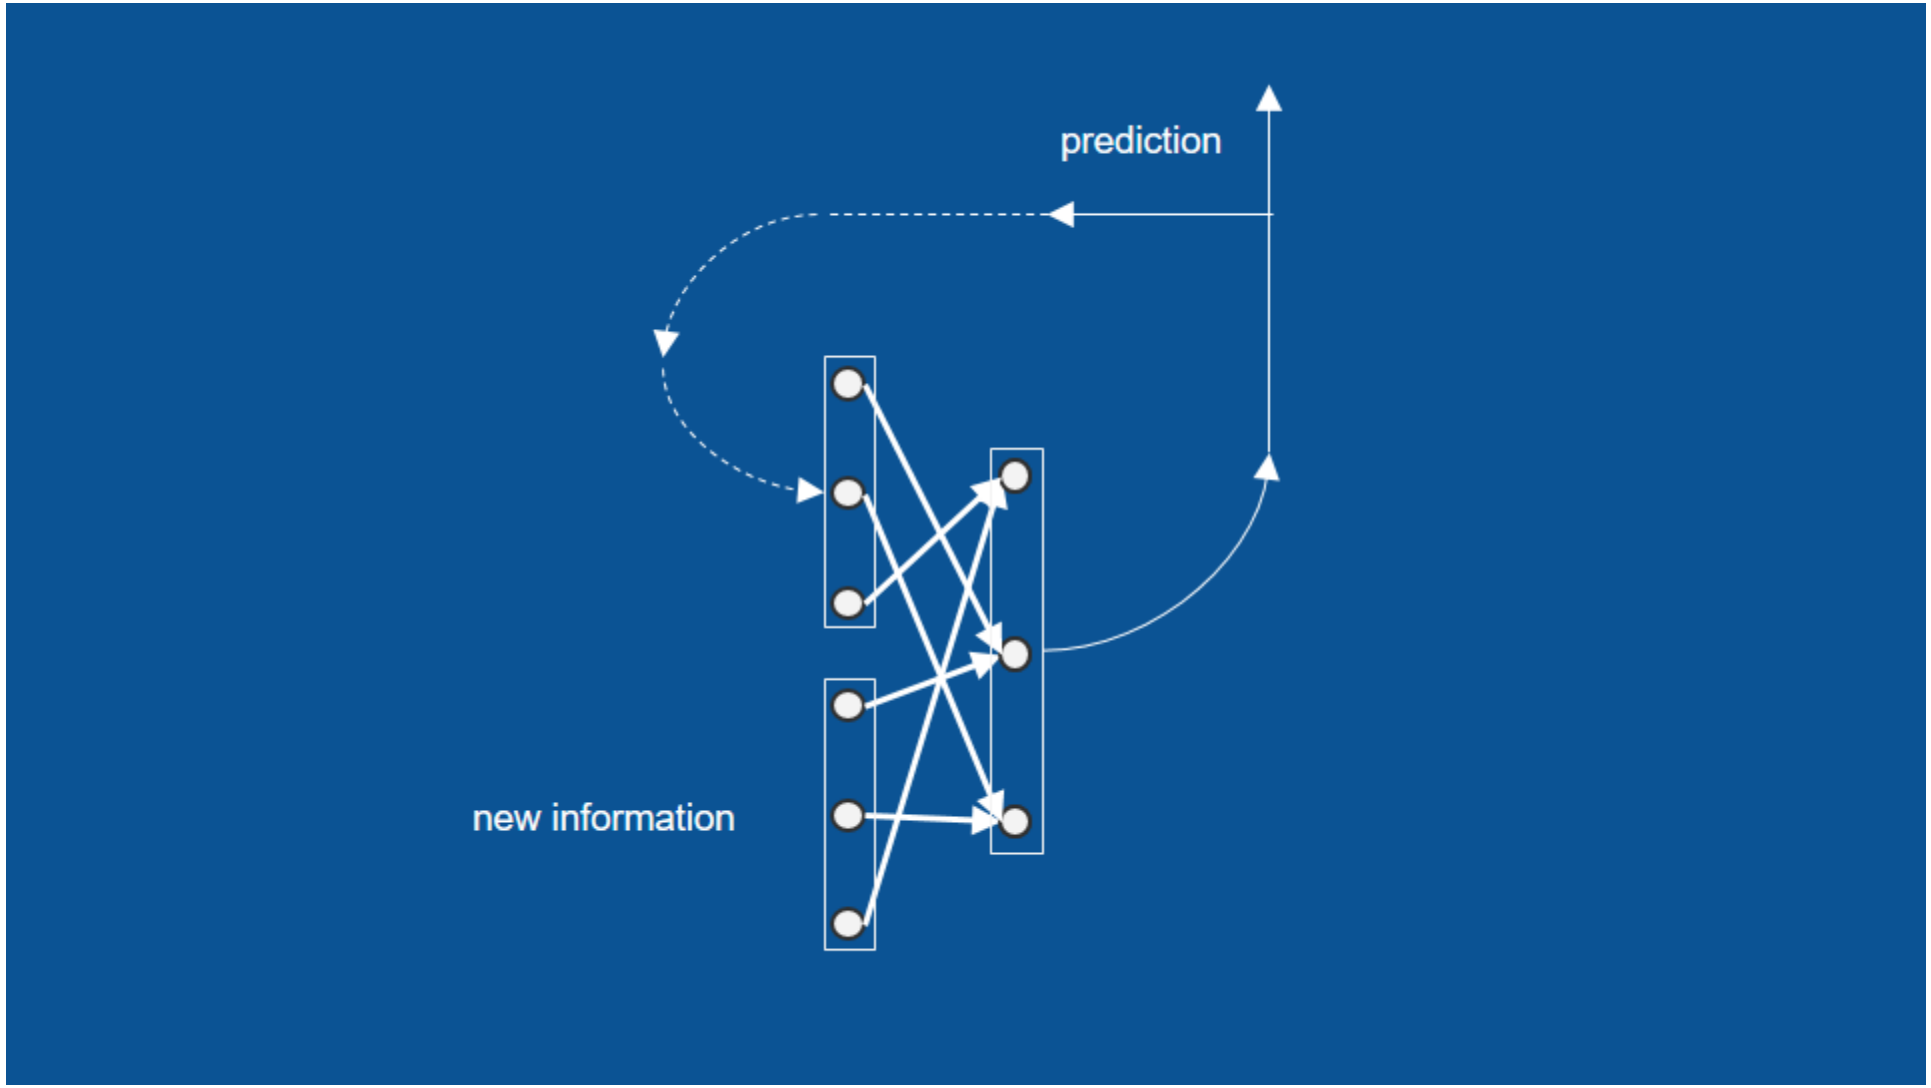

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

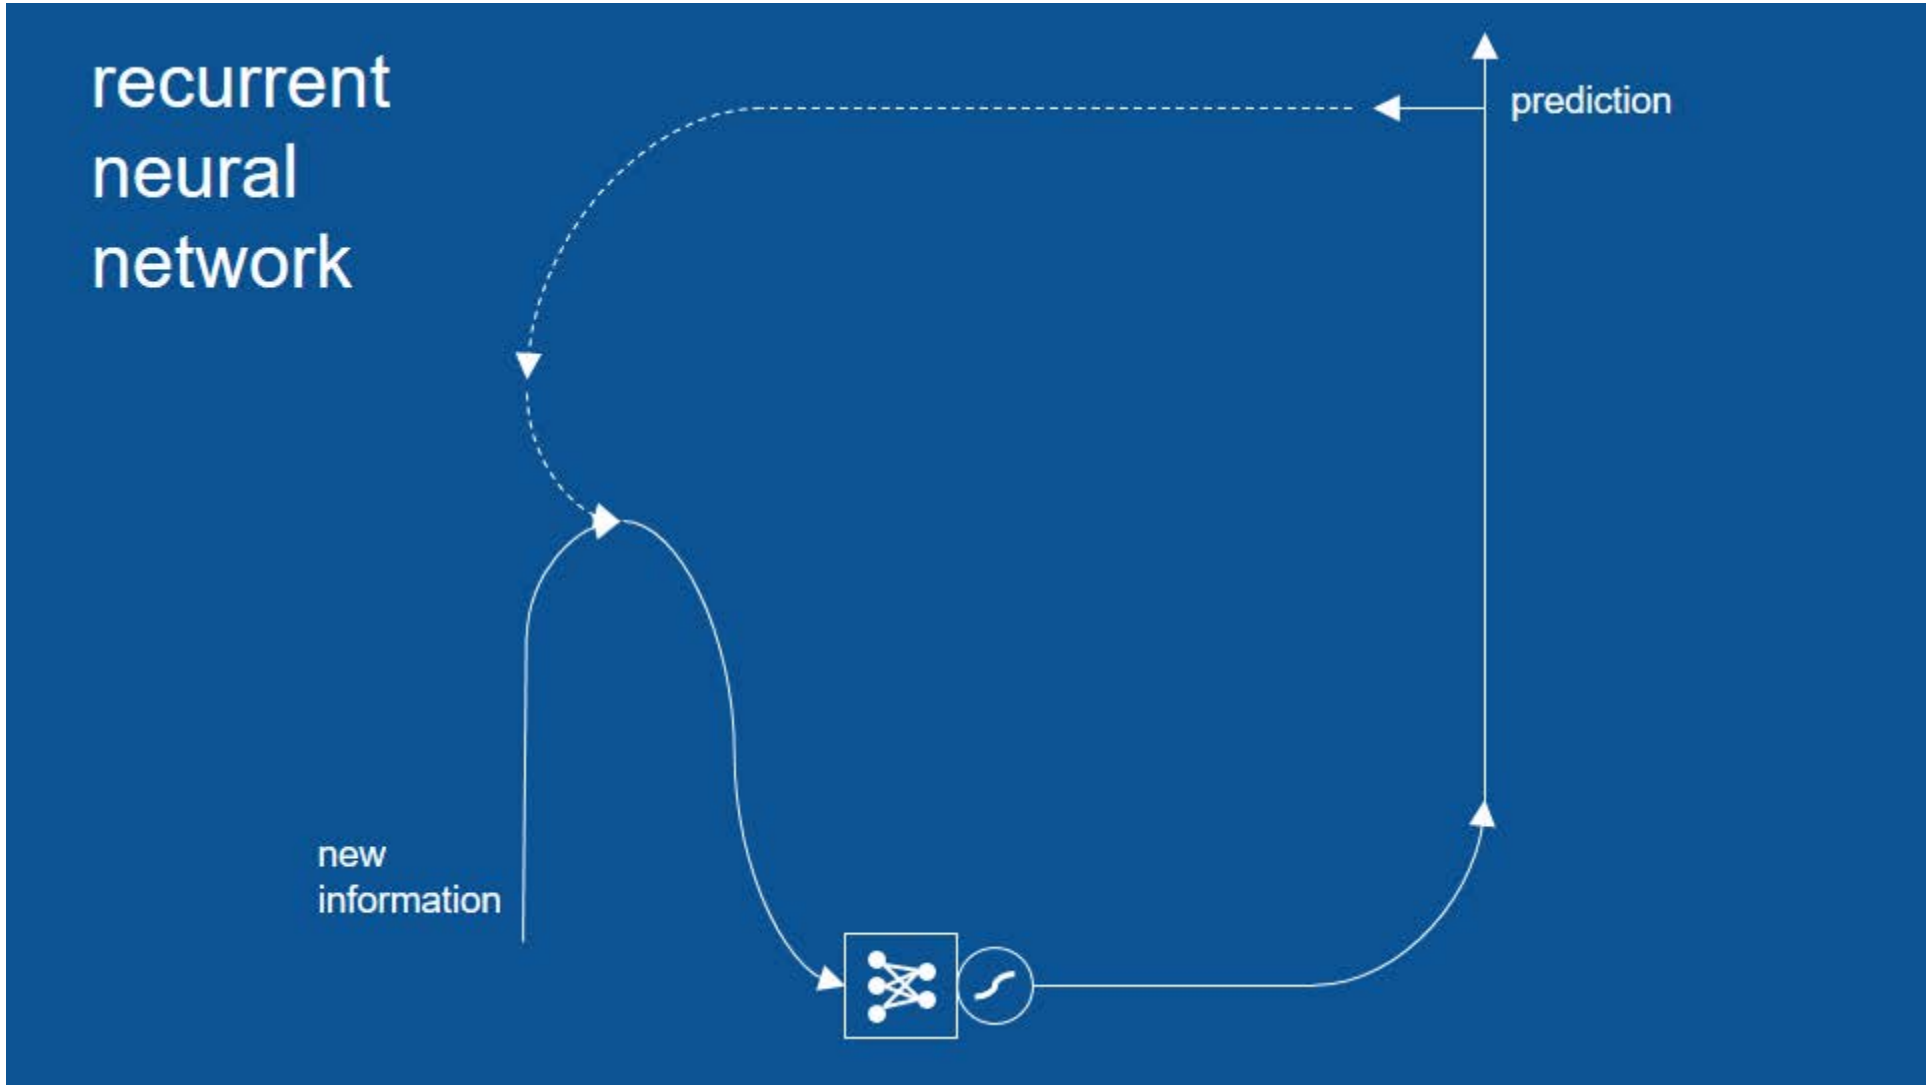

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

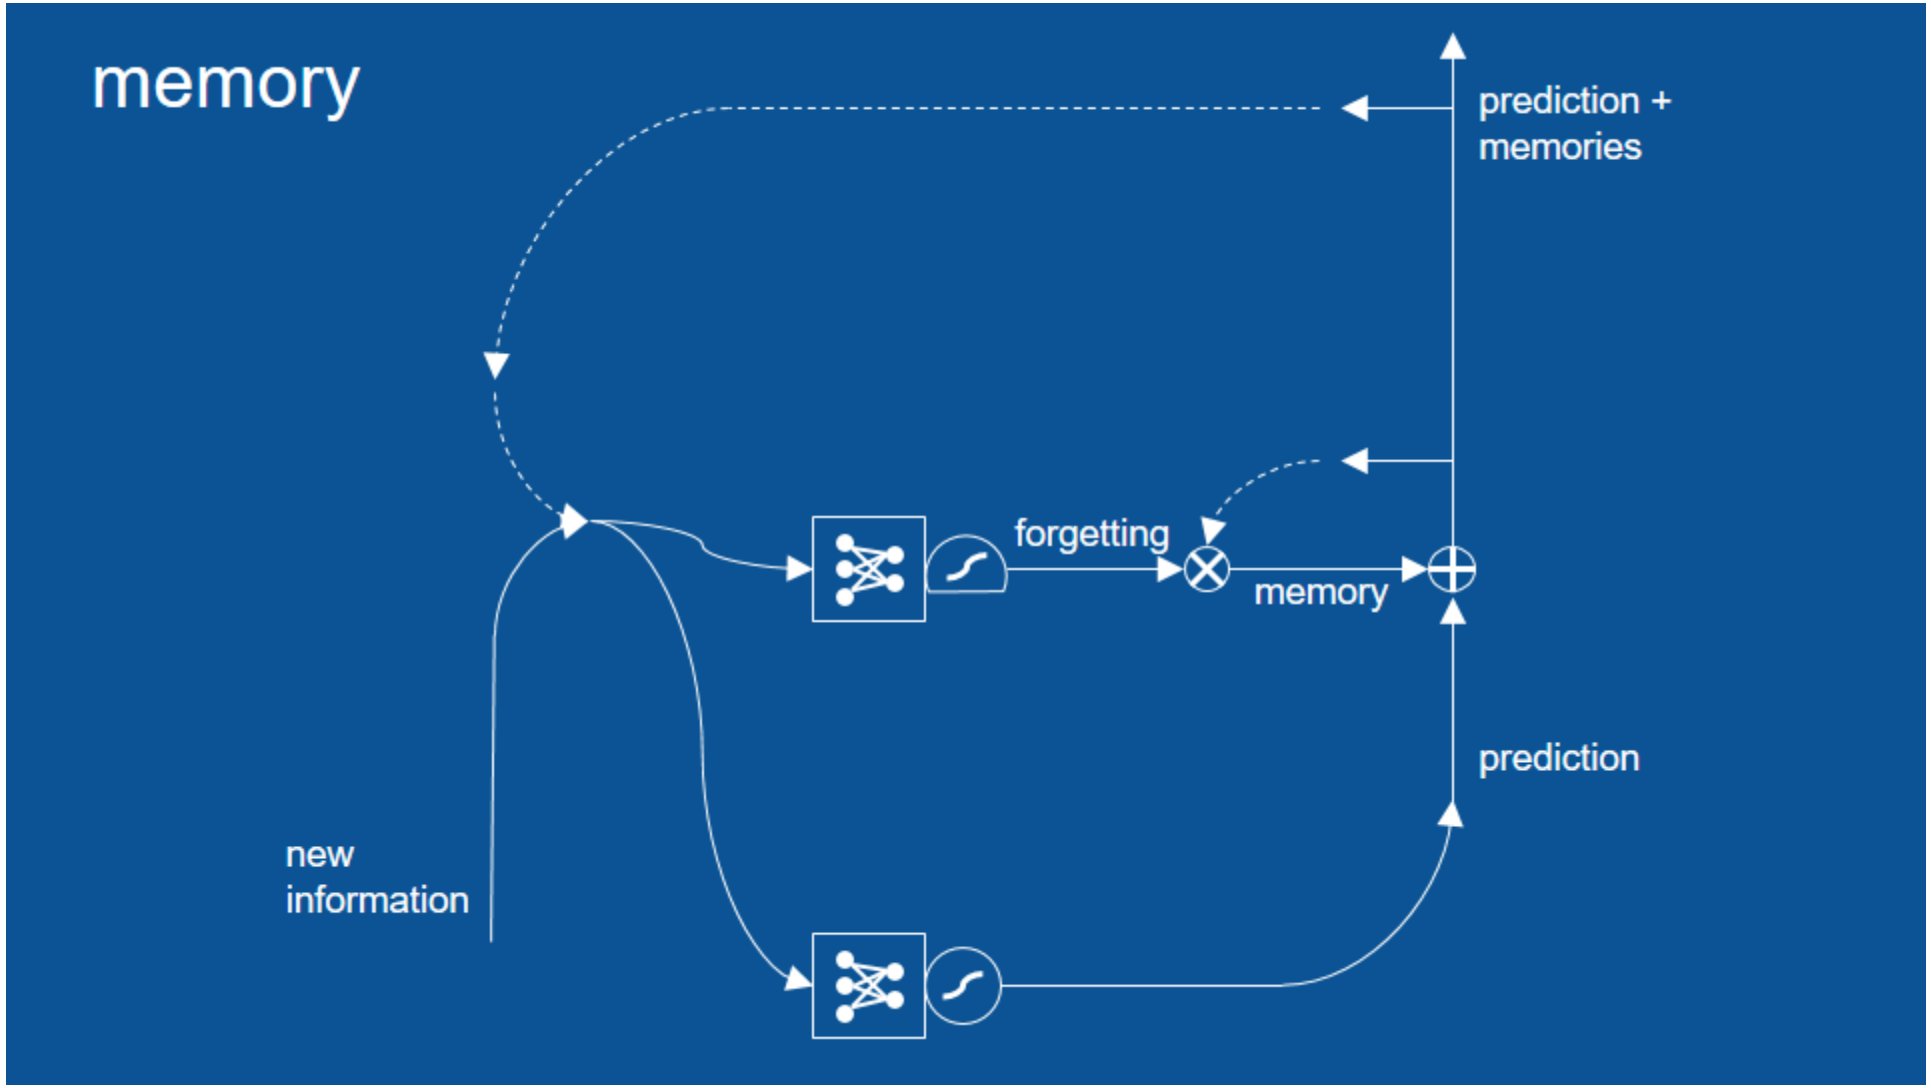

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

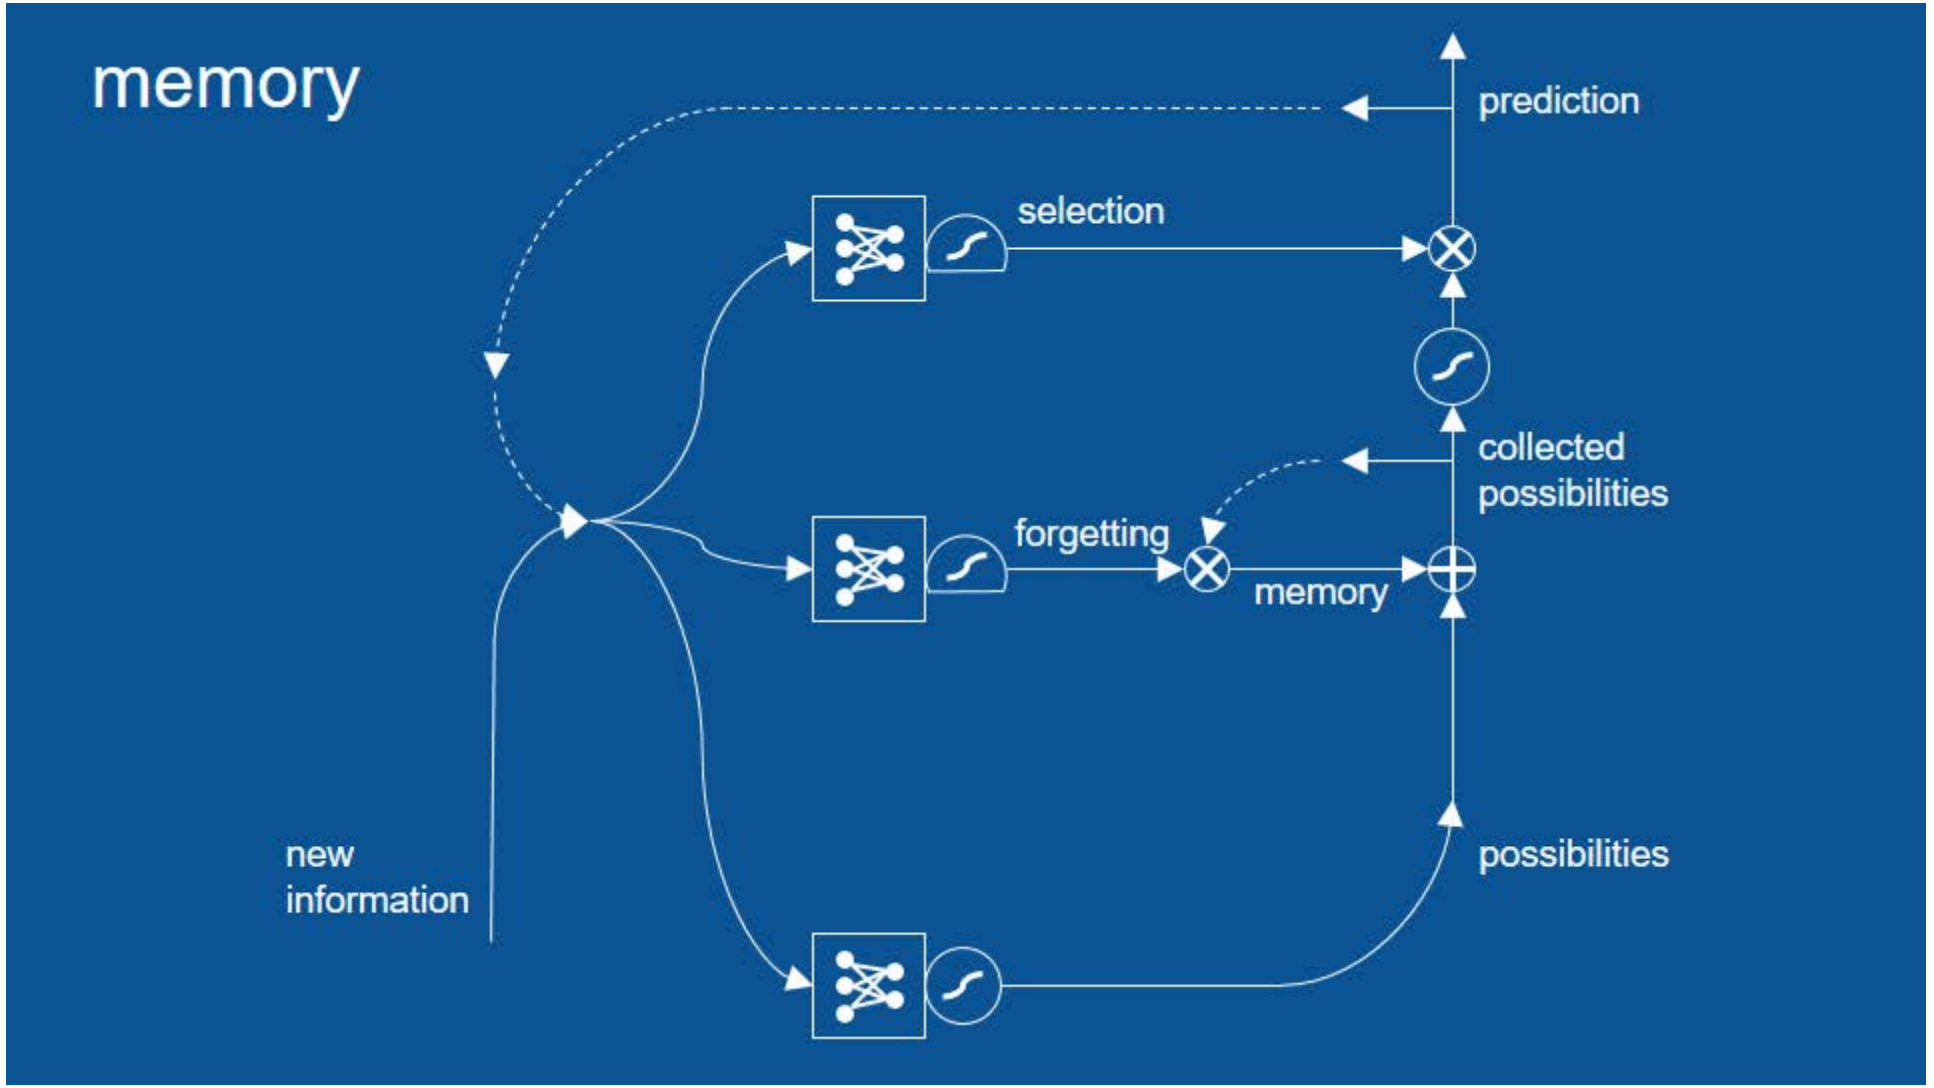

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

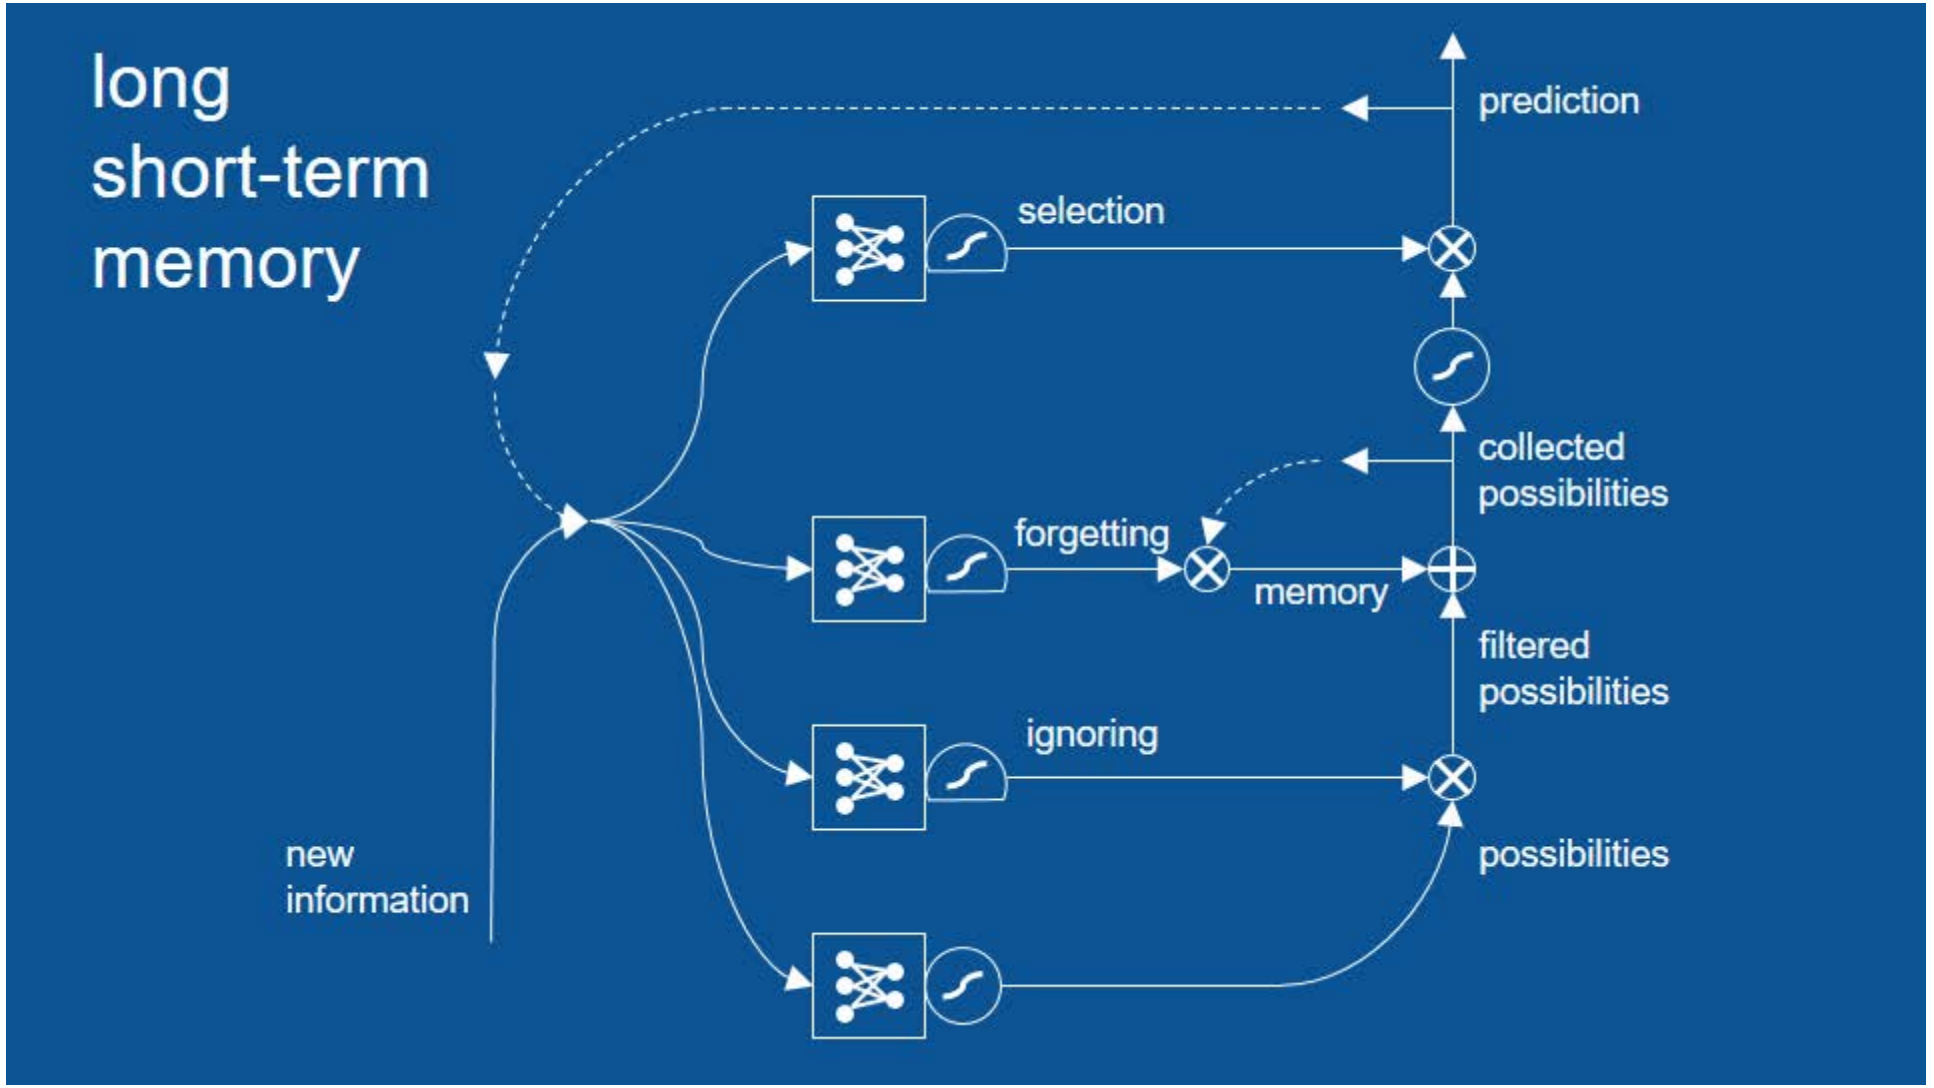

https://elham-khanche.github.io/blog/RNNs_and_LSTM/

Long Short-Term Memory

- Forget gate controls erase operation
- Input gate controls write operation
- Output gate controls the read operation

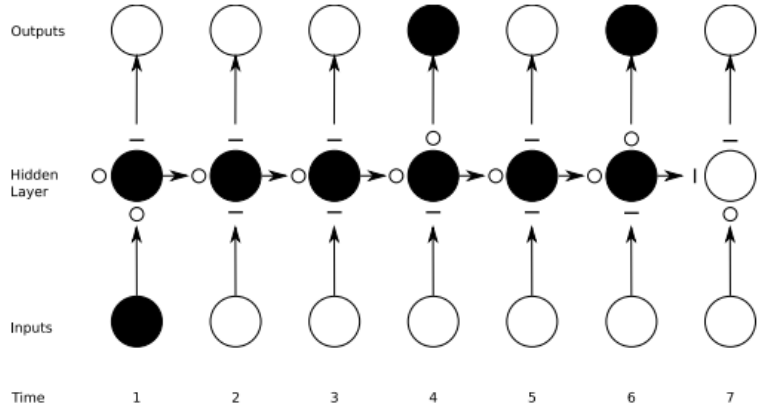

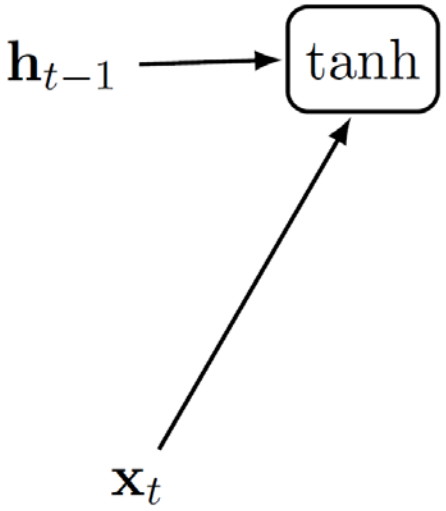

Long Short-Term Memory

- Forget gate controls erase operation
- Input gate controls write operation
- Output gate controls the read operation

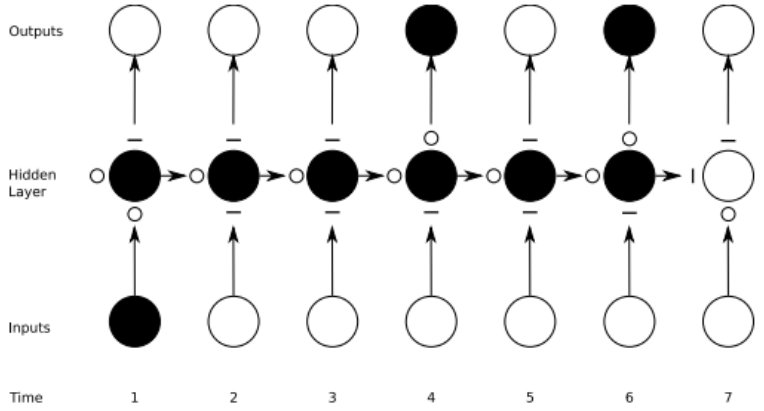

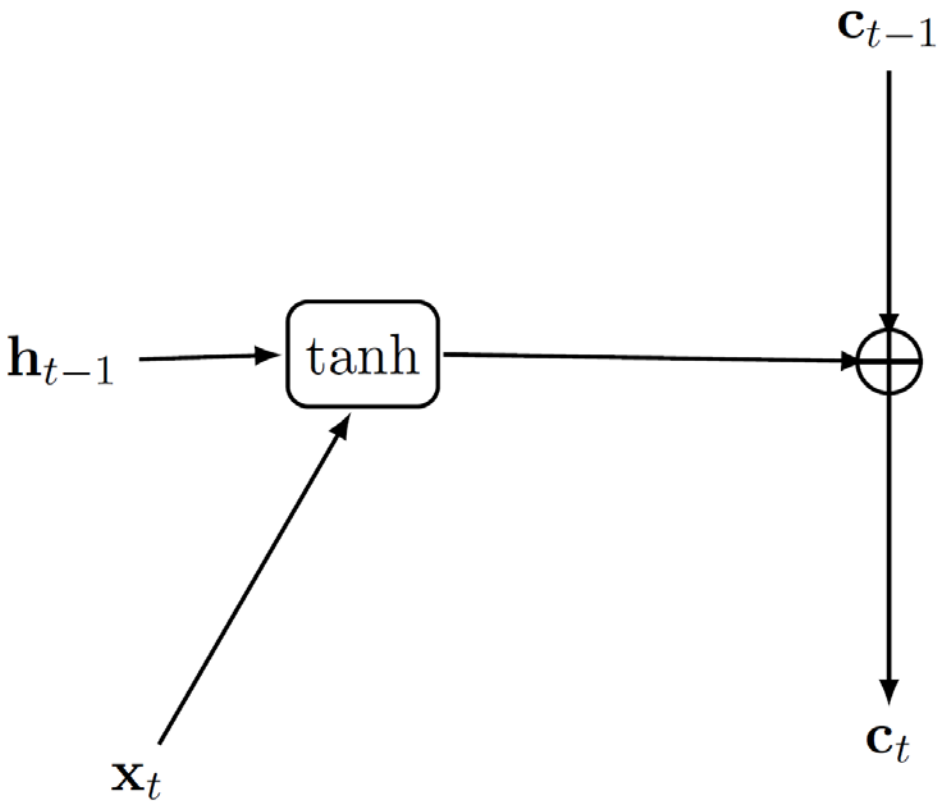

$$\tilde{\mathbf{c}}_t = \tanh(W\mathbf{h}_{t-1} + U\mathbf{x}_t)$$
$$\mathbf{c}_t = \mathbf{c}_{t-1} + \tilde{c}_t$$

Long Short-Term Memory

- Forget gate controls erase operation
- Input gate controls write operation
- Output gate controls the read operation

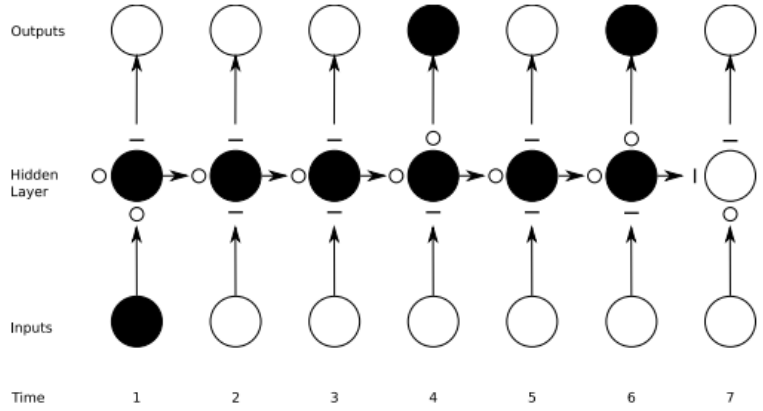

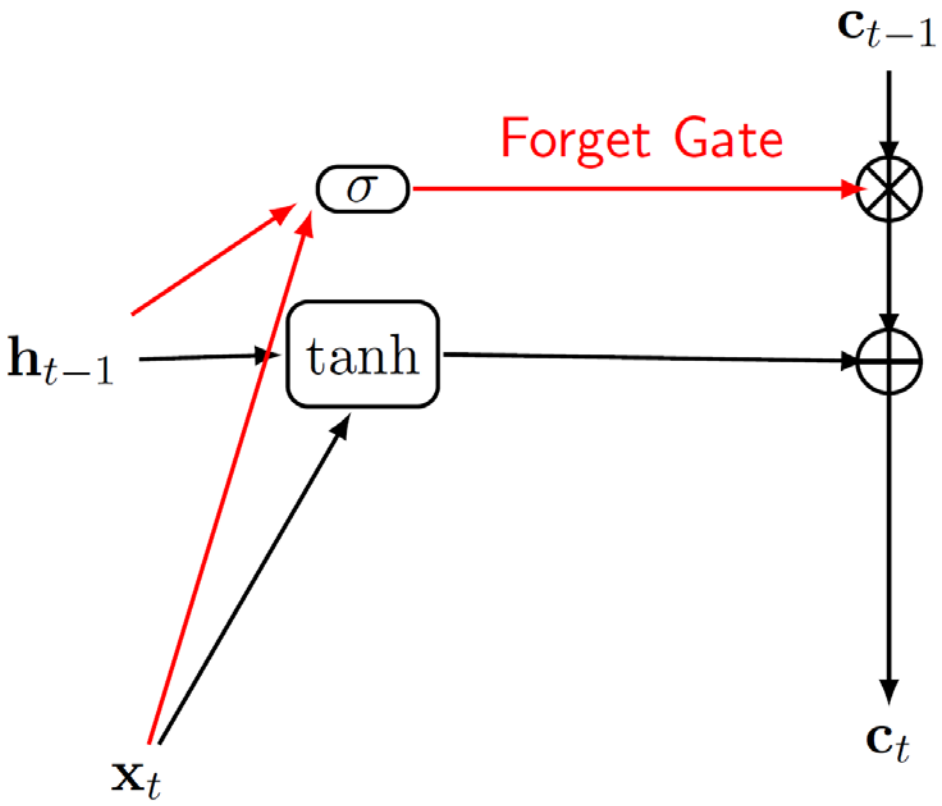

$$f_t = \sigma(W_f \mathbf{h}_{t-1} + U_f \mathbf{x}_t)$$

$$\tilde{\mathbf{c}}_t = \tanh(W \mathbf{h}_{t-1} + U \mathbf{x}_t)$$

$$\mathbf{c}_t = f_t \odot \mathbf{c}_{t-1} + \tilde{c}_t$$

Long Short-Term Memory

- Forget gate controls erase operation
- Input gate controls write operation
- Output gate controls the read operation

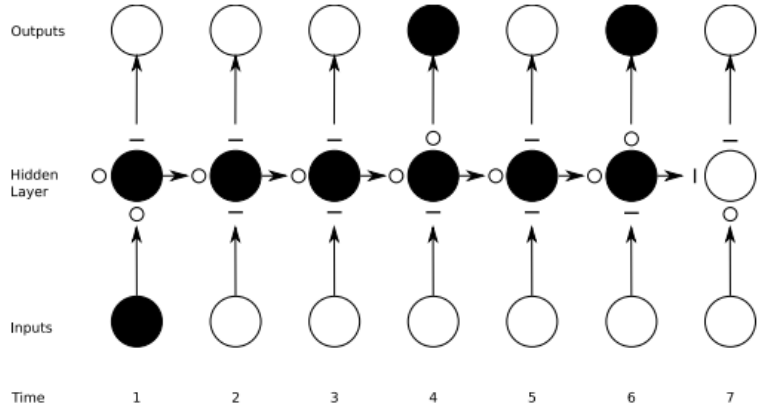

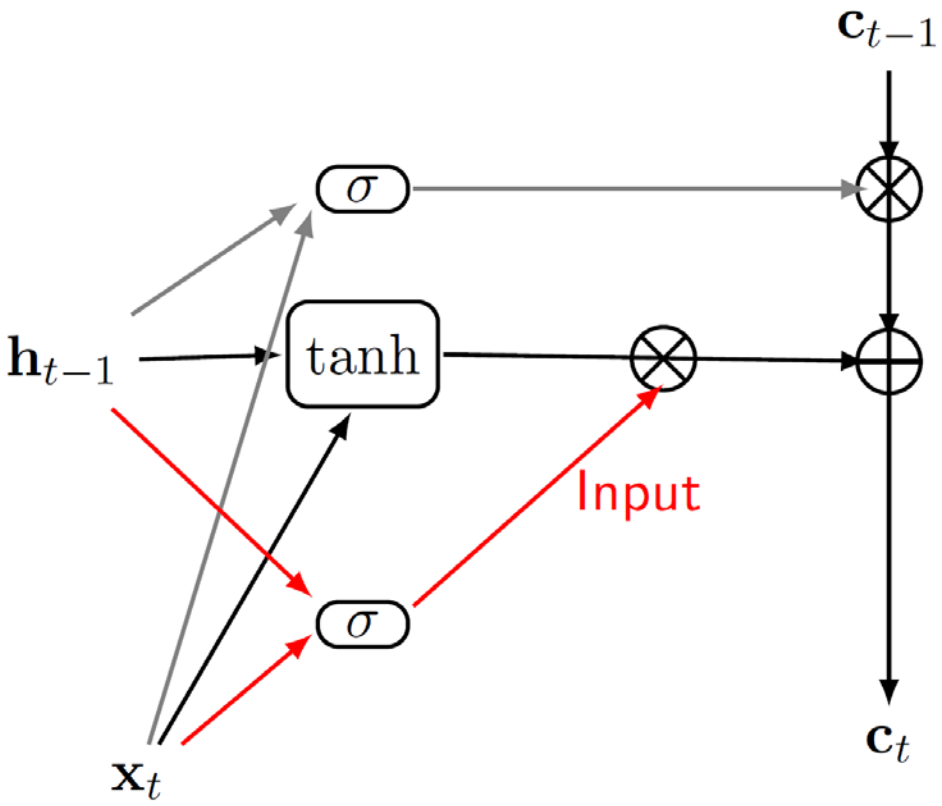

$$f_t = \sigma(W_f \mathbf{h}_{t-1} + U_f \mathbf{x}_t)$$
$$i_t = \sigma(W_i \mathbf{h}_{t-1} + U_i \mathbf{x}_t)$$

$$\tilde{\mathbf{c}}_t = \tanh(W \mathbf{h}_{t-1} + U \mathbf{x}_t)$$
$$\mathbf{c}_t = f_t \odot \mathbf{c}_{t-1} + i_t \odot \tilde{c}_t$$

Long Short-Term Memory

- Forget gate controls erase operation
- Input gate controls write operation
- Output gate controls the read operation

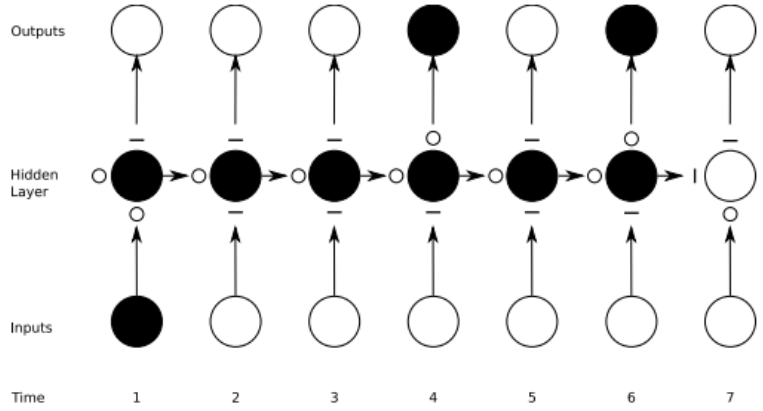

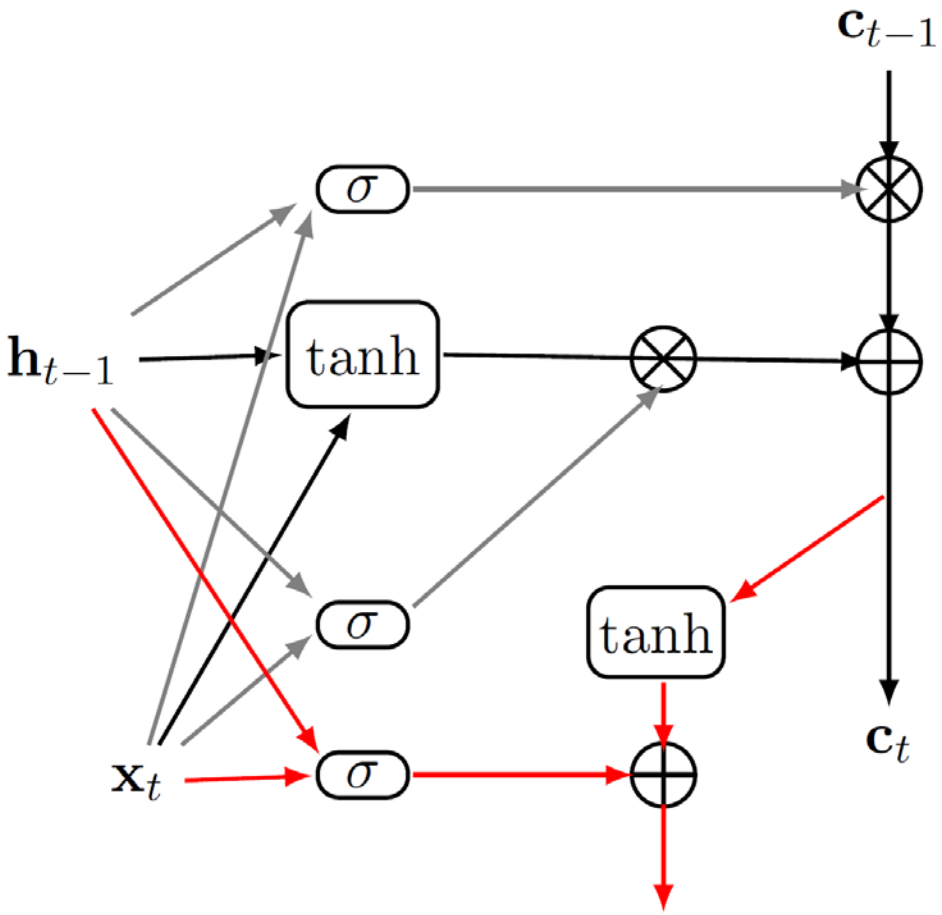

$$f_t = \sigma(W_f \mathbf{h}_{t-1} + U_f \mathbf{x}_t)$$
$$i_t = \sigma(W_i \mathbf{h}_{t-1} + U_i \mathbf{x}_t)$$
$$o_t = \sigma(W_o \mathbf{h}_{t-1} + U_o \mathbf{x}_t)$$

$$\tilde{\mathbf{c}}_t = \tanh(W \mathbf{h}_{t-1} + U \mathbf{x}_t)$$
$$\mathbf{c}_t = f_t \odot \mathbf{c}_{t-1} + i_t \odot \tilde{c}_t$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tanh(\mathbf{c}_t)$$

Weakness of RNNs

- Sequential computation is slow
- Vanishing and exploding gradients are still problematic
- Long-term credit assignment is difficult

LSTM Implementation

Time Series Data and RNN

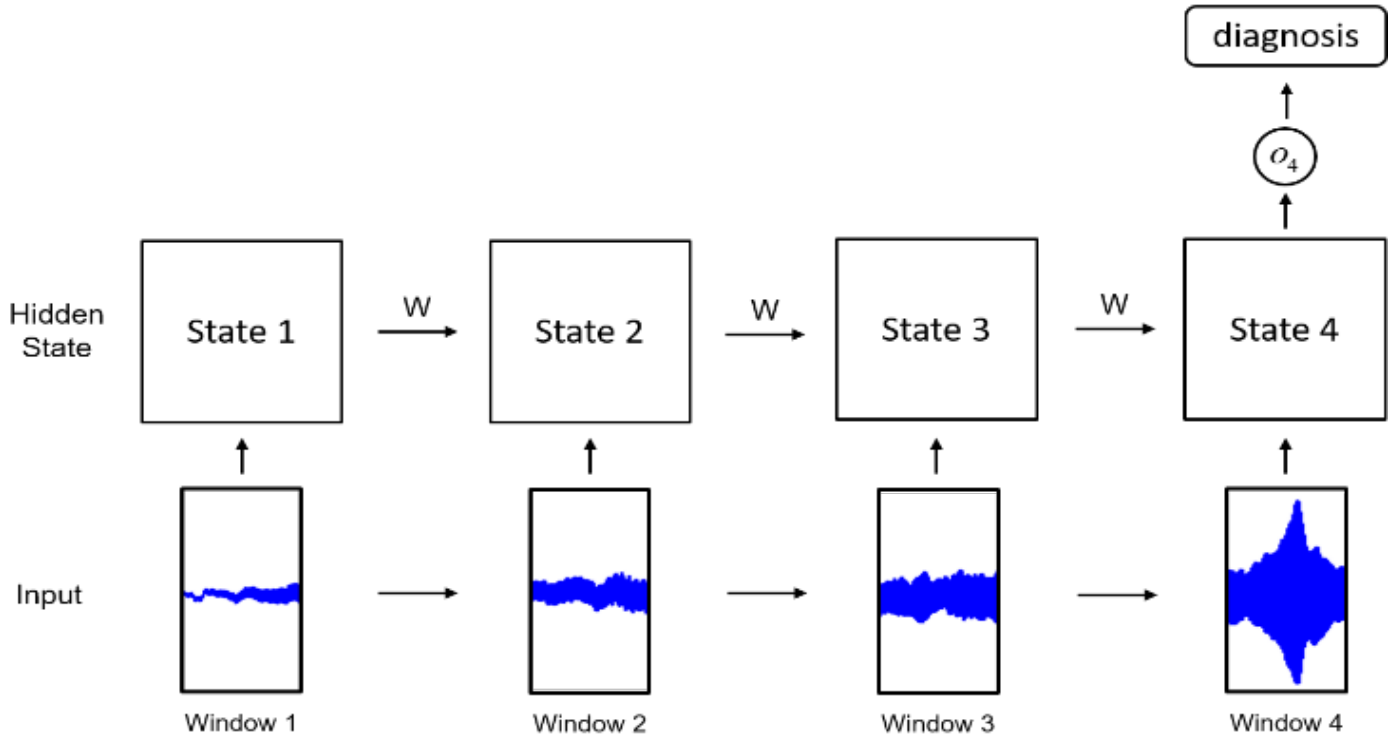

RNN for Classification

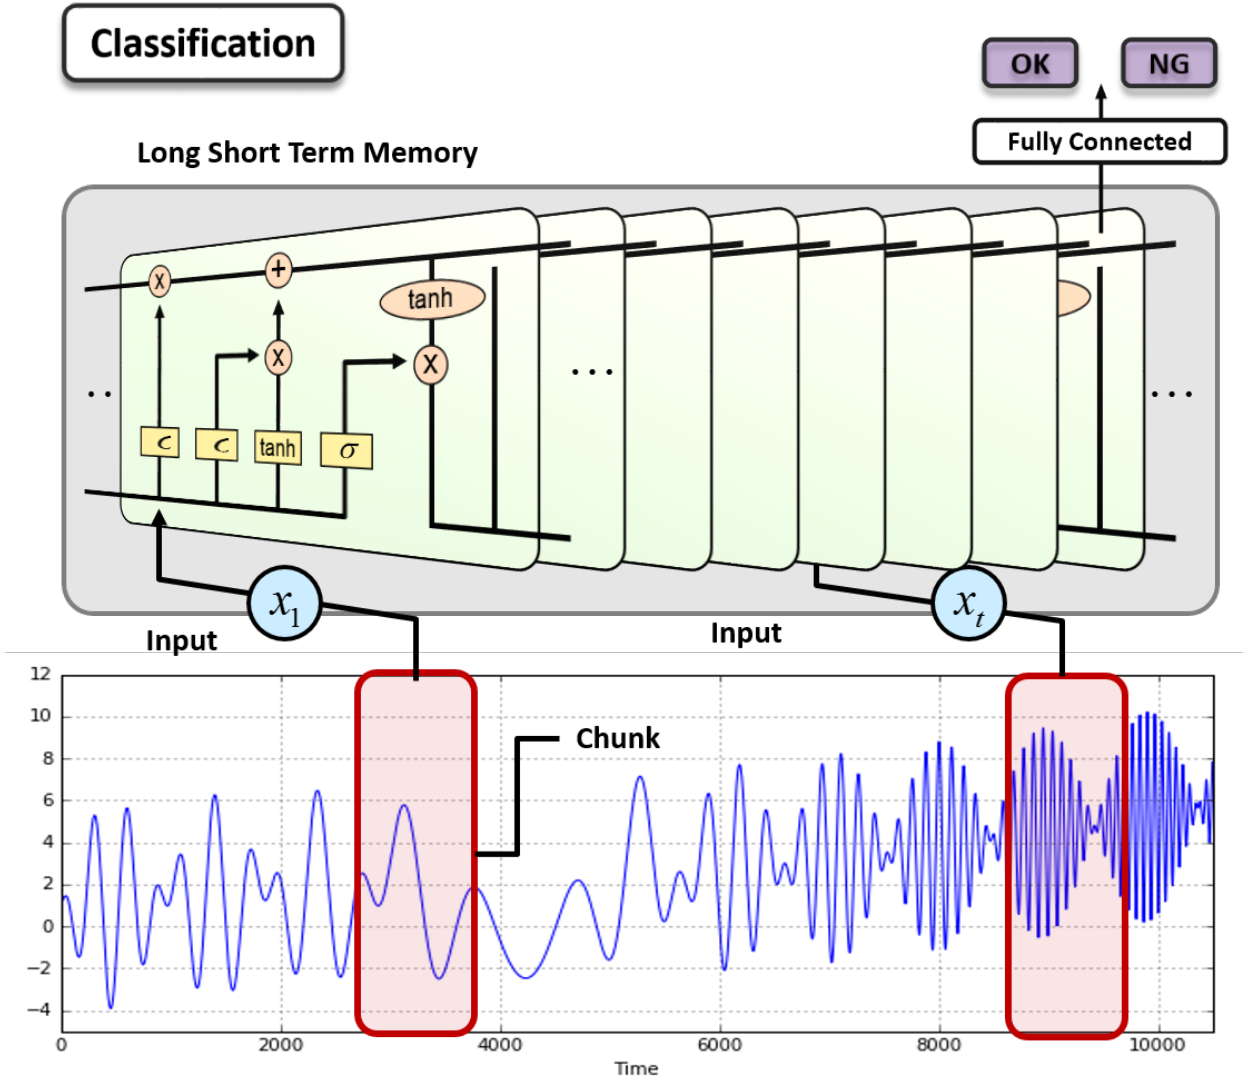

RNN for Prediction

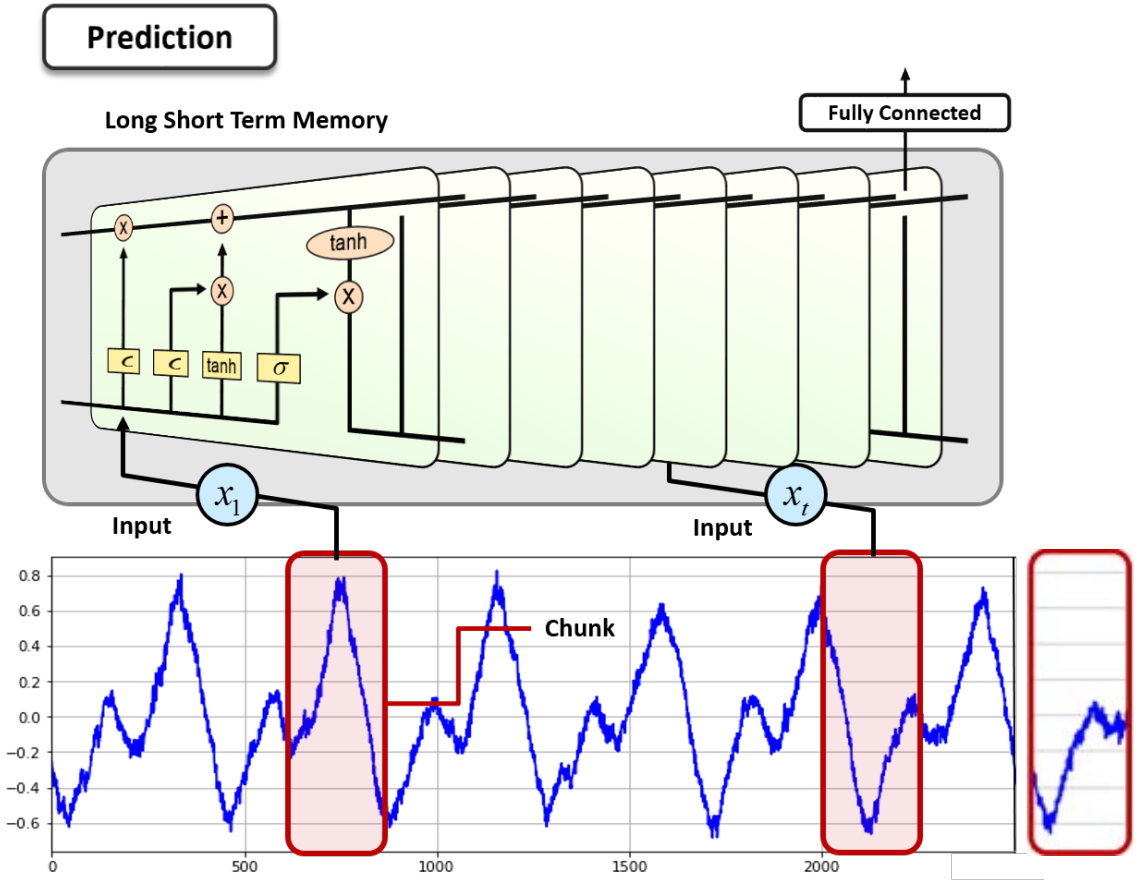

LSTM with TensorFlow

- An example for predicting a next piece of an acceleration signal
- Regression problem

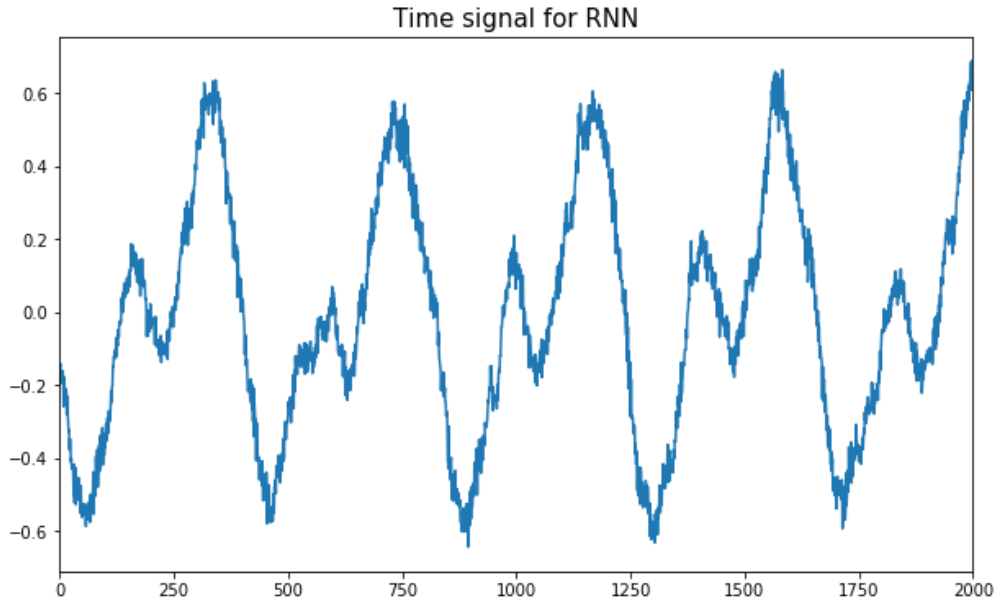

RNN Structure

```
n_step = 25
n_input = 100

## LSTM shape
n_lstm1 = 100
n_lstm2 = 100

## Fully connected
n_hidden = 100
n_output = 100
```

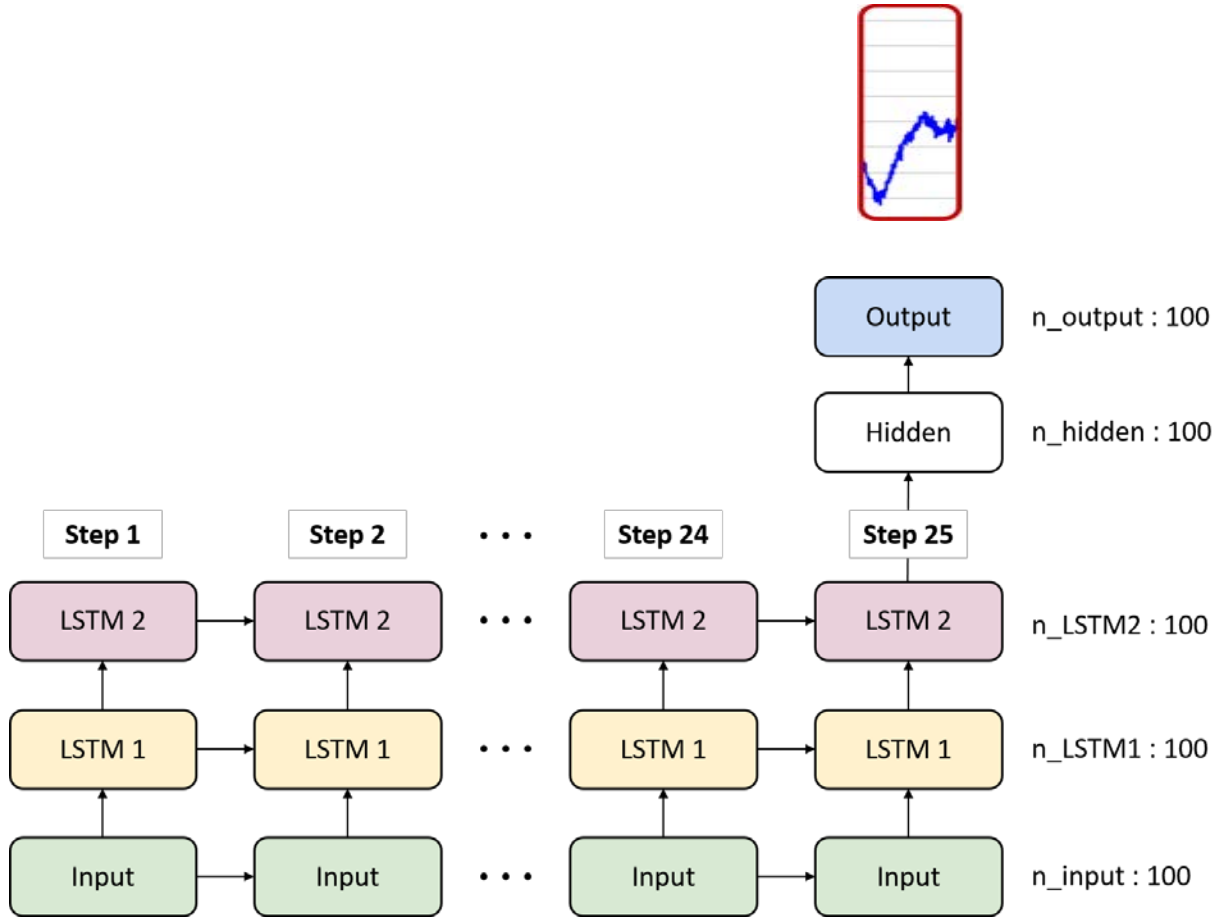

LSTM, Weights and Biases

- LSTM Cell
  - Do not need to define weights and biases of LSTM cells

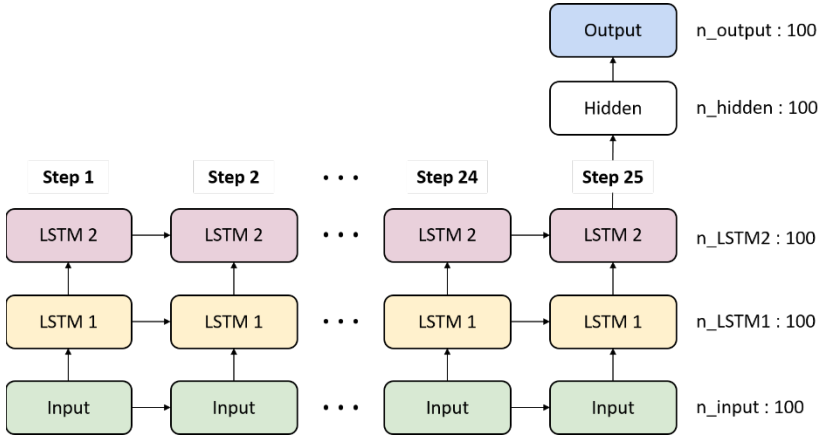

- Fully connected
  - Define parameters based on the predefined layer size
  - Initialize with a normal distribution with $\mu = 0$ and $\sigma = 0.01$

```
weights = {
    'hidden' : tf.Variable(tf.random_normal([n_lstm2, n_hidden], stddev = 0.01)),
    'output' : tf.Variable(tf.random_normal([n_hidden, n_output], stddev = 0.01))
}

biases = {
    'hidden' : tf.Variable(tf.random_normal([n_hidden], stddev = 0.01)),
    'output' : tf.Variable(tf.random_normal([n_output], stddev = 0.01))
}

x = tf.placeholder(tf.float32, [None, n_step, n_input])
y = tf.placeholder(tf.float32, [None, n_output])
```

Build a Model

- First, define the LSTM cells
- Second, compute hidden state (h) and LSTM cell (c) with the predefined LSTM cell and input

```
def build_model(x, weights, biases):
    with tf.variable_scope('rnn'):
        # Build RNN network
        with tf.variable_scope('lstm1'):
            lstm1 = tf.nn.rnn_cell.LSTMCell(n_lstm1)
            h1, c1 = tf.nn.dynamic_rnn(lstm1, x, dtype = tf.float32)
        with tf.variable_scope('lstm2'):
            lstm2 = tf.nn.rnn_cell.LSTMCell(n_lstm2)
            h2, c2 = tf.nn.dynamic_rnn(lstm2, h1, dtype = tf.float32)

        # Build classifier
        hidden = tf.add(tf.matmul(h2[:,-1,:], weights['hidden']), biases['hidden'])
        hidden = tf.nn.relu(hidden)
        output = tf.add(tf.matmul(hidden, weights['output']), biases['output'])
        return output
```

Cost, Initializer and Optimizer

- Loss
  - Regression: Squared loss

$$\frac{1}{N}\sum_{i=1}^{N}(\hat{y}^{(i)} - y^{(i)})^2$$

- Initializer
  - Initialize all the empty variables
- Optimizer
  - AdamOptimizer: the most popular optimize

```
LR = 0.0001

pred = build_model(x, weights, biases)
loss = tf.square(tf.subtract(y, pred))
loss = tf.reduce_mean(loss)

optm = tf.train.AdamOptimizer(LR).minimize(loss)
init = tf.global_variables_initializer()

sess = tf.Session()
```

Prediction Example

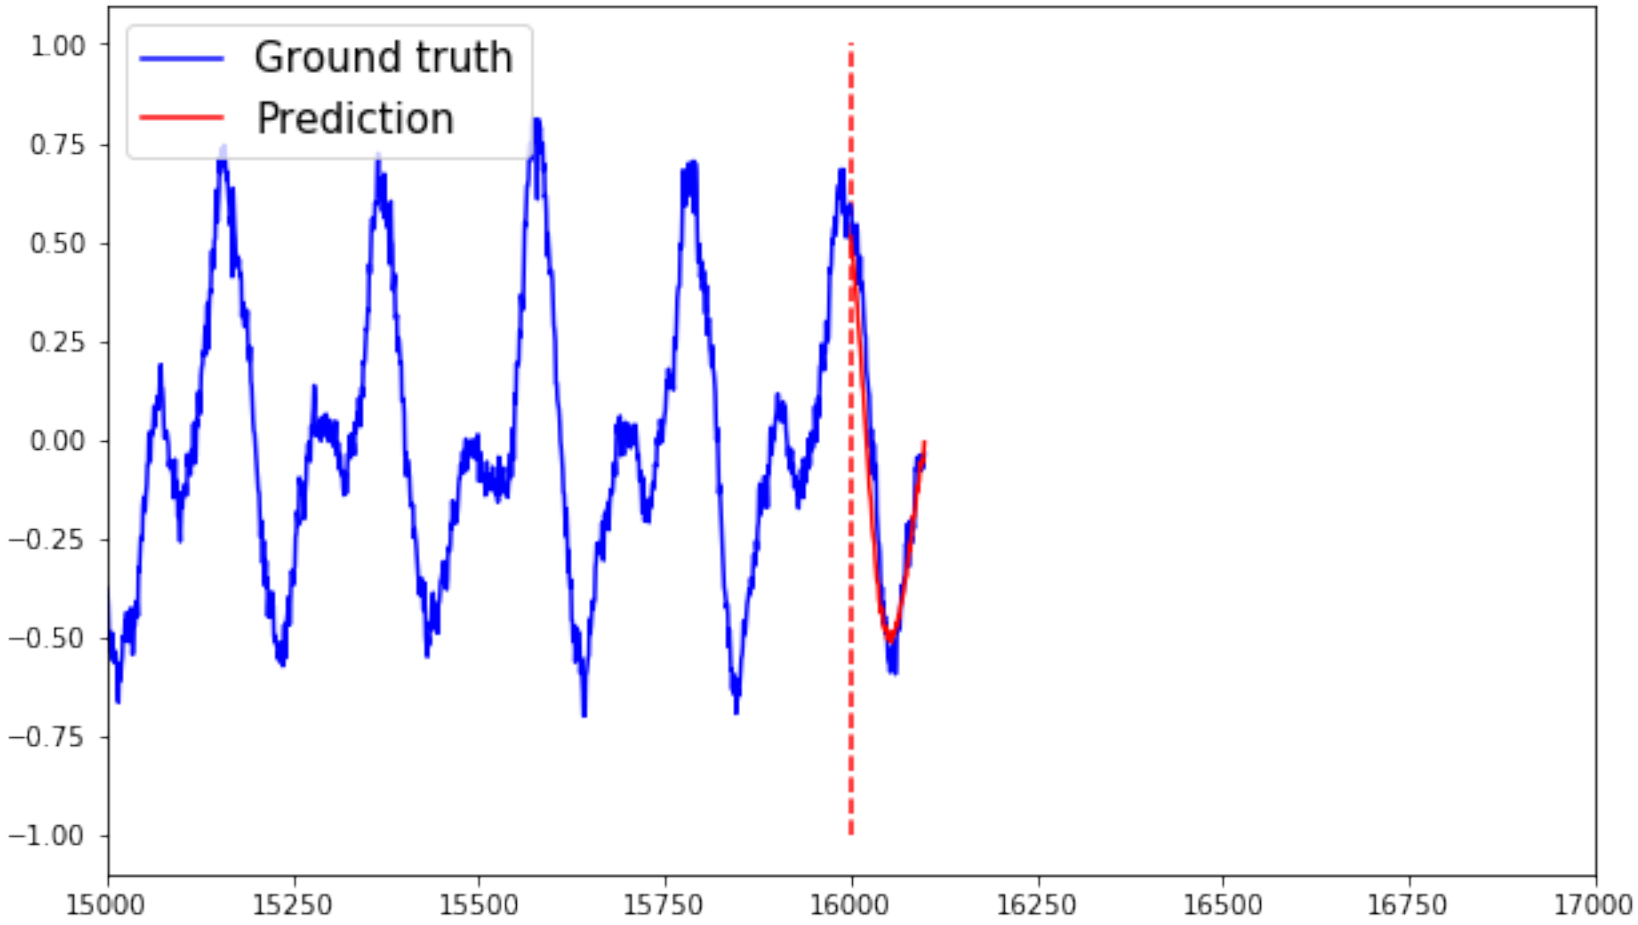
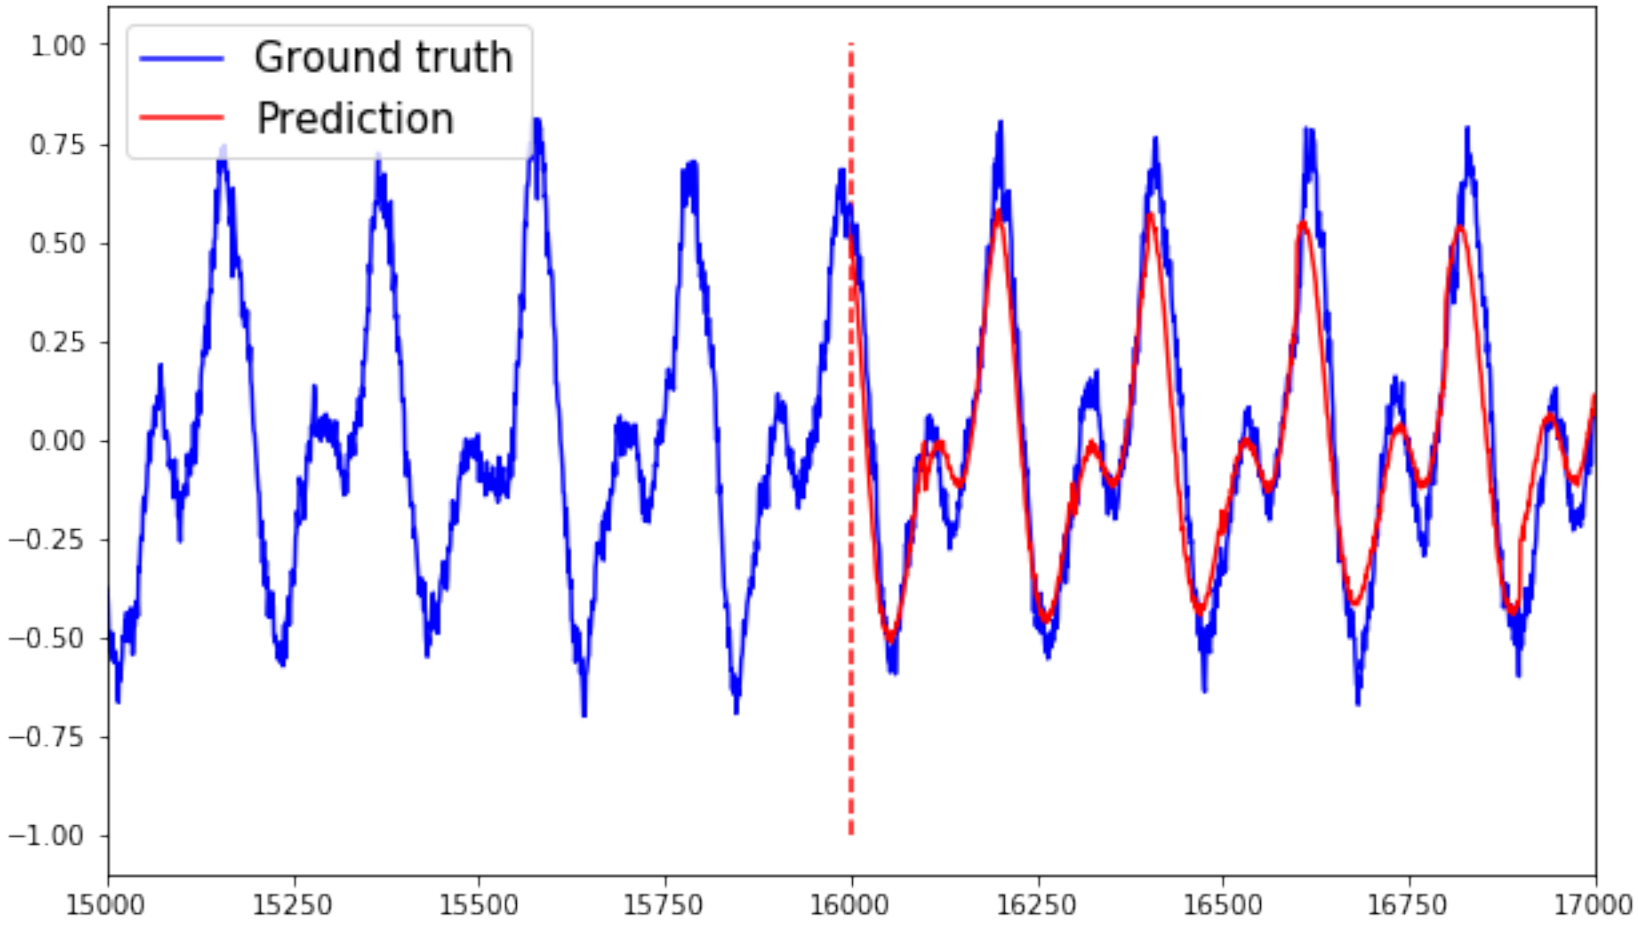
Prediction Example